

MicroBlaze Processor Reference Guide

***Embedded Development Kit
EDK 12.1***

UG081 (v11.0)



Xilinx is providing this product documentation, hereinafter "Information," to you "AS IS" with no warranty of any kind, express or implied. Xilinx makes no representation that the Information, or any particular implementation thereof, is free from any claims of infringement. You are responsible for obtaining any rights you may require for any implementation based on the Information. All specifications are subject to change without notice.

XILINX EXPRESSLY DISCLAIMS ANY WARRANTY WHATSOEVER WITH RESPECT TO THE ADEQUACY OF THE INFORMATION OR ANY IMPLEMENTATION BASED THEREON, INCLUDING BUT NOT LIMITED TO ANY WARRANTIES OR REPRESENTATIONS THAT THIS IMPLEMENTATION IS FREE FROM CLAIMS OF INFRINGEMENT AND ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Except as stated herein, none of the Information may be copied, reproduced, distributed, republished, downloaded, displayed, posted, or transmitted in any form or by any means including, but not limited to, electronic, mechanical, photocopying, recording, or otherwise, without the prior written consent of Xilinx.

© 2010 Xilinx, Inc. XILINX, the Xilinx logo, Virtex, Spartan, ISE, and other designated brands included herein are trademarks of Xilinx in the United States and other countries. All other trademarks are the property of their respective owners.

MicroBlaze Processor Reference Guide

UG081 (v11.0)

The following table shows the revision history for this document.

| Date | Version | Revision |
|----------|---------|---|
| 10/01/02 | 1.0 | Xilinx EDK 3.1 release |
| 03/11/03 | 2.0 | Xilinx EDK 3.2 release |
| 09/24/03 | 3.0 | Xilinx EDK 6.1 release |
| 02/20/04 | 3.1 | Xilinx EDK 6.2 release |
| 08/24/04 | 4.0 | Xilinx EDK 6.3 release |
| 09/21/04 | 4.1 | Minor corrections for EDK 6.3 SP1 release |
| 11/18/04 | 4.2 | Minor corrections for EDK 6.3 SP2 release |
| 01/20/05 | 5.0 | Xilinx EDK 7.1 release |
| 04/02/05 | 5.1 | Minor corrections for EDK 7.1 SP1 release |
| 05/09/05 | 5.2 | Minor corrections for EDK 7.1 SP2 release |
| 10/05/05 | 5.3 | Minor corrections for EDK 8.1 release |
| 02/21/06 | 5.4 | Corrections for EDK 8.1 SP2 release |
| 06/01/06 | 6.0 | Xilinx EDK 8.2 release |
| 07/24/06 | 6.1 | Minor corrections for EDK 8.2 SP1 release |
| 08/21/06 | 6.2 | Minor corrections for EDK 8.2 SP2 release |
| 08/29/06 | 6.3 | Minor corrections for EDK 8.2 SP2 release |
| 09/15/06 | 7.0 | Xilinx EDK 9.1 release |
| 02/22/07 | 7.1 | Minor corrections for EDK 9.1 SP1 release |

| Date | Version | Revision |
|----------|---------|--|
| 03/27/07 | 7.2 | Minor corrections for EDK 9.1 SP2 release |
| 06/25/07 | 8.0 | Xilinx EDK 9.2 release |
| 10/12/07 | 8.1 | Minor corrections for EDK 9.2 SP2 release |
| 01/17/08 | 9.0 | Xilinx EDK 10.1 release |
| 03/04/08 | 9.1 | Minor corrections for EDK 10.1 SP1 release |
| 05/14/08 | 9.2 | Minor corrections for EDK 10.1 SP2 release |
| 07/14/08 | 9.3 | Minor corrections for EDK 10.1 SP3 release |
| 02/04/09 | 10.0 | Xilinx EDK 11.1 release |
| 04/15/09 | 10.1 | Xilinx EDK 11.2 release |
| 05/28/09 | 10.2 | Xilinx EDK 11.3 release |
| 10/26/09 | 10.3 | Xilinx EDK 11.4 release |
| 04/19/10 | 11.0 | Xilinx EDK 12.1 release |

Table of Contents

Preface: About This Guide

| | |
|-----------------------|----|
| Guide Contents | 9 |
| Conventions | 10 |
| Typographical | 10 |
| Online Document | 11 |

Chapter 1: MicroBlaze Architecture

| | |
|---|----|
| Overview | 14 |
| Features | 14 |
| Data Types and Endianness | 16 |
| Instructions | 17 |
| Instruction Summary | 17 |
| Semaphore Synchronization | 25 |
| Registers | 27 |
| General Purpose Registers | 27 |
| Special Purpose Registers | 28 |
| Pipeline Architecture | 51 |
| Three Stage Pipeline | 51 |
| Five Stage Pipeline | 51 |
| Branches | 51 |
| Memory Architecture | 53 |
| Privileged Instructions | 54 |
| Virtual-Memory Management | 55 |
| Real Mode | 55 |
| Virtual Mode | 56 |
| Translation Look-Aside Buffer | 57 |
| Access Protection | 62 |
| UTLB Management | 63 |
| Recording Page Access and Page Modification | 64 |
| Reset, Interrupts, Exceptions, and Break | 65 |
| Reset | 65 |
| Hardware Exceptions | 66 |
| Breaks | 68 |
| Interrupt | 69 |
| User Vector (Exception) | 70 |
| Instruction Cache | 71 |
| Overview | 71 |
| General Instruction Cache Functionality | 71 |
| Instruction Cache Operation | 72 |
| Instruction Cache Software Support | 73 |
| Data Cache | 73 |
| Overview | 73 |
| General Data Cache Functionality | 73 |
| Data Cache Operation | 74 |

| | |
|---------------------------------------|-----------|
| Data Cache Software Support | 75 |
| Floating Point Unit (FPU)..... | 76 |
| Overview | 76 |
| Format | 76 |
| Rounding..... | 77 |
| Operations | 77 |
| Exceptions | 77 |
| Software Support | 78 |
| Fast Simplex Link (FSL)..... | 80 |
| Hardware Acceleration using FSL | 80 |
| Debug and Trace | 81 |
| Debug Overview | 81 |
| Trace Overview | 81 |

Chapter 2: MicroBlaze Signal Interface Description

| | |
|--|------------|
| Overview | 83 |
| Features | 83 |
| MicroBlaze I/O Overview | 84 |
| Processor Local Bus (PLB) Interface Description | 88 |
| Local Memory Bus (LMB) Interface Description | 88 |
| LMB Signal Interface | 89 |
| LMB Transactions | 91 |
| Read and Write Data Steering | 93 |
| Fast Simplex Link (FSL) Interface Description | 94 |
| Master FSL Signal Interface | 94 |
| Slave FSL Signal Interface | 94 |
| FSL Transactions | 95 |
| Direct FSL Connections | 95 |
| Xilinx CacheLink (XCL) Interface Description | 96 |
| CacheLink Signal Interface | 97 |
| CacheLink Transactions | 98 |
| Debug Interface Description | 101 |
| Trace Interface Description | 101 |
| MicroBlaze Core Configurability | 104 |

Chapter 3: MicroBlaze Application Binary Interface

| | |
|---|------------|
| Data Types | 111 |
| Register Usage Conventions | 112 |
| Stack Convention | 113 |
| Calling Convention | 115 |
| Memory Model | 115 |
| Small Data Area | 115 |
| Data Area | 115 |
| Common Un-Initialized Area | 115 |
| Literals or Constants | 115 |
| Interrupt and Exception Handling | 116 |

Chapter 4: MicroBlaze Instruction Set Architecture

| | |
|---------------------------|-----|
| Notation | 117 |
| Formats | 119 |
| Instructions | 119 |

About This Guide

The MicroBlaze™ Processor Reference Guide provides information about the 32-bit soft processor, MicroBlaze, which is part of the Embedded Processor Development Kit (EDK). The document is intended as a guide to the MicroBlaze hardware architecture.

Guide Contents

This guide contains the following chapters:

- Chapter 1, “[MicroBlaze Architecture](#)”, contains an overview of MicroBlaze features as well as information on Big-Endian bit-reversed format, 32-bit general purpose registers, cache software support, and Fast Simplex Link interfaces.
- Chapter 2, “[MicroBlaze Signal Interface Description](#)”, describes the types of signal interfaces that can be used to connect MicroBlaze.
- Chapter 3, “[MicroBlaze Application Binary Interface](#)”, describes the Application Binary Interface important for developing software in assembly language for the soft processor.
- Chapter 4, “[MicroBlaze Instruction Set Architecture](#)”, provides notation, formats, and instructions for the Instruction Set Architecture of MicroBlaze.

For additional information, go to <http://www.xilinx.com/support/mysupport.htm>. The following table lists some of the resources you can access directly using the provided URLs.

| Resource | Description/URL |
|-----------------|--|
| Tutorials | Tutorials covering Xilinx® design flows, from design entry to verification and debugging. http://www.xilinx.com/support/techsup/tutorials/index.htm |
| Answer Browser | Database of Xilinx solution records. http://www.xilinx.com/support/answers/index.htm |
| Problem Solvers | Interactive tools that allow you to troubleshoot your design issues. http://www.xilinx.com/support/troubleshoot/psolvers.htm |
| GNU Manuals | The entire set of GNU manuals. http://www.gnu.org/manual |

Conventions

This document uses the following conventions. An example illustrates each convention.

Typographical

The following typographical conventions are used in this document:

| Convention | Meaning or Use | Example |
|-------------------------|---|--|
| Courier font | Messages, prompts, and program files that the system displays. | <code>speed grade: - 100</code> |
| Courier bold | Literal commands that you enter in a syntactical statement. | <code>ngdbuild design_name</code> |
| Helvetica bold | Commands that you select from a menu. | File → Open |
| | Keyboard shortcuts | Ctrl+C |
| Italic font | Variables in a syntax statement for which you must supply values. | <code>ngdbuild design_name</code> |
| | References to other manuals. | See the <i>Development System Reference Guide</i> for more information. |
| | Emphasis in text. | If a wire is drawn so that it overlaps the pin of a symbol, the two nets are <i>not</i> connected. |
| Square brackets [] | An optional entry or parameter. However, in bus specifications, such as bus [7 : 0] , they are required. | <code>ngdbuild [option_name] design_name</code> |
| Braces { } | A list of items from which you must choose one or more. | <code>lowpwr ={on off}</code> |
| Vertical bar | Separates items in a list of choices. | <code>lowpwr ={on off}</code> |
| Vertical ellipsis . . . | Repetitive material that has been omitted | IOB #1: Name = QOUT' IOB #2: Name = CLKIN' . . . |
| Horizontal ellipsis ... | Repetitive material that has been omitted | <code>allow block block_name loc1 loc2 ... locn;</code> |

Online Document

The following conventions are used in this document:

| Convention | Meaning or Use | Example |
|--|--|---|
| Blue text | Cross-reference link to a location in the current document | See the section “ Additional Resources ” for details. Refer to “ Title Formats ” in Chapter 1 for details. |
| <u>Blue, underlined text</u> | Hyperlink to a web-site (URL) | Go to http://www.xilinx.com for the latest speed files. |

MicroBlaze Architecture

This chapter contains an overview of MicroBlaze™ features and detailed information on MicroBlaze architecture including Big-Endian bit-reversed format, 32-bit general purpose registers, virtual-memory management, cache software support, and Fast Simplex Link (FSL) interfaces.

This chapter has the following sections:

- “Overview”
- “Data Types and Endianess”
- “Instructions”
- “Registers”
- “Pipeline Architecture”
- “Memory Architecture”
- “Privileged Instructions”
- “Virtual-Memory Management”
- “Reset, Interrupts, Exceptions, and Break”
- “Instruction Cache”
- “Data Cache”
- “Floating Point Unit (FPU)”
- “Fast Simplex Link (FSL)”
- “Debug and Trace”

Overview

The MicroBlaze™ embedded processor soft core is a reduced instruction set computer (RISC) optimized for implementation in Xilinx® Field Programmable Gate Arrays (FPGAs). Figure 1-1 shows a functional block diagram of the MicroBlaze core.

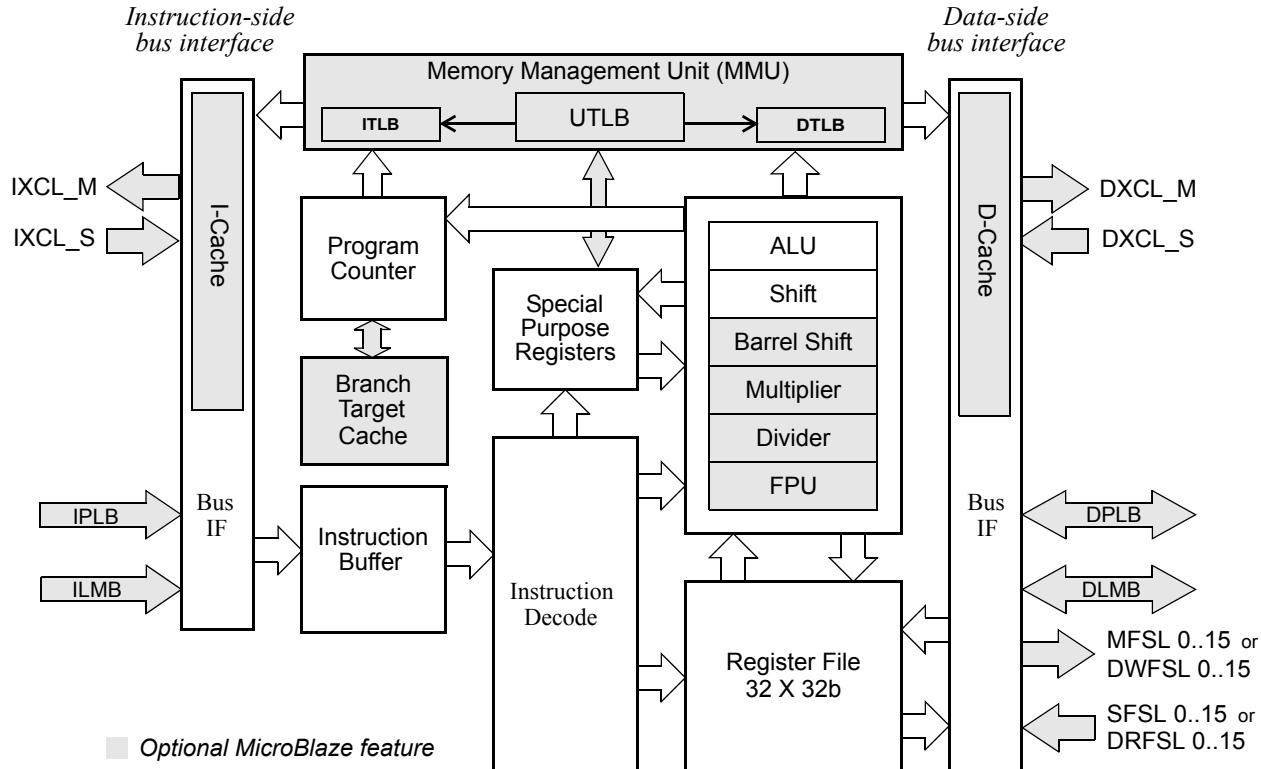


Figure 1-1: MicroBlaze Core Block Diagram

Features

The MicroBlaze soft core processor is highly configurable, allowing you to select a specific set of features required by your design.

The fixed feature set of the processor includes:

- Thirty-two 32-bit general purpose registers
- 32-bit instruction word with three operands and two addressing modes
- 32-bit address bus
- Single issue pipeline

In addition to these fixed features, the MicroBlaze processor is parameterized to allow selective enabling of additional functionality. Older (deprecated) versions of MicroBlaze support a subset of the optional features described in this manual. Only the latest (preferred) version of MicroBlaze (v7.30) supports all options.

Xilinx recommends that all new designs use the latest **preferred** version of the MicroBlaze processor.

Table 1-1, page 15 provides an overview of the configurable features by MicroBlaze versions.

Table 1-1: Configurable Feature Overview by MicroBlaze Version

| Feature | MicroBlaze Versions | | | |
|---|---------------------|----------|------------|-----------|
| | v7.00 | v7.10 | v7.20 | v7.30 |
| Version Status | obsolete | obsolete | deprecated | preferred |
| Processor pipeline depth | 3/5 | 3/5 | 3/5 | 3/5 |
| On-chip Peripheral Bus (OPB) data side interface | option | option | option | No |
| On-chip Peripheral Bus (OPB) instruction side interface | option | option | option | No |
| Local Memory Bus (LMB) data side interface | option | option | option | option |
| Local Memory Bus (LMB) instruction side interface | option | option | option | option |
| Hardware barrel shifter | option | option | option | option |
| Hardware divider | option | option | option | option |
| Hardware debug logic | option | option | option | option |
| Fast Simplex Link (FSL) interfaces | 0-15 | 0-15 | 0-15 | 0-15 |
| Machine status set and clear instructions | option | option | option | option |
| Instruction cache over IOPB interface | No | No | No | No |
| Data cache over DOPB interface | No | No | No | No |
| Instruction cache over Cache Link (IXCL) interface | option | option | option | option |
| Data cache over Cache Link (DXCL) interface | option | option | option | option |
| 4 or 8-word cache line | option | option | option | option |
| Hardware exception support | option | option | option | option |
| Pattern compare instructions | option | option | option | option |
| Floating point unit (FPU) | option | option | option | option |
| Disable hardware multiplier ¹ | option | option | option | option |
| Hardware debug readable ESR and EAR | Yes | Yes | Yes | Yes |
| Processor Version Register (PVR) | option | option | option | option |
| Area or speed optimized | option | option | option | option |
| Hardware multiplier 64-bit result | option | option | option | option |
| LUT cache memory | option | option | option | option |
| Processor Local Bus (PLB) data side interface | option | option | option | option |
| Processor Local Bus (PLB) instruction side interface | option | option | option | option |
| Floating point conversion and square root instructions | option | option | option | option |
| Memory Management Unit (MMU) | option | option | option | option |

Table 1-1: Configurable Feature Overview by MicroBlaze Version

| Feature | MicroBlaze Versions | | | |
|---|---------------------|--------|--------|--------|
| | v7.00 | v7.10 | v7.20 | v7.30 |
| Extended Fast Simplex Link (FSL) instructions | option | option | option | option |
| Use Xilinx Cache Link for All I-Cache Memory Accesses | - | option | option | option |
| Use Xilinx Cache Link for All D-Cache Memory Accesses | - | option | option | option |
| Use Write-back Caching Policy for D-Cache | - | - | option | option |
| Cache Link (DXCL) protocol for D-Cache | - | - | option | option |
| Cache Link (IXCL) protocol for I-Cache | - | - | option | option |
| Branch Target Cache (BTC) | - | - | - | option |
| Streams for I-Cache | - | - | - | option |
| Victim handling for I-Cache | - | - | - | option |
| Victim handling for D-Cache | - | - | - | option |

1. Used in Virtex®-4 and subsequent families, for saving MUL18 and DSP48 primitives.

Data Types and Endianness

MicroBlaze uses Big-Endian bit-reversed format to represent data. The hardware supported data types for MicroBlaze are word, half word, and byte. The bit and byte organization for each type is shown in the following tables.

Table 1-2: Word Data Type

| | | | | |
|-------------------|--------|-----|-----|--------|
| Byte address | n | n+1 | n+2 | n+3 |
| Byte label | 0 | 1 | 2 | 3 |
| Byte significance | MSByte | | | LSByte |
| Bit label | 0 | | | |
| Bit significance | MSBit | | | |

Table 1-3: Half Word Data Type

| | | |
|-------------------|--------|--------|
| Byte address | n | n+1 |
| Byte label | 0 | 1 |
| Byte significance | MSByte | LSByte |
| Bit label | 0 15 | |
| Bit significance | MSBit | LSBit |

Table 1-4: Byte Data Type

| | |
|------------------|------------------|
| Byte address | n |
| Bit label | 0 7 |
| Bit significance | MSBit LSBit |

Instructions

Instruction Summary

All MicroBlaze instructions are 32 bits and are defined as either Type A or Type B. Type A instructions have up to two source register operands and one destination register operand. Type B instructions have one source register and a 16-bit immediate operand (which can be extended to 32 bits by preceding the Type B instruction with an imm instruction). Type B instructions have a single destination register operand. Instructions are provided in the following functional categories: arithmetic, logical, branch, load/store, and special. [Table 1-6](#) lists the MicroBlaze instruction set. Refer to [Chapter 4, “MicroBlaze Instruction Set Architecture”](#) for more information on these instructions. [Table 1-5](#) describes the instruction set nomenclature used in the semantics of each instruction.

Table 1-5: Instruction Set Nomenclature

| Symbol | Description |
|-----------|---|
| Ra | R0 - R31, General Purpose Register, source operand a |
| Rb | R0 - R31, General Purpose Register, source operand b |
| Rd | R0 - R31, General Purpose Register, destination operand |
| SPR[x] | Special Purpose Register number x |
| MSR | Machine Status Register = SPR[1] |
| ESR | Exception Status Register = SPR[5] |
| EAR | Exception Address Register = SPR[3] |
| FSR | Floating Point Unit Status Register = SPR[7] |
| PVRx | Processor Version Register, where x is the register number = SPR[8192 + x] |
| BTR | Branch Target Register = SPR[11] |
| PC | Execute stage Program Counter = SPR[0] |
| $x[y]$ | Bit y of register x |
| $x[y:z]$ | Bit range y to z of register x |
| \bar{x} | Bit inverted value of register x |
| Imm | 16 bit immediate value |
| Imm x | x bit immediate value |

Table 1-5: Instruction Set Nomenclature (Continued)

| Symbol | Description |
|----------------------|--|
| FSL x | 4 bit Fast Simplex Link (FSL) port designator where x is the port number |
| C | Carry flag, MSR[29] |
| Sa | Special Purpose Register, source operand |
| Sd | Special Purpose Register, destination operand |
| s(x) | Sign extend argument x to 32-bit value |
| *Addr | Memory contents at location Addr (data-size aligned) |
| $:$ = | Assignment operator |
| $=$ | Equality comparison |
| $!=$ | Inequality comparison |
| $>$ | Greater than comparison |
| \geq | Greater than or equal comparison |
| $<$ | Less than comparison |
| \leq | Less than or equal comparison |
| $+$ | Arithmetic add |
| $*$ | Arithmetic multiply |
| $/$ | Arithmetic divide |
| $>> x$ | Bit shift right x bits |
| $<< x$ | Bit shift left x bits |
| and | Logic AND |
| or | Logic OR |
| xor | Logic exclusive OR |
| op1 if cond else op2 | Perform $op1$ if condition $cond$ is true, else perform $op2$ |
| $\&$ | Concatenate. E.g. “0000100 & Imm7” is the concatenation of the fixed field “0000100” and a 7 bit immediate value. |
| signed | Operation performed on signed integer data type. All arithmetic operations are performed on signed word operands, unless otherwise specified |
| unsigned | Operation performed on unsigned integer data type |
| float | Operation performed on floating point data type |

Table 1-6: MicroBlaze Instruction Set Summary

| Type A | 0-5 | 6-10 | 11-15 | 16-20 | 21-31 | Semantics |
|-------------------|--------|------|-------|---------------------|--------------|--|
| Type B | 0-5 | 6-10 | 11-15 | 16-31 | | |
| ADD Rd,Ra,Rb | 000000 | Rd | Ra | Rb | 000000000000 | $Rd := Rb + Ra$ |
| RSUB Rd,Ra,Rb | 000001 | Rd | Ra | Rb | 000000000000 | $Rd := Rb + \overline{Ra} + 1$ |
| ADDC Rd,Ra,Rb | 000010 | Rd | Ra | Rb | 000000000000 | $Rd := Rb + Ra + C$ |
| RSUBC Rd,Ra,Rb | 000011 | Rd | Ra | Rb | 000000000000 | $Rd := Rb + \overline{Ra} + C$ |
| ADDK Rd,Ra,Rb | 000100 | Rd | Ra | Rb | 000000000000 | $Rd := Rb + Ra$ |
| RSUBK Rd,Ra,Rb | 000101 | Rd | Ra | Rb | 000000000000 | $Rd := Rb + \overline{Ra} + 1$ |
| ADDKC Rd,Ra,Rb | 000110 | Rd | Ra | Rb | 000000000000 | $Rd := Rb + Ra + C$ |
| RSUBKC Rd,Ra,Rb | 000111 | Rd | Ra | Rb | 000000000000 | $Rd := Rb + \overline{Ra} + C$ |
| CMP Rd,Ra,Rb | 000101 | Rd | Ra | Rb | 000000000001 | $Rd := Rb + \overline{Ra} + 1$ $Rd[0] := 0$ if ($Rb \geq Ra$) else $Rd[0] := 1$ |
| CMPU Rd,Ra,Rb | 000101 | Rd | Ra | Rb | 000000000011 | $Rd := Rb + \overline{Ra} + 1$ (unsigned) $Rd[0] := 0$ if ($Rb \geq Ra$, unsigned) else $Rd[0] := 1$ |
| ADDI Rd,Ra,Imm | 001000 | Rd | Ra | Imm | | $Rd := s(Imm) + Ra$ |
| RSUBI Rd,Ra,Imm | 001001 | Rd | Ra | Imm | | $Rd := s(Imm) + \overline{Ra} + 1$ |
| ADDIC Rd,Ra,Imm | 001010 | Rd | Ra | Imm | | $Rd := s(Imm) + Ra + C$ |
| RSUBIC Rd,Ra,Imm | 001011 | Rd | Ra | Imm | | $Rd := s(Imm) + \overline{Ra} + C$ |
| ADDIK Rd,Ra,Imm | 001100 | Rd | Ra | Imm | | $Rd := s(Imm) + Ra$ |
| RSUBIK Rd,Ra,Imm | 001101 | Rd | Ra | Imm | | $Rd := s(Imm) + \overline{Ra} + 1$ |
| ADDIKC Rd,Ra,Imm | 001110 | Rd | Ra | Imm | | $Rd := s(Imm) + Ra + C$ |
| RSUBIKC Rd,Ra,Imm | 001111 | Rd | Ra | Imm | | $Rd := s(Imm) + \overline{Ra} + C$ |
| MUL Rd,Ra,Rb | 010000 | Rd | Ra | Rb | 000000000000 | $Rd := Ra * Rb$ |
| MULH Rd,Ra,Rb | 010000 | Rd | Ra | Rb | 000000000001 | $Rd := (Ra * Rb) \gg 32$ (signed) |
| MULHU Rd,Ra,Rb | 010000 | Rd | Ra | Rb | 000000000011 | $Rd := (Ra * Rb) \gg 32$ (unsigned) |
| MULHSU Rd,Ra,Rb | 010000 | Rd | Ra | Rb | 000000000010 | $Rd := (Ra, signed * Rb, unsigned) \gg 32$ (signed) |
| BSRA Rd,Ra,Rb | 010001 | Rd | Ra | Rb | 010000000000 | $Rd := s(Ra \gg Rb)$ |
| BSLL Rd,Ra,Rb | 010001 | Rd | Ra | Rb | 100000000000 | $Rd := (Ra \ll Rb) \& 0$ |
| MULI Rd,Ra,Imm | 011000 | Rd | Ra | Imm | | $Rd := Ra * s(Imm)$ |
| BSRLI Rd,Ra,Imm | 011001 | Rd | Ra | 000000000000 & Imm5 | | $Rd := 0 \& (Ra \gg Imm5)$ |

Table 1-6: MicroBlaze Instruction Set Summary (Continued)

| Type A | 0-5 | 6-10 | 11-15 | 16-20 | 21-31 | Semantics |
|------------------|--------|-------|-------|--------------------|-----------------|---|
| Type B | 0-5 | 6-10 | 11-15 | 16-31 | | |
| BSRAI Rd,Ra,Imm | 011001 | Rd | Ra | 00000010000 & Imm5 | | Rd := s(Ra >> Imm5) |
| BSLLI Rd,Ra,Imm | 011001 | Rd | Ra | 00000100000 & Imm5 | | Rd := (Ra << Imm5) & 0 |
| IDIV Rd,Ra,Rb | 010010 | Rd | Ra | Rb | 00000000000 | Rd := Rb/Ra |
| IDIVU Rd,Ra,Rb | 010010 | Rd | Ra | Rb | 00000000010 | Rd := Rb/Ra, unsigned |
| TNEAGETD Rd,Rb | 010011 | Rd | 00000 | Rb | 0N0TAE 00000 | Rd := FSL Rb[28:31] (data read) MSR[FSL] := 1 if (FSL_S_Control = 1) MSR[C] := not FSL_S_Exists if N = 1 |
| TNAPUTD Ra,Rb | 010011 | 00000 | Ra | Rb | 0N0TA0 00000 | FSL Rb[28:31] := Ra (data write) MSR[C] := FSL_M_Full if N = 1 |
| TNECAGETD Rd,Rb | 010011 | Rd | 00000 | Rb | 0N1TAE 00000 | Rd := FSL Rb[28:31] (control read) MSR[FSL] := 1 if (FSL_S_Control = 0) MSR[C] := not FSL_S_Exists if N = 1 |
| TNCAPUTD Ra,Rb | 010011 | 00000 | Ra | Rb | 0N1TA0 00000 | FSL Rb[28:31] := Ra (control write) MSR[C] := FSL_M_Full if N = 1 |
| FADD Rd,Ra,Rb | 010110 | Rd | Ra | Rb | 00000000000 | Rd := Rb+Ra, float ¹ |
| FRSUB Rd,Ra,Rb | 010110 | Rd | Ra | Rb | 00010000000 | Rd := Rb-Ra, float ¹ |
| FMUL Rd,Ra,Rb | 010110 | Rd | Ra | Rb | 00100000000 | Rd := Rb*Ra, float ¹ |
| FDIV Rd,Ra,Rb | 010110 | Rd | Ra | Rb | 00110000000 | Rd := Rb/Ra, float ¹ |
| FCMP.UN Rd,Ra,Rb | 010110 | Rd | Ra | Rb | 01000000000 | Rd := 1 if (Rb = NaN or Ra = NaN, float ¹) else Rd := 0 |
| FCMP.LT Rd,Ra,Rb | 010110 | Rd | Ra | Rb | 01000010000 | Rd := 1 if (Rb < Ra, float ¹) else Rd := 0 |
| FCMP.EQ Rd,Ra,Rb | 010110 | Rd | Ra | Rb | 01000100000 | Rd := 1 if (Rb = Ra, float ¹) else Rd := 0 |
| FCMP.LE Rd,Ra,Rb | 010110 | Rd | Ra | Rb | 01000110000 | Rd := 1 if (Rb <= Ra, float ¹) else Rd := 0 |
| FCMP.GT Rd,Ra,Rb | 010110 | Rd | Ra | Rb | 01001000000 | Rd := 1 if (Rb > Ra, float ¹) else Rd := 0 |
| FCMP.NE Rd,Ra,Rb | 010110 | Rd | Ra | Rb | 01001010000 | Rd := 1 if (Rb != Ra, float ¹) else Rd := 0 |
| FCMP.GE Rd,Ra,Rb | 010110 | Rd | Ra | Rb | 01001100000 | Rd := 1 if (Rb >= Ra, float ¹) else Rd := 0 |
| FLT Rd,Ra | 010110 | Rd | Ra | 0 | 01010000000 | Rd := float(Ra) ¹ |
| FINT Rd,Ra | 010110 | Rd | Ra | 0 | 01100000000 | Rd := int(Ra) ¹ |

Table 1-6: MicroBlaze Instruction Set Summary (Continued)

| Type A | 0-5 | 6-10 | 11-15 | 16-20 | 21-31 | Semantics |
|------------------|--------|-------|-------|---------------------|--------------|---|
| Type B | 0-5 | 6-10 | 11-15 | 16-31 | | |
| FSQRT Rd,Ra | 010110 | Rd | Ra | 0 | 011100000000 | Rd := sqrt (Ra) ¹ |
| TNEAGET Rd,FSLx | 011011 | Rd | 00000 | 0N0TAE000000 & FSLx | | Rd := FSLx (data read, blocking if N = 0) MSR[FSL] := 1 if (FSLx_S_Control = 1) MSR[C] := not FSLx_S_Exists if N = 1 |
| TNAPUT Ra,FSLx | 011011 | 00000 | Ra | 1N0TA0000000 & FSLx | | FSLx := Ra (data write, blocking if N = 0) MSR[C] := FSLx_M_Full if N = 1 |
| TNECAGET Rd,FSLx | 011011 | Rd | 00000 | 0N1TAE000000 & FSLx | | Rd := FSLx (control read, blocking if N = 0) MSR[FSL] := 1 if (FSLx_S_Control = 0) MSR[C] := not FSLx_S_Exists if N = 1 |
| TNCAPUT Ra,FSLx | 011011 | 00000 | Ra | 1N1TA0000000 & FSLx | | FSLx := Ra (control write, blocking if N = 0) MSR[C] := FSLx_M_Full if N = 1 |
| OR Rd,Ra,Rb | 100000 | Rd | Ra | Rb | 000000000000 | Rd := Ra or Rb |
| AND Rd,Ra,Rb | 100001 | Rd | Ra | Rb | 000000000000 | Rd := Ra and Rb |
| XOR Rd,Ra,Rb | 100010 | Rd | Ra | Rb | 000000000000 | Rd := Ra xor Rb |
| ANDN Rd,Ra,Rb | 100011 | Rd | Ra | Rb | 000000000000 | Rd := Ra and \overline{Rb} |
| PCMPBF Rd,Ra,Rb | 100000 | Rd | Ra | Rb | 100000000000 | Rd := 1 if (Rb[0:7] = Ra[0:7]) else Rd := 2 if (Rb[8:15] = Ra[8:15]) else Rd := 3 if (Rb[16:23] = Ra[16:23]) else Rd := 4 if (Rb[24:31] = Ra[24:31]) else Rd := 0 |
| PCMPEQ Rd,Ra,Rb | 100010 | Rd | Ra | Rb | 100000000000 | Rd := 1 if (Rd = Ra) else Rd := 0 |
| PCMPNE Rd,Ra,Rb | 100011 | Rd | Ra | Rb | 100000000000 | Rd := 1 if (Rd != Ra) else Rd := 0 |
| SRA Rd,Ra | 100100 | Rd | Ra | 0000000000000001 | | Rd := s(Ra >> 1) C := Ra[31] |
| SRC Rd,Ra | 100100 | Rd | Ra | 0000000000100001 | | Rd := C & (Ra >> 1) C := Ra[31] |
| SRL Rd,Ra | 100100 | Rd | Ra | 0000000001000001 | | Rd := 0 & (Ra >> 1) C := Ra[31] |
| SEXT8 Rd,Ra | 100100 | Rd | Ra | 0000000001100000 | | Rd := s(Ra[24:31]) |
| SEXT16 Rd,Ra | 100100 | Rd | Ra | 0000000001100001 | | Rd := s(Ra[16:31]) |
| WIC Ra,Rb | 100100 | 00000 | Ra | Rb | 00001101000 | ICache_Line[Ra >> 4].Tag := 0 if (C_ICACHE_LINE_LEN = 4) ICache_Line[Ra >> 5].Tag := 0 if (C_ICACHE_LINE_LEN = 8) |

Table 1-6: MicroBlaze Instruction Set Summary (Continued)

| Type A | 0-5 | 6-10 | 11-15 | 16-20 | 21-31 | Semantics |
|-----------------|--------|-------|-------|------------|--------------|---|
| Type B | 0-5 | 6-10 | 11-15 | 16-31 | | |
| WDC Ra,Rb | 100100 | 00000 | Ra | Rb | 00001100100 | Cache line is cleared, discarding stored data. DCache_Line[Ra >> 4].Tag := 0 if (C_DCACHE_LINE_LEN = 4) DCache_Line[Ra >> 5].Tag := 0 if (C_DCACHE_LINE_LEN = 8) |
| WDC.FLUSH Ra,Rb | 100100 | 00000 | Ra | Rb | 00001110100 | Cache line is flushed, writing stored data to memory, and then cleared. Used when C_DCACHE_USE_WRITEBACK = 1. |
| WDC.CLEAR Ra,Rb | 100100 | 00000 | Ra | Rb | 00001110110 | Cache line with matching address is cleared, discarding stored data. Used when C_DCACHE_USE_WRITEBACK = 1. |
| MTS Sd,Ra | 100101 | 00000 | Ra | 11 & Sd | | SPR[Sd] := Ra, where: <ul style="list-style-type: none">• SPR[0x0001] is MSR• SPR[0x0007] is FSR• SPR[0x1000] is PID• SPR[0x1001] is ZPR• SPR[0x1002] is TLBX• SPR[0x1003] is TLBLO• SPR[0x1004] is TLBHI• SPR[0x1005] is TLBSX |
| MFS Rd,Sa | 100101 | Rd | 00000 | 10 & Sa | | Rd := SPR[Sa], where: <ul style="list-style-type: none">• SPR[0x0000] is PC• SPR[0x0001] is MSR• SPR[0x0003] is EAR• SPR[0x0005] is ESR• SPR[0x0007] is FSR• SPR[0x000B] is BTR• SPR[0x000D] is EDR• SPR[0x1000] is PID• SPR[0x1001] is ZPR• SPR[0x1002] is TLBX• SPR[0x1003] is TLBLO• SPR[0x1004] is TLBHI• SPR[0x2000 to 0x200B] is PVR[0 to 11] |
| MSRCLR Rd,Imm | 100101 | Rd | 00001 | 00 & Imm14 | | Rd := MSR MSR := MSR and <u>Imm14</u> |
| MSRSET Rd,Imm | 100101 | Rd | 00000 | 00 & Imm14 | | Rd := MSR MSR := MSR or Imm14 |
| BR Rb | 100110 | 00000 | 00000 | Rb | 000000000000 | PC := PC + Rb |
| BRD Rb | 100110 | 00000 | 10000 | Rb | 000000000000 | PC := PC + Rb |

Table 1-6: MicroBlaze Instruction Set Summary (Continued)

| Type A | 0-5 | 6-10 | 11-15 | 16-20 | 21-31 | Semantics |
|-----------------|--------|-------|-------|-------|--------------|--|
| Type B | 0-5 | 6-10 | 11-15 | 16-31 | | |
| BRLD Rd,Rb | 100110 | Rd | 10100 | Rb | 000000000000 | PC := PC + Rb Rd := PC |
| BRA Rb | 100110 | 00000 | 01000 | Rb | 000000000000 | PC := Rb |
| BRAD Rb | 100110 | 00000 | 11000 | Rb | 000000000000 | PC := Rb |
| BRALD Rd,Rb | 100110 | Rd | 11100 | Rb | 000000000000 | PC := Rb Rd := PC |
| BRK Rd,Rb | 100110 | Rd | 01100 | Rb | 000000000000 | PC := Rb Rd := PC MSR[BIP] := 1 |
| BEQ Ra,Rb | 100111 | 00000 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra = 0 |
| BNE Ra,Rb | 100111 | 00001 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra != 0 |
| BLT Ra,Rb | 100111 | 00010 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra < 0 |
| BLE Ra,Rb | 100111 | 00011 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra <= 0 |
| BGT Ra,Rb | 100111 | 00100 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra > 0 |
| BGE Ra,Rb | 100111 | 00101 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra >= 0 |
| BEQD Ra,Rb | 100111 | 10000 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra = 0 |
| BNED Ra,Rb | 100111 | 10001 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra != 0 |
| BLTD Ra,Rb | 100111 | 10010 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra < 0 |
| BLED Ra,Rb | 100111 | 10011 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra <= 0 |
| BGTD Ra,Rb | 100111 | 10100 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra > 0 |
| BGED Ra,Rb | 100111 | 10101 | Ra | Rb | 000000000000 | PC := PC + Rb if Ra >= 0 |
| ORI Rd,Ra,Imm | 101000 | Rd | Ra | Imm | | Rd := Ra or s(Imm) |
| ANDI Rd,Ra,Imm | 101001 | Rd | Ra | Imm | | Rd := Ra and s(Imm) |
| XORI Rd,Ra,Imm | 101010 | Rd | Ra | Imm | | Rd := Ra xor s(Imm) |
| ANDNI Rd,Ra,Imm | 101011 | Rd | Ra | Imm | | Rd := Ra and $\overline{s(Imm)}$ |
| IMM Imm | 101100 | 00000 | 00000 | Imm | | Imm[0:15] := Imm |
| RTSD Ra,Imm | 101101 | 10000 | Ra | Imm | | PC := Ra + s(Imm) |
| RTID Ra,Imm | 101101 | 10001 | Ra | Imm | | PC := Ra + s(Imm) MSR[IE] := 1 |
| RTBD Ra,Imm | 101101 | 10010 | Ra | Imm | | PC := Ra + s(Imm) MSR[BIP] := 0 |
| RTED Ra,Imm | 101101 | 10100 | Ra | Imm | | PC := Ra + s(Imm) MSR[EE] := 1, MSR[EIP] := 0 ESR := 0 |

Table 1-6: MicroBlaze Instruction Set Summary (Continued)

| Type A | 0-5 | 6-10 | 11-15 | 16-20 | 21-31 | Semantics |
|---------------|--------|-------|-------|-------|--------------|--|
| Type B | 0-5 | 6-10 | 11-15 | 16-31 | | |
| BRI Imm | 101110 | 00000 | 00000 | | Imm | PC := PC + s(Imm) |
| BRID Imm | 101110 | 00000 | 10000 | | Imm | PC := PC + s(Imm) |
| BRLID Rd,Imm | 101110 | Rd | 10100 | | Imm | PC := PC + s(Imm) Rd := PC |
| BRAI Imm | 101110 | 00000 | 01000 | | Imm | PC := s(Imm) |
| BRAID Imm | 101110 | 00000 | 11000 | | Imm | PC := s(Imm) |
| BRALID Rd,Imm | 101110 | Rd | 11100 | | Imm | PC := s(Imm) Rd := PC |
| BRKI Rd,Imm | 101110 | Rd | 01100 | | Imm | PC := s(Imm) Rd := PC MSR[BIP] := 1 |
| BEQI Ra,Imm | 101111 | 00000 | Ra | | Imm | PC := PC + s(Imm) if Ra = 0 |
| BNEI Ra,Imm | 101111 | 00001 | Ra | | Imm | PC := PC + s(Imm) if Ra != 0 |
| BLTI Ra,Imm | 101111 | 00010 | Ra | | Imm | PC := PC + s(Imm) if Ra < 0 |
| BLEI Ra,Imm | 101111 | 00011 | Ra | | Imm | PC := PC + s(Imm) if Ra <= 0 |
| BGTI Ra,Imm | 101111 | 00100 | Ra | | Imm | PC := PC + s(Imm) if Ra > 0 |
| BGEI Ra,Imm | 101111 | 00101 | Ra | | Imm | PC := PC + s(Imm) if Ra >= 0 |
| BEQID Ra,Imm | 101111 | 10000 | Ra | | Imm | PC := PC + s(Imm) if Ra = 0 |
| BNEID Ra,Imm | 101111 | 10001 | Ra | | Imm | PC := PC + s(Imm) if Ra != 0 |
| BLTID Ra,Imm | 101111 | 10010 | Ra | | Imm | PC := PC + s(Imm) if Ra < 0 |
| BLEID Ra,Imm | 101111 | 10011 | Ra | | Imm | PC := PC + s(Imm) if Ra <= 0 |
| BGTID Ra,Imm | 101111 | 10100 | Ra | | Imm | PC := PC + s(Imm) if Ra > 0 |
| BGEID Ra,Imm | 101111 | 10101 | Ra | | Imm | PC := PC + s(Imm) if Ra >= 0 |
| LBU Rd,Ra,Rb | 110000 | Rd | Ra | Rb | 000000000000 | Addr := Ra + Rb Rd[0:23] := 0 Rd[24:31] := *Addr[0:7] |
| LHU Rd,Ra,Rb | 110001 | Rd | Ra | Rb | 000000000000 | Addr := Ra + Rb Rd[0:15] := 0 Rd[16:31] := *Addr[0:15] |
| LW Rd,Ra,Rb | 110010 | Rd | Ra | Rb | 000000000000 | Addr := Ra + Rb Rd := *Addr |
| LWX Rd,Ra,Rb | 110010 | Rd | Ra | Rb | 100000000000 | Addr := Ra + Rb Rd := *Addr Reservation := 1 |

Table 1-6: MicroBlaze Instruction Set Summary (Continued)

| Type A | 0-5 | 6-10 | 11-15 | 16-20 | 21-31 | Semantics |
|----------------|--------|------|-------|-------|--------------|---|
| Type B | 0-5 | 6-10 | 11-15 | 16-31 | | |
| SB Rd,Ra,Rb | 110100 | Rd | Ra | Rb | 000000000000 | Addr := Ra + Rb *Addr[0:8] := Rd[24:31] |
| SH Rd,Ra,Rb | 110101 | Rd | Ra | Rb | 000000000000 | Addr := Ra + Rb *Addr[0:16] := Rd[16:31] |
| SW Rd,Ra,Rb | 110110 | Rd | Ra | Rb | 000000000000 | Addr := Ra + Rb *Addr := Rd |
| SWX Rd,Ra,Rb | 110110 | Rd | Ra | Rb | 100000000000 | Addr := Ra + Rb *Addr := Rd if Reservation = 1 Reservation := 0 |
| LBUI Rd,Ra,Imm | 111000 | Rd | Ra | Imm | | Addr := Ra + s(Imm) Rd[0:23] := 0 Rd[24:31] := *Addr[0:7] |
| LHUI Rd,Ra,Imm | 111001 | Rd | Ra | Imm | | Addr := Ra + s(Imm) Rd[0:15] := 0 Rd[16:31] := *Addr[0:15] |
| LWI Rd,Ra,Imm | 111010 | Rd | Ra | Imm | | Addr := Ra + s(Imm) Rd := *Addr |
| SBI Rd,Ra,Imm | 111100 | Rd | Ra | Imm | | Addr := Ra + s(Imm) *Addr[0:7] := Rd[24:31] |
| SHI Rd,Ra,Imm | 111101 | Rd | Ra | Imm | | Addr := Ra + s(Imm) *Addr[0:15] := Rd[16:31] |
| SWI Rd,Ra,Imm | 111110 | Rd | Ra | Imm | | Addr := Ra + s(Imm) *Addr := Rd |

1. Due to the many different corner cases involved in floating point arithmetic, only the normal behavior is described. A full description of the behavior can be found in [Chapter 4, “MicroBlaze Instruction Set Architecture.”](#)

Semaphore Synchronization

The LWX and SWX instructions are used to implement common semaphore operations, including test and set, compare and swap, exchange memory, and fetch and add. They are also used to implement spinlocks.

These instructions are typically used by system programs and are called by application programs as needed. Generally, a program uses LWX to load a semaphore from memory, causing the reservation to be set (the processor maintains the reservation internally). The program can compute a result based on the semaphore value and conditionally store the result back to the same memory location using the SWX instruction. The conditional store is performed based on the existence of the reservation established by the preceding LWX instruction. If the reservation exists when the store is executed, the store is performed and MSR[C] is cleared to 0. If the reservation does not exist when the store is executed, the target memory location is not modified and MSR[C] is set to 1.

If the store is successful, the sequence of instructions from the semaphore load to the semaphore store appear to be executed atomically—no other device modified the semaphore location between

the read and the update. Other devices can read from the semaphore location during the operation. For a semaphore operation to work properly, the LWX instruction must be paired with an SWX instruction, and both must specify identical addresses. The reservation granularity in MicroBlaze is a word. For both instructions, the address must be word aligned. No unaligned exceptions are generated for these instructions.

The conditional store is always performed when a reservation exists, even if the store address does not match the load address that set the reservation.

Only one reservation can be maintained at a time. The address associated with the reservation can be changed by executing a subsequent LWX instruction. The conditional store is performed based upon the reservation established by the last LWX instruction executed. Executing an SWX instruction always clears a reservation held by the processor, whether the address matches that established by the LWX or not.

Reset, interrupts, exceptions, and breaks (including the BRK and BRKI instructions) all clear the reservation.

The following provides general guidelines for using the LWX and SWX instructions:

- The LWX and SWX instructions should be paired and use the same address.
- An unpaired SWX instruction to an arbitrary address can be used to clear any reservation held by the processor.
- A conditional sequence begins with an LWX instruction. It can be followed by memory accesses and/or computations on the loaded value. The sequence ends with an SWX instruction. In most cases, failure of the SWX instruction should cause a branch back to the LWX for a repeated attempt.
- An LWX instruction can be left unpaired when executing certain synchronization primitives if the value loaded by the LWX is not zero. An implementation of Test and Set exemplifies this:

```
loop: lwx r5,r3,r0 ; load and reserve
      bnei r5,next ; branch if not equal to zero
      addik r5,r5,1 ; increment value
      swx r5,r3,r0 ; try to store non-zero value
      addic r5,r0,0 ; check reservation
      bnei r5,loop ; loop if reservation lost
next:
```

- Performance can be improved by minimizing looping on an LWX instruction that fails to return a desired value. Performance can also be improved by using an ordinary load instruction to do the initial value check. An implementation of a spinlock exemplifies this:

```
loop: lw r5,r3,r0 ; load the word
      bnei r5,loop ; loop back if word not equal to 0
      lwx r5,r3,r0 ; try reserving again
      bnei r5,loop ; likely that no branch is needed
      addik r5,r5,1 ; increment value
      swx r5,r3,r0 ; try to store non-zero value
      addic r5,r0,0 ; check reservation
      bnei r5,loop ; loop if reservation lost
```

- Minimizing the looping on an LWX/SWX instruction pair increases the likelihood that forward progress is made. The old value should be tested before attempting the store. If the order is reversed (store before load), more SWX instructions are executed and reservations are more likely to be lost between the LWX and SWX instructions.

Registers

MicroBlaze has an orthogonal instruction set architecture. It has thirty-two 32-bit general purpose registers and up to eighteen 32-bit special purpose registers, depending on configured options.

General Purpose Registers

The thirty-two 32-bit General Purpose Registers are numbered R0 through R31. The register file is reset on bit stream download (reset value is 0x00000000). [Figure 1-2](#) is a representation of a General Purpose Register and [Table 1-7](#) provides a description of each register and the register reset value (if existing).

Note: The register file is **not** reset by the external reset inputs: `Reset` and `Debug_Rst`.

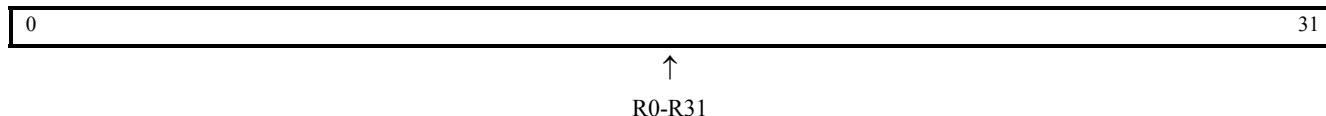


Figure 1-2: R0-R31

Table 1-7: General Purpose Registers (R0-R31)

| Bits | Name | Description | Reset Value |
|------|-----------------|---|-------------|
| 0:31 | R0 | Always has a value of zero. Anything written to R0 is discarded | 0x00000000 |
| 0:31 | R1 through R13 | 32-bit general purpose registers | - |
| 0:31 | R14 | 32-bit register used to store return addresses for interrupts. | - |
| 0:31 | R15 | 32-bit general purpose register. Recommended for storing return addresses for user vectors. | - |
| 0:31 | R16 | 32-bit register used to store return addresses for breaks. | - |
| 0:31 | R17 | If MicroBlaze is configured to support hardware exceptions, this register is loaded with the address of the instruction following the instruction causing the HW exception, except for exceptions in delay slots that use BTR instead (see “ Branch Target Register (BTR) ”); if not, it is a general purpose register. | - |
| 0:31 | R18 through R31 | R18 through R31 are 32-bit general purpose registers. | - |

Refer to [Table 3-2](#) for software conventions on general purpose register usage.

Special Purpose Registers

Program Counter (PC)

The Program Counter (PC) is the 32-bit address of the execution instruction. It can be read with an MFS instruction, but it cannot be written with an MTS instruction. When used with the MFS instruction the PC register is specified by setting Sa = 0x0000. [Figure 1-3](#) illustrates the PC and [Table 1-8](#) provides a description and reset value.

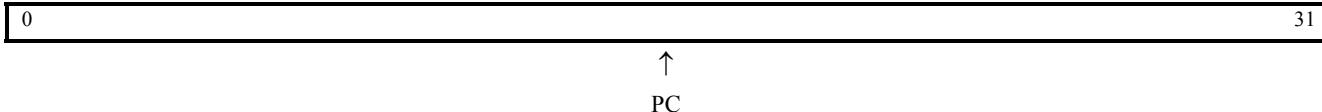


Figure 1-3: PC

Table 1-8: Program Counter (PC)

| Bits | Name | Description | Reset Value |
|------|------|--|-------------|
| 0:31 | PC | Program Counter Address of executing instruction, that is, “mfs r2 0” stores the address of the mfs instruction itself in R2. | 0x00000000 |

Machine Status Register (MSR)

The Machine Status Register contains control and status bits for the processor. It can be read with an MFS instruction. When reading the MSR, bit 29 is replicated in bit 0 as the carry copy. MSR can be written using either an MTS instruction or the dedicated MSRSET and MSRCLR instructions.

When writing to the MSR using MSRSET or MSRCLR, the Carry bit takes effect immediately and the remaining bits take effect one clock cycle later. When writing using MTS, all bits take effect one clock cycle later. Any value written to bit 0 is discarded.

When used with an MTS or MFS instruction, the MSR is specified by setting Sx = 0x0001.

[Figure 1-4](#) illustrates the MSR register and [Table 1-9](#) provides the bit description and reset values.



Figure 1-4: MSR

Table 1-9: Machine Status Register (MSR)

| Bits | Name | Description | Reset Value |
|------|----------|---|--------------------------|
| 0 | CC | Arithmetic Carry Copy Copy of the Arithmetic Carry (bit 29). CC is always the same as bit C. | 0 |
| 1:16 | Reserved | | |
| 17 | VMS | Virtual Protected Mode Save Only available when configured with an MMU (if C_USE_MMU > 1 and C_AREA_OPTIMIZED = 0) Read/Write | 0 |
| 18 | VM | Virtual Protected Mode 0 = MMU address translation and access protection disabled, with C_USE_MMU = 3. Access protection disabled, with C_USE_MMU = 2. 1 = MMU address translation and access protection enabled, with C_USE_MMU = 3. Access protection enabled, with C_USE_MMU = 2. Only available when configured with an MMU (if C_USE_MMU > 1 and C_AREA_OPTIMIZED = 0) Read/Write | 0 |
| 19 | UMS | User Mode Save Only available when configured with an MMU (if C_USE_MMU > 0 and C_AREA_OPTIMIZED = 0) Read/Write | 0 |
| 20 | UM | User Mode 0 = Privileged Mode, all instructions are allowed 1 = User Mode, certain instructions are not allowed Only available when configured with an MMU (if C_USE_MMU > 0 and C_AREA_OPTIMIZED = 0) Read/Write | 0 |
| 21 | PVR | Processor Version Register exists 0 = No Processor Version Register 1 = Processor Version Register exists Read only | Based on parameter C_PVR |
| 22 | EIP | Exception In Progress 0 = No hardware exception in progress 1 = Hardware exception in progress Only available if configured with exception support (C_*_EXCEPTION or C_USE_MMU) Read/Write | 0 |

Table 1-9: Machine Status Register (MSR) (Continued)

| Bits | Name | Description | Reset Value |
|------|------|--|-------------|
| 23 | EE | Exception Enable 0 = Hardware exceptions disabled ¹ 1 = Hardware exceptions enabled Only available if configured with exception support (C_*_EXCEPTION or C_USE_MMU) Read/Write | 0 |
| 24 | DCE | Data Cache Enable 0 = Data Cache disabled 1 = Data Cache enabled Only available if configured to use data cache (C_USE_DCACHE = 1) Read/Write | 0 |
| 25 | DZO | Division by Zero or Division Overflow ² 0 = No division by zero or division overflow has occurred 1 = Division by zero or division overflow has occurred Only available if configured to use hardware divider (C_USE_DIV = 1) Read/Write | 0 |
| 26 | ICE | Instruction Cache Enable 0 = Instruction Cache disabled 1 = Instruction Cache enabled Only available if configured to use instruction cache (C_USE_ICACHE = 1) Read/Write | 0 |
| 27 | FSL | FSL Error 0 = FSL get/getd/put/putd had no error 1 = FSL get/getd/put/putd control type mismatch Only available if configured to use FSL links (C_FSL_LINKS > 0) Read/Write | 0 |
| 28 | BIP | Break in Progress 0 = No Break in Progress 1 = Break in Progress Break Sources can be software break instruction or hardware break from Ext_Brk or Ext_NM_Brk pin. Read/Write | 0 |

Table 1-9: Machine Status Register (MSR) (Continued)

| Bits | Name | Description | Reset Value |
|-------------|-------------|---|--------------------|
| 29 | C | Arithmetic Carry 0 = No Carry (Borrow) 1 = Carry (No Borrow) Read/Write | 0 |
| 30 | IE | Interrupt Enable 0 = Interrupts disabled 1 = Interrupts enabled Read/Write | 0 |
| 31 | - | Reserved | 0 |

1. The MMU exceptions (Data Storage Exception, Instruction Storage Exception, Data TLB Miss Exception, Instruction TLB Miss Exception) cannot be disabled, and are not affected by this bit.
2. This bit is only used for integer divide-by-zero or divide overflow signaling. There is a floating point equivalent in the FSR. The DZO-bit flags divide by zero or divide overflow conditions regardless if the processor is configured with exception handling or not.

Exception Address Register (EAR)

The Exception Address Register stores the full load/store address that caused the exception for the following:

- An unaligned access exception that means the unaligned access address
- A DPLB exception that specifies the failing PLB data access address
- A data storage exception that specifies the (virtual) effective address accessed
- An instruction storage exception that specifies the (virtual) effective address read
- A data TLB miss exception that specifies the (virtual) effective address accessed
- An instruction TLB miss exception that specifies the (virtual) effective address read

The contents of this register is undefined for all other exceptions. When read with the MFS instruction, the EAR is specified by setting Sa = 0x0003. The EAR register is illustrated in [Figure 1-5](#) and [Table 1-10](#) provides bit descriptions and reset values.

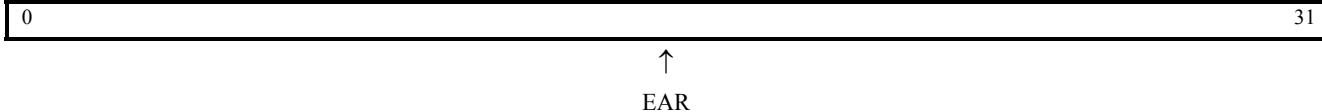


Figure 1-5: EAR

Table 1-10: Exception Address Register (EAR)

| Bits | Name | Description | Reset Value |
|------|------|----------------------------|-------------|
| 0:31 | EAR | Exception Address Register | 0x00000000 |

Exception Status Register (ESR)

The Exception Status Register contains status bits for the processor. When read with the MFS instruction, the ESR is specified by setting Sa = 0x0005. The ESR register is illustrated in Figure 1-6, Table 1-11 provides bit descriptions and reset values, and Table 1-12 provides the Exception Specific Status (ESS).

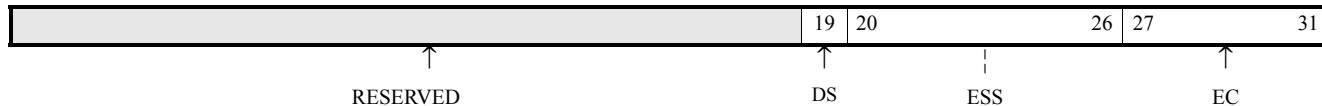


Figure 1-6: ESR

Table 1-11: Exception Status Register (ESR)

| Bits | Name | Description | Reset Value |
|-------|----------|---|----------------|
| 0:18 | Reserved | | |
| 19 | DS | Delay Slot Exception. 0 = not caused by delay slot instruction 1 = caused by delay slot instruction Read-only | 0 |
| 20:26 | ESS | Exception Specific Status For details refer to Table 1-12. Read-only | See Table 1-12 |
| 27:31 | EC | Exception Cause 00000 = Fast Simplex Link exception 00001 = Unaligned data access exception 00010 = Illegal op-code exception 00011 = Instruction bus error exception 00100 = Data bus error exception 00101 = Divide exception 00110 = Floating point unit exception 00111 = Privileged instruction exception 10000 = Data storage exception 10001 = Instruction storage exception 10010 = Data TLB miss exception 10011 = Instruction TLB miss exception Read-only | 0 |

Table 1-12: Exception Specific Status (ESS)

| Exception Cause | Bits | Name | Description | Reset Value |
|------------------------|-------|----------|--|-------------|
| Unaligned Data Access | 20 | W | Word Access Exception 0 = unaligned halfword access 1 = unaligned word access | 0 |
| | 21 | S | Store Access Exception 0 = unaligned load access 1 = unaligned store access | 0 |
| | 22:26 | Rx | Source/Destination Register General purpose register used as source (Store) or destination (Load) in unaligned access | 0 |
| Illegal Instruction | 20:26 | Reserved | | 0 |
| Instruction bus error | 20:26 | Reserved | | 0 |
| Data bus error | 20:26 | Reserved | | 0 |
| Divide | 20 | DEC | Divide - Division exception cause 0 = Divide-By-Zero 1 = Division Overflow | 0 |
| | 21:26 | Reserved | | 0 |
| Floating point unit | 20:26 | Reserved | | 0 |
| Privileged instruction | 20:26 | Reserved | | 0 |
| Fast Simplex Link | 20:22 | Reserved | | 0 |
| | 23:26 | FSL | Fast Simplex Link index that caused the exception | 0 |
| Data storage | 20 | DIZ | Data storage - Zone protection 0 = Did not occur 1 = Occurred | 0 |
| | 21 | S | Data storage - Store instruction 0 = Did not occur 1 = Occurred | 0 |
| | 22:26 | Reserved | | 0 |

Table 1-12: Exception Specific Status (ESS) (Continued)

| Exception Cause | Bits | Name | Description | Reset Value |
|----------------------|-------|----------|--|-------------|
| Instruction storage | 20 | DIZ | Instruction storage - Zone protection 0 = Did not occur 1 = Occurred | 0 |
| | 21:26 | Reserved | | 0 |
| Data TLB miss | 20 | Reserved | | 0 |
| | 21 | S | Data TLB miss - Store instruction 0 = Did not occur 1 = Occurred | 0 |
| | 22:26 | Reserved | | 0 |
| Instruction TLB miss | 20:26 | Reserved | | 0 |

Branch Target Register (BTR)

The Branch Target Register only exists if the MicroBlaze processor is configured to use exceptions. The register stores the branch target address for all delay slot branch instructions executed while MSR[EIP] = 0. If an exception is caused by an instruction in a delay slot (that is, ESR[DS]=1), the exception handler should return execution to the address stored in BTR instead of the normal exception return address stored in R17. When read with the MFS instruction, the BTR is specified by setting Sa = 0x0000B. The BTR register is illustrated in [Figure 1-7](#) and [Table 1-13](#) provides bit descriptions and reset values.

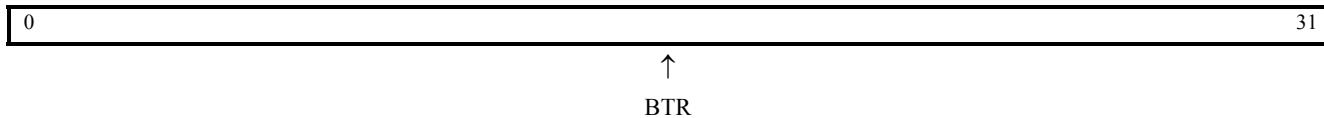


Figure 1-7: BTR

Table 1-13: Branch Target Register (BTR)

| Bits | Name | Description | Reset Value |
|------|------|---|-------------|
| 0:31 | BTR | Branch target address used by handler when returning from an exception caused by an instruction in a delay slot. Read-only | 0x00000000 |

Floating Point Status Register (FSR)

The Floating Point Status Register contains status bits for the floating point unit. It can be read with an MFS, and written with an MTS instruction. When read or written, the register is specified by setting Sa = 0x0007. The bits in this register are sticky – floating point instructions can only set bits in the register, and the only way to clear the register is by using the MTS instruction. [Figure 1-8](#) illustrates the FSR register and [Table 1-14](#) provides bit descriptions and reset values.

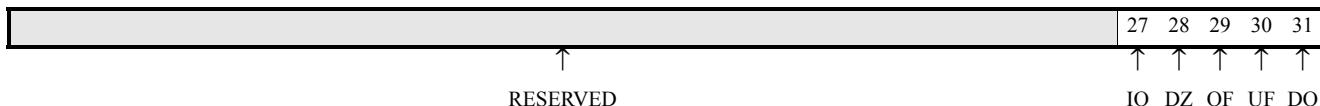


Figure 1-8: FSR

Table 1-14: Floating Point Status Register (FSR)

| Bits | Name | Description | Reset Value |
|------|----------|----------------------------|-------------|
| 0:26 | Reserved | | undefined |
| 27 | IO | Invalid operation | 0 |
| 28 | DZ | Divide-by-zero | 0 |
| 29 | OF | Overflow | 0 |
| 30 | UF | Underflow | 0 |
| 31 | DO | Denormalized operand error | 0 |

Exception Data Register (EDR)

The Exception Data Register stores data read on an FSL link that caused an FSL exception.

The contents of this register is undefined for all other exceptions. When read with the MFS instruction, the EDR is specified by setting Sa = 0x000D. [Figure 1-9](#) illustrates the EDR register and [Table 1-15](#) provides bit descriptions and reset values.

Note: The register is only implemented if C_FSL_LINKS is greater than 0 and C_FSL_EXCEPTION is set to 1.

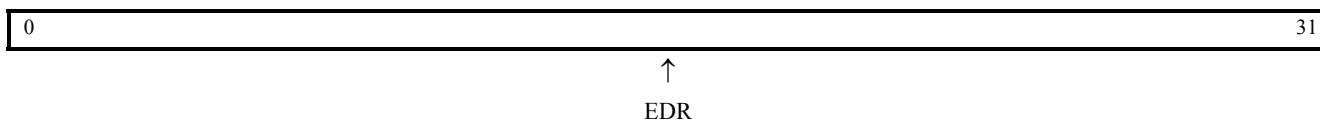


Figure 1-9: EDR

Table 1-15: Exception Data Register (EDR)

| Bits | Name | Description | Reset Value |
|------|------|-------------------------|-------------|
| 0:31 | EDR | Exception Data Register | 0x00000000 |

Process Identifier Register (PID)

The Process Identifier Register is used to uniquely identify a software process during MMU address translation. It is controlled by the C_USE_MMU configuration option on MicroBlaze. The register is only implemented if C_USE_MMU is greater than 1 and C_AREA_OPTIMIZED is set to 0. When accessed with the MFS and MTS instructions, the PID is specified by setting Sa = 0x1000. The register is accessible according to the memory management special registers parameter C_MMU_TLB_ACCESS.

PID is also used when accessing a TLB entry:

- When writing Translation Look-Aside Buffer High (TLBHI) the value of PID is stored in the TID field of the TLB entry
- When reading TLBHI and MSR[UM] is not set, the value in the TID field is stored in PID

[Figure 1-10](#) illustrates the PID register and [Table 1-16](#) provides bit descriptions and reset values.

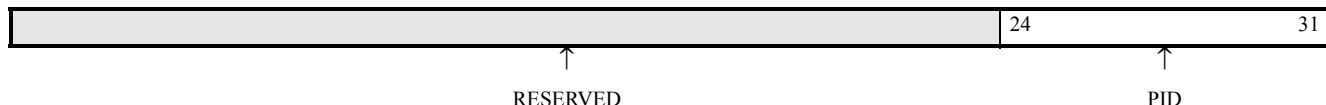


Figure 1-10: PID

Table 1-16: Process Identifier Register (PID)

| Bits | Name | Description | Reset Value |
|-------|----------|--|-------------|
| 0:23 | Reserved | | |
| 24:31 | PID | Used to uniquely identify a software process during MMU address translation. Read/Write | 0x00 |

Zone Protection Register (ZPR)

The Zone Protection Register is used to override MMU memory protection defined in TLB entries. It is controlled by the C_USE_MMU configuration option on MicroBlaze. The register is only implemented if C_USE_MMU is greater than 1, C_AREA_OPTIMIZED is set to 0, and if the number of specified memory protection zones is greater than zero (C_MMU_ZONES > 0). The implemented register bits depend on the number of specified memory protection zones (C_MMU_ZONES). When accessed with the MFS and MTS instructions, the ZPR is specified by setting Sa = 0x1001. The register is accessible according to the memory management special registers parameter C_MMU_TLB_ACCESS. Figure 1-11 illustrates the ZPR register and Table 1-17 provides bit descriptions and reset values.

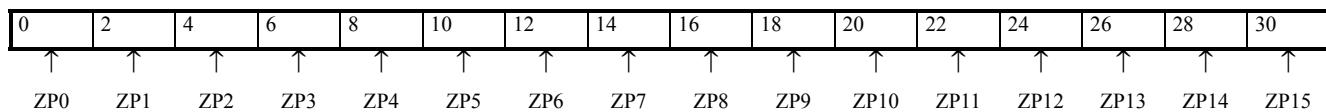


Figure 1-11: ZPR

Table 1-17: Zone Protection Register (ZPR)

| Bits | Name | Description | Reset Value |
|-------|------|--|-------------|
| 0:1 | ZP0 | Zone Protect | 0x00000000 |
| 2:3 | ZP1 | User mode (MSR[UM] = 1): | |
| ... | ... | 00 = Override V in TLB entry. No access to the page is allowed | |
| 30:31 | ZP15 | 01 = No override. Use V, WR and EX from TLB entry 10 = No override. Use V, WR and EX from TLB entry 11 = Override WR and EX in TLB entry. Access the page as writable and executable Privileged mode (MSR[UM] = 0): 00 = No override. Use V, WR and EX from TLB entry 01 = No override. Use V, WR and EX from TLB entry 10 = Override WR and EX in TLB entry. Access the page as writable and executable 11 = Override WR and EX in TLB entry. Access the page as writable and executable Read/Write | |

Translation Look-Aside Buffer Low Register (TLBLO)

The Translation Look-Aside Buffer Low Register is used to access MMU Unified Translation Look-Aside Buffer (UTLB) entries. It is controlled by the C_USE_MMU configuration option on MicroBlaze. The register is only implemented if C_USE_MMU is greater than 1, and C_AREA_OPTIMIZED is set to 0. When accessed with the MFS and MTS instructions, the TLBLO is specified by setting Sa = 0x1003. When reading or writing TLBLO, the UTLB entry indexed by the TLBX register is accessed. The register is readable according to the memory management special registers parameter C_MMU_TLB_ACCESS.

The UTLB is reset on bit stream download (reset value is 0x00000000 for all TLBLO entries).

Note: The UTLB is **not** reset by the external reset inputs: Reset and Debug_Rst.

Figure 1-12 illustrates the TLBLO register and Table 1-18 provides bit descriptions and reset values.

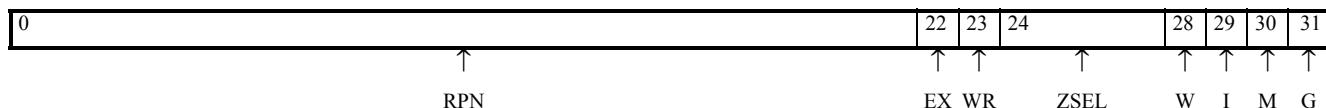


Figure 1-12: TLBLO

Table 1-18: Translation Look-Aside Buffer Low Register (TLBLO)

| Bits | Name | Description | Reset Value |
|------|------|---|-------------|
| 0:21 | RPN | <p>Real Page Number or Physical Page Number</p> <p>When a TLB hit occurs, this field is read from the TLB entry and is used to form the physical address. Depending on the value of the SIZE field, some of the RPN bits are not used in the physical address. Software must clear unused bits in this field to zero.</p> <p>Only defined when C_USE_MMU=3.</p> <p>Read/Write</p> | 0x0000000 |
| 22 | EX | <p>Executable</p> <p>When bit is set to 1, the page contains executable code, and instructions can be fetched from the page. When bit is cleared to 0, instructions cannot be fetched from the page. Attempts to fetch instructions from a page with a clear EX bit cause an instruction-storage exception.</p> <p>Read/Write</p> | 0 |

Table 1-18: Translation Look-Aside Buffer Low Register (TLBLO) (Continued)

| Bits | Name | Description | Reset Value |
|-------|------|---|-------------|
| 23 | WR | <p>Writable</p> <p>When bit is set to 1, the page is writable and store instructions can be used to store data at addresses within the page.</p> <p>When bit is cleared to 0, the page is read-only (not writable). Attempts to store data into a page with a clear WR bit cause a data storage exception.</p> <p>Read/Write</p> | 0 |
| 24:27 | ZSEL | <p>Zone Select</p> <p>This field selects one of 16 zone fields (Z0-Z15) from the zone-protection register (ZPR).</p> <p>For example, if ZSEL 0x5, zone field Z5 is selected. The selected ZPR field is used to modify the access protection specified by the TLB entry EX and WR fields. It is also used to prevent access to a page by overriding the TLB V (valid) field.</p> <p>Read/Write</p> | 0x0 |
| 28 | W | <p>Write Through</p> <p>When the parameter C_DCACHE_USE_WRITEBACK is set to 1, this bit controls caching policy. A write-through policy is selected when set to 1, and a write-back policy is selected otherwise.</p> <p>This bit is fixed to 1, and write-through is always used, when C_DCACHE_USE_WRITEBACK is cleared to 0.</p> <p>Read/Write</p> | 0/1 |
| 29 | I | <p>Inhibit Caching</p> <p>When bit is set to 1, accesses to the page are not cached (caching is inhibited).</p> <p>When cleared to 0, accesses to the page are cacheable.</p> <p>Read/Write</p> | 0 |
| 30 | M | <p>Memory Coherent</p> <p>This bit is fixed to 0, because memory coherence is not implemented on MicroBlaze.</p> <p>Read Only</p> | 0 |
| 31 | G | <p>Guarded</p> <p>When bit is set to 1, speculative page accesses are not allowed (memory is guarded).</p> <p>When cleared to 0, speculative page accesses are allowed.</p> <p>The G attribute can be used to protect memory-mapped I/O devices from inappropriate instruction accesses.</p> <p>Read/Write</p> | 0 |

Translation Look-Aside Buffer High Register (TLBHI)

The Translation Look-Aside Buffer High Register is used to access MMU Unified Translation Look-Aside Buffer (UTLB) entries. It is controlled by the C_USE_MMU configuration option on MicroBlaze. The register is only implemented if C_USE_MMU is greater than 1, and C_AREA_OPTIMIZED is set to 0. When accessed with the MFS and MTS instructions, the TLBHI is specified by setting Sa = 0x1004. When reading or writing TLBHI, the UTLB entry indexed by the TLBX register is accessed. The register is readable according to the memory management special registers parameter C_MMU_TLB_ACCESS.

PID is also used when accessing a TLB entry:

- When writing TLBHI the value of PID is stored in the TID field of the TLB entry
- When reading TLBHI and MSR[UM] is not set, the value in the TID field is stored in PID

The UTLB is reset on bit stream download (reset value is 0x00000000 for all TLBHI entries).

Note: The UTLB is **not** reset by the external reset inputs: Reset and Debug_Rst.

Figure 1-13 illustrates the TLBHI register and Table 1-19 provides bit descriptions and reset values.

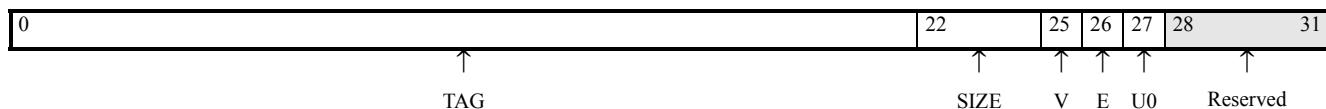


Figure 1-13: TLBHI

Table 1-19: Translation Look-Aside Buffer High Register (TLBHI)

| Bits | Name | Description | Reset Value |
|-------|------|---|-------------|
| 0:21 | TAG | TLB-entry tag Is compared with the page number portion of the virtual memory address under the control of the SIZE field. Read/Write | 0x000000 |
| 22:24 | SIZE | Size Specifies the page size. The SIZE field controls the bit range used in comparing the TAG field with the page number portion of the virtual memory address. The page sizes defined by this field are listed in Table 1-34. Read/Write | 000 |
| 25 | V | Valid When bit is set to 1, the TLB entry is valid and contains a page-translation entry. When cleared to 0, the TLB entry is invalid. Read/Write | 0 |

**Table 1-19: Translation Look-Aside Buffer High Register (TLBHI) (Continued)
(Continued)**

| Bits | Name | Description | Reset Value |
|-------|----------|--|-------------|
| 26 | E | Endian This bit is fixed to 0, since accesses to pages on MicroBlaze are always big endian. Read Only | 0 |
| 27 | U0 | User Defined This bit is fixed to 0, since there are no user defined storage attributes on MicroBlaze. Read Only | 0 |
| 28:31 | Reserved | | |

Translation Look-Aside Buffer Index Register (TLBX)

The Translation Look-Aside Buffer Index Register is used as an index to the Unified Translation Look-Aside Buffer (UTLB) when accessing the TLBLO and TLBHI registers. It is controlled by the C_USE_MMU configuration option on MicroBlaze. The register is only implemented if C_USE_MMU is greater than 1, and C_AREA_OPTIMIZED is set to 0. When accessed with the MFS and MTS instructions, the TLBX is specified by setting Sa = 0x1002. [Figure 1-14](#) illustrates the TLBX register and [Table 1-20](#) provides bit descriptions and reset values.

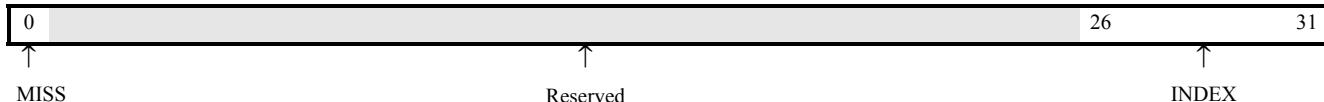


Figure 1-14: TLBX

Table 1-20: Translation Look-Aside Buffer Index Register (TLBX)

| Bits | Name | Description | Reset Value |
|-------|----------|---|-------------|
| 0 | MISS | <p>TLB Miss</p> <p>This bit is cleared to 0 when the TLBSX register is written with a virtual address, and the virtual address is found in a TLB entry.</p> <p>The bit is set to 1 if the virtual address is not found. It is also cleared when the TLBX register itself is written.</p> <p>Read Only</p> <p>Can be read if the memory management special registers parameter C_MMU_TLB_ACCESS > 0.</p> | 0 |
| 1:25 | Reserved | | |
| 26:31 | INDEX | <p>TLB Index</p> <p>This field is used to index the Translation Look-Aside Buffer entry accessed by the TLBLO and TLBHI registers. The field is updated with a TLB index when the TLBSX register is written with a virtual address, and the virtual address is found in the corresponding TLB entry.</p> <p>Read/Write</p> <p>Can be read and written if the memory management special registers parameter C_MMU_TLB_ACCESS > 0.</p> | 000000 |

Translation Look-Aside Buffer Search Index Register (TLBSX)

The Translation Look-Aside Buffer Search Index Register is used to search for a virtual page number in the Unified Translation Look-Aside Buffer (UTLB). It is controlled by the C_USE_MMU configuration option on MicroBlaze. The register is only implemented if C_USE_MMU is greater than 1, and C_AREA_OPTIMIZED is set to 0. When written with the MTS instruction, the TLBSX is specified by setting Sa = 0x1005. [Figure 1-15](#) illustrates the TLBSX register and [Table 1-21](#) provides bit descriptions and reset values.

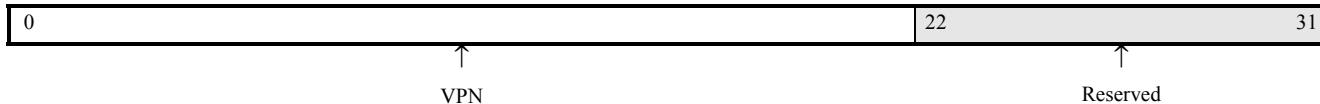


Figure 1-15: TLBSX

Table 1-21: Translation Look-Aside Buffer Index Search Register (TLBSX)

| Bits | Name | Description | Reset Value |
|-------|----------|---|-------------|
| 0:21 | VPN | <p>Virtual Page Number</p> <p>This field represents the page number portion of the virtual memory address. It is compared with the page number portion of the virtual memory address under the control of the SIZE field, in each of the Translation Look-Aside Buffer entries that have the V bit set to 1.</p> <p>If the virtual page number is found, the TLBX register is written with the index of the TLB entry and the MISS bit in TLBX is cleared to 0. If the virtual page number is not found in any of the TLB entries, the MISS bit in the TLBX register is set to 1.</p> <p>Write Only</p> | |
| 22:31 | Reserved | | |

Processor Version Register (PVR)

The Processor Version Register is controlled by the C_PVR configuration option on MicroBlaze.

- When C_PVR is set to 0 the processor does not implement any PVR and MSR[PVR]=0.
- When C_PVR is set to 1, MicroBlaze implements only the first register: PVR0, and if set to 2, all 12 PVR registers (PVR0 to PVR11) are implemented.

When read with the MFS instruction the PVR is specified by setting Sa = 0x200x, with x being the register number between 0x0 and 0xB.

[Table 1-22](#) through [Table 1-33](#) provide bit descriptions and values.

Table 1-22: Processor Version Register 0 (PVR0)

| Bits | Name | Description | Value |
|-------|----------|--|---|
| 0 | CFG | PVR implementation: 0=basic, 1=full | Based on C_PVR |
| 1 | BS | Use barrel shifter | C_USE_BARREL |
| 2 | DIV | Use divider | C_USE_DIV |
| 3 | MUL | Use hardware multiplier | C_USE_HW_MUL > 0 |
| 4 | FPU | Use FPU | C_USE_FPU > 0 |
| 5 | EXC | Use any type of exceptions | Based on C_*_EXCEPTION Also set if C_USE_MMU > 0 |
| 6 | ICU | Use instruction cache | C_USE_ICACHE |
| 7 | DCU | Use data cache | C_USE_DCACHE |
| 8 | MMU | Use MMU | C_USE_MMU > 0 |
| 9 | BTC | Use branch target cache | C_USE_BRANCH_TARGET_CACHE |
| 10:15 | Reserved | | 0 |
| 16:23 | MBV | MicroBlaze release version code 0x1 = v5.00.a 0x2 = v5.00.b 0x3 = v5.00.c 0x4 = v6.00.a 0x6 = v6.00.b 0x5 = v7.00.a 0x7 = v7.00.b 0x8 = v7.10.a 0x9 = v7.10.b 0xA = v7.10.c 0xB = v7.10.d 0xC = v7.20.a 0xD = v7.20.b 0xE = v7.20.c 0xF = v7.20.d 0x10 = v7.30.a | Release Specific |
| 24:31 | USR1 | User configured value 1 | C_PVR_USER1 |

Table 1-23: Processor Version Register 1 (PVR1)

| Bits | Name | Description | Value |
|-------------|-------------|-------------------------|--------------|
| 0:31 | USR2 | User configured value 2 | C_PVR_USER2 |

Table 1-24: Processor Version Register 2 (PVR2)

| Bits | Name | Description | Value |
|-------------|-------------|--|--------------------------|
| 0 | Reserved | | 0 |
| 1 | DLMB | Data side LMB in use | C_D_LMB |
| 2 | Reserved | | 0 |
| 3 | ILMB | Instruction side LMB in use | C_I_LMB |
| 4 | IRQEDGE | Interrupt is edge triggered | C_INTERRUPT_IS_EDGE |
| 5 | IRQPOS | Interrupt edge is positive | C_EDGE_IS_POSITIVE |
| 6 | DPLB | Data side PLB in use | C_D_PLB |
| 7 | IPLB | Instruction side PLB in use | C_I_PLB |
| 8 | Reserved | | 1 |
| 9:11 | Reserved | | |
| 12 | FSL | Use extended FSL instructions | C_USE_EXTENDED_FSL_INSTR |
| 13 | FSLEXC | Generate exception for FSL control bit mismatch | C_FSL_EXCEPTION |
| 14 | MSR | Use msrset and msrclr instructions | C_USE_MSR_INSTR |
| 15 | PCMP | Use pattern compare instructions | C_USE_PCMP_INSTR |
| 16 | AREA | Select implementation to optimize area with lower instruction throughput | C_AREA_OPTIMIZED |
| 17 | BS | Use barrel shifter | C_USE_BARREL |
| 18 | DIV | Use divider | C_USE_DIV |
| 19 | MUL | Use hardware multiplier | C_USE_HW_MUL > 0 |
| 20 | FPU | Use FPU | C_USE_FPU > 0 |
| 21 | MUL64 | Use 64-bit hardware multiplier | C_USE_HW_MUL = 2 |
| 22 | FPU2 | Use floating point conversion and square root instructions | C_USE_FPU = 2 |

Table 1-24: Processor Version Register 2 (PVR2) (Continued)

| Bits | Name | Description | Value |
|-------------|-------------|--|------------------------|
| 23 | IPLBEXC | Generate exception for IPLB error | C_IPLB_BUS_EXCEPTION |
| 24 | DPLBEXC | Generate exception for DPLB error | C_DPLB_BUS_EXCEPTION |
| 25 | OP0EXC | Generate exception for 0x0 illegal opcode | C_OPCODE_0x0_ILLEGAL |
| 26 | UNEXC | Generate exception for unaligned data access | C_UNALIGNED_EXCEPTIONS |
| 27 | OPEXC | Generate exception for any illegal opcode | C_ILL_OPCODE_EXCEPTION |
| 28 | Reserved | | 0 |
| 29 | Reserved | | 0 |
| 30 | DIVEXC | Generate exception for division by zero or division overflow | C_DIV_ZERO_EXCEPTION |
| 31 | FPUEXC | Generate exceptions from FPU | C_FPU_EXCEPTION |

Table 1-25: Processor Version Register 3 (PVR3)

| Bits | Name | Description | Value |
|-------------|-------------|-------------------------------------|----------------------------|
| 0 | DEBUG | Use debug logic | C_DEBUG_ENABLED |
| 1:2 | Reserved | | |
| 3:6 | PCBRK | Number of PC breakpoints | C_NUMBER_OF_PC_BRK |
| 7:9 | Reserved | | |
| 10:12 | RDADDR | Number of read address breakpoints | C_NUMBER_OF_RD_ADDR_BRK |
| 13:15 | Reserved | | |
| 16:18 | WRADDR | Number of write address breakpoints | C_NUMBER_OF_WR_ADDR_BRK |
| 19 | Reserved | | |
| 20:24 | FSL | Number of FSLs | C_FSL_LINKS |
| 25:28 | Reserved | | |
| 29:31 | BTC_SIZE | Branch Target Cache Size | C_BRANCH_TARGET_CACHE_SIZE |

Table 1-26: Processor Version Register 4 (PVR4)

| Bits | Name | Description | Value |
|-------------|-------------|---|---------------------------------|
| 0 | ICU | Use instruction cache | C_USE_ICACHE |
| 1:5 | ICTS | Instruction cache tag size | C_ADDR_TAG_BITS |
| 6 | Reserved | | 1 |
| 7 | ICW | Allow instruction cache write | C_ALLOW_ICACHE_WR |
| 8:10 | ICLL | The base two logarithm of the instruction cache line length | log2(C_ICACHE_LINE_LEN) |
| 11:15 | ICBS | The base two logarithm of the instruction cache byte size | log2(C_CACHE_BYTE_SIZE) |
| 16 | IAU | The instruction cache is used for all memory accesses | C_ICACHE_ALWAYS_USED |
| 17 | Reserved | | 0 |
| 18 | ICI | Instruction cache XCL protocol | C_ICACHE_INTERFACE |
| 19:21 | ICV | Instruction cache victims | 0-3: C_ICACHE_VICTIMS = 0,2,4,8 |
| 22:23 | ICS | Instruction cache streams | C_ICACHE_STREAMS |
| 24:31 | Reserved | | 0 |

Table 1-27: Processor Version Register 5 (PVR5)

| Bits | Name | Description | Value |
|-------------|-------------|--|---------------------------------|
| 0 | DCU | Use data cache | C_USE_DCACHE |
| 1:5 | DCTS | Data cache tag size | C_DCACHE_ADDR_TAG |
| 6 | Reserved | | 1 |
| 7 | DCW | Allow data cache write | C_ALLOW_DCACHE_WR |
| 8:10 | DCLL | The base two logarithm of the data cache line length | log2(C_DCACHE_LINE_LEN) |
| 11:15 | DCBS | The base two logarithm of the data cache byte size | log2(C_DCACHE_BYTE_SIZE) |
| 16 | DAU | The data cache is used for all memory accesses | C_DCACHE_ALWAYS_USED |
| 17 | DWB | Data cache policy is write-back | C_DCACHE_USE_WRITEBACK |
| 18 | DCI | Data cache XCL protocol | C_DCACHE_INTERFACE |
| 19:21 | DCV | Data cache victims | 0-3: C_DCACHE_VICTIMS = 0,2,4,8 |
| 22:31 | Reserved | | 0 |

Table 1-28: Processor Version Register 6 (PVR6)

| Bits | Name | Description | Value |
|-------------|-------------|--------------------------------|-------------------|
| 0:31 | ICBA | Instruction Cache Base Address | C_ICACHE_BASEADDR |

Table 1-29: Processor Version Register 7 (PVR7)

| Bits | Name | Description | Value |
|-------------|-------------|--------------------------------|-------------------|
| 0:31 | ICHA | Instruction Cache High Address | C_ICACHE_HIGHADDR |

Table 1-30: Processor Version Register 8 (PVR8)

| Bits | Name | Description | Value |
|-------------|-------------|-------------------------|-------------------|
| 0:31 | DCBA | Data Cache Base Address | C_DCACHE_BASEADDR |

Table 1-31: Processor Version Register 9 (PVR9)

| Bits | Name | Description | Value |
|-------------|-------------|-------------------------|-------------------|
| 0:31 | DCHA | Data Cache High Address | C_DCACHE_HIGHADDR |

Table 1-32: Processor Version Register 10 (PVR10)

| Bits | Name | Description | Value |
|-------------|-------------|---|----------------------------------|
| 0:7 | ARCH | Target architecture: 0x6 = Spartan®-3, Automotive Spartan-3 0x7 = Virtex-4, QPro Virtex-4 Hi-Rel, QPro Virtex-4 Rad Tolerant 0x8 = Virtex-5 0x9 = Spartan-3E, Automotive Spartan-3E 0xA = Spartan-3A, Automotive Spartan-3A 0xB = Spartan-3AN 0xC = Spartan-3A DSP, Automotive Spartan-3A DSP 0xD = Spartan-6, Automotive Spartan-6, QPro Spartan-6 Hi-Rel 0xE = Virtex-6 | Defined by parameter C_FAMILY |
| 8:31 | Reserved | | 0 |

Table 1-33: Processor Version Register 11 (PVR11)

| Bits | Name | Description | Value |
|-------|----------|--|--------------------------------------|
| 0:1 | MMU | Use MMU: 0 = None 1 = User Mode 2 = Protection 3 = Virtual | C_USE_MMU |
| 2:4 | ITLB | Instruction Shadow TLB size | $\log_2(C_{\text{MMU_ITLB_SIZE}})$ |
| 5:7 | DTLB | Data Shadow TLB size | $\log_2(C_{\text{MMU_DTLB_SIZE}})$ |
| 8:9 | TLBACC | TLB register access: 0 = Minimal 1 = Read 2 = Write 3 = Full | C_MMU_TLB_ACCESS |
| 10:14 | ZONES | Number of memory protection zones | C_MMU_ZONES |
| 15:16 | Reserved | Reserved for future use | 0 |
| 17:31 | RSTMSR | Reset value for MSR | C_RESET_MSR |

Pipeline Architecture

MicroBlaze instruction execution is pipelined. For most instructions, each stage takes one clock cycle to complete. Consequently, the number of clock cycles necessary for a specific instruction to complete is equal to the number of pipeline stages, and one instruction is completed on every cycle. A few instructions require multiple clock cycles in the execute stage to complete. This is achieved by stalling the pipeline.

When executing from slower memory, instruction fetches may take multiple cycles. This additional latency directly affects the efficiency of the pipeline. MicroBlaze implements an instruction prefetch buffer that reduces the impact of such multi-cycle instruction memory latency. While the pipeline is stalled by a multi-cycle instruction in the execution stage, the prefetch buffer continues to load sequential instructions. When the pipeline resumes execution, the fetch stage can load new instructions directly from the prefetch buffer instead of waiting for the instruction memory access to complete. If instructions are modified during execution (e.g. with self-modifying code), the prefetch buffer should be emptied before executing the modified instructions, to ensure that it does not contain the old unmodified instructions. This can be done with a synchronizing branch instruction, for example BRI 4.

Three Stage Pipeline

With C_AREA_OPTIMIZED set to 1, the pipeline is divided into three stages to minimize hardware cost: Fetch, Decode, and Execute.

| | cycle 1 | cycle 2 | cycle 3 | cycle4 | cycle5 | cycle6 | cycle7 |
|---------------|---------|---------|---------|---------|---------|---------|---------|
| instruction 1 | Fetch | Decode | Execute | | | | |
| instruction 2 | | Fetch | Decode | Execute | Execute | Execute | |
| instruction 3 | | | Fetch | Decode | Stall | Stall | Execute |

Five Stage Pipeline

With C_AREA_OPTIMIZED set to 0, the pipeline is divided into five stages to maximize performance: Fetch (IF), Decode (OF), Execute (EX), Access Memory (MEM), and Writeback (WB).

| | cycle 1 | cycle 2 | cycle 3 | cycle4 | cycle5 | cycle6 | cycle7 | cycle8 | cycle9 |
|---------------|------------|------------|------------|--------|--------|--------|--------|--------|--------|
| instruction 1 | IF | OF | EX | MEM | WB | | | | |
| instruction 2 | | IF | OF | EX | MEM | MEM | MEM | WB | |
| instruction 3 | | | IF | OF | EX | Stall | Stall | MEM | WB |

Branches

Normally the instructions in the fetch and decode stages (as well as prefetch buffer) are flushed when executing a taken branch. The fetch pipeline stage is then reloaded with a new instruction from the calculated branch address. A taken branch in MicroBlaze takes three clock cycles to execute, two of which are required for refilling the pipeline. To reduce this latency overhead, MicroBlaze supports branches with delay slots.

Delay Slots

When executing a taken branch with delay slot, only the fetch pipeline stage in MicroBlaze is flushed. The instruction in the decode stage (branch delay slot) is allowed to complete. This technique effectively reduces the branch penalty from two clock cycles to one. Branch instructions with delay slots have a D appended to the instruction mnemonic. For example, the BNE instruction does not execute the subsequent instruction (does not have a delay slot), whereas BNED executes the next instruction before control is transferred to the branch location.

A delay slot must not contain the following instructions: IMM, branch, or break. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

Instructions that could cause recoverable exceptions (e.g. unaligned word or halfword load and store) are allowed in the delay slot. If an exception is caused in a delay slot the ESR[DS] bit is set, and the exception handler is responsible for returning the execution to the branch target (stored in the special purpose register BTR). If the ESR[DS] bit is set, register R17 is not valid (otherwise it contains the address following the instruction causing the exception).

Branch Target Cache

To improve branch performance, MicroBlaze provides a Branch Target Cache (BTC) coupled with a branch prediction scheme. With the BTC enabled, a correctly predicted immediate branch or return instruction incurs no overhead.

The BTC operates by saving the target address of each immediate branch and return instruction the first time the instruction is encountered. The next time it is encountered, it is usually found in the Branch Target Cache, and the Instruction Fetch Program Counter is then simply changed to the saved target address, in case the branch should be taken. Unconditional branches and return instructions are always taken, whereas conditional branches use branch prediction, to avoid taking a branch that should not have been taken and vice versa.

There are three cases where the branch prediction can cause a mispredict, namely:

- A conditional branch that should not have been taken, is actually taken,
- A conditional branch that should actually have been taken, is not taken,
- The target address of a return instruction is incorrect, which may occur when returning from a function called from different places in the code.

All of these cases are detected and corrected when the branch or return instruction reaches the execute stage, and the branch prediction bits or target address are updated in the BTC, to reflect the actual instruction behavior. This correction incurs a penalty of two clock cycles.

The size of the BTC can be selected with C_BRANCH_TARGET_CACHE_SIZE. The default recommended setting uses one block RAM, and provides either 512 entries (for Virtex-5 or Virtex-6) or 256 entries (for all other families). When selecting 64 entries or below, distributed RAM is used to implement the BTC, otherwise block RAM is used.

The Branch Target Cache is available when C_USE_BRANCH_TARGET_CACHE is set to 1 and C_AREA_OPTIMIZED is set to 0.

Memory Architecture

MicroBlaze is implemented with a Harvard memory architecture; instruction and data accesses are done in separate address spaces. Each address space has a 32-bit range (that is, handles up to 4-GB of instructions and data memory respectively). The instruction and data memory ranges can be made to overlap by mapping them both to the same physical memory. The latter is useful for software debugging.

Both instruction and data interfaces of MicroBlaze are 32 bits wide and use big endian, bit-reversed format. MicroBlaze supports word, halfword, and byte accesses to data memory.

Data accesses must be aligned (word accesses must be on word boundaries, halfword on halfword boundaries), unless the processor is configured to support unaligned exceptions. All instruction accesses must be word aligned.

MicroBlaze does not separate data accesses to I/O and memory (it uses memory mapped I/O). The processor has up to three interfaces for memory accesses:

- Local Memory Bus (LMB)
- Processor Local Bus (PLB)
- Xilinx CacheLink (XCL)

The LMB memory address range must not overlap with PLB or XCL ranges.

MicroBlaze has a single cycle latency for accesses to local memory (LMB) and for cache read hits, except with C_AREA_OPTIMIZED set to 1, when data side accesses and data cache read hits require two clock cycles. The data cache write latency depends on C_DCACHE_USE_WRITEBACK. When C_DCACHE_USE_WRITEBACK is set to 1, the write latency normally is one cycle (more if the cache needs to do memory accesses). When C_DCACHE_USE_WRITEBACK is cleared to 0, the write latency normally is two cycles (more if the posted-write buffer in the memory controller is full).

The MicroBlaze instruction and data caches can be configured to use 4 or 8 word cache lines. When using a longer cache line, more bytes are prefetched, which generally improves performance for software with sequential access patterns. However, for software with a more random access pattern the performance can instead decrease for a given cache size. This is caused by a reduced cache hit rate due to fewer available cache lines.

For details on the different memory interfaces refer to [Chapter 2, “MicroBlaze Signal Interface Description”](#).

Privileged Instructions

The following MicroBlaze instructions are privileged:

- GET, PUT, NGET, NPUT, CGET, CPUT, NCGET, NCPUT
- WIC, WDC
- MTS
- MSRCLR, MSRSET (except when only the C bit is affected)
- BRK
- RTID, RTBD, RTED
- BRKI (except when jumping to physical address 0x8 or 0x18)

Attempted use of these instructions when running in user mode causes a privileged instruction exception.

There are six ways to leave user mode and virtual mode:

1. Hardware generated reset (including debug reset)
2. Hardware exception
3. Non-maskable break or hardware break
4. Interrupt
5. Executing the instruction "BRALID Re , 0x8" to perform a user vector exception
6. Executing the software break instructions "BRKI" jumping to physical address 0x8 or 0x18

In all of these cases, except hardware generated reset, the user mode and virtual mode status is saved in the MSR UMS and VMS bits.

Application (user-mode) programs transfer control to system-service routines (privileged mode programs) using the BRALID or BRKI instruction, jumping to physical address 0x8. Executing this instruction causes a system-call exception to occur. The exception handler determines which system-service routine to call and whether the calling application has permission to call that service. If permission is granted, the exception handler performs the actual procedure call to the system-service routine on behalf of the application program.

The execution environment expected by the system-service routine requires the execution of prologue instructions to set up that environment. Those instructions usually create the block of storage that holds procedural information (the activation record), update and initialize pointers, and save volatile registers (registers the system-service routine uses). Prologue code can be inserted by the linker when creating an executable module, or it can be included as stub code in either the system-call interrupt handler or the system-library routines.

Returns from the system-service routine reverse the process described above. Epilog code is executed to unwind and deallocate the activation record, restore pointers, and restore volatile registers. The interrupt handler executes a return from exception instruction (RTED) to return to the application.

Virtual-Memory Management

Programs running on MicroBlaze use effective addresses to access a flat 4 GB address space. The processor can interpret this address space in one of two ways, depending on the translation mode:

- In real mode, effective addresses are used to directly access physical memory
- In virtual mode, effective addresses are translated into physical addresses by the virtual-memory management hardware in the processor

Virtual mode provides system software with the ability to relocate programs and data anywhere in the physical address space. System software can move inactive programs and data out of physical memory when space is required by active programs and data.

Relocation can make it appear to a program that more memory exists than is actually implemented by the system. This frees the programmer from working within the limits imposed by the amount of physical memory present in a system. Programmers do not need to know which physical-memory addresses are assigned to other software processes and hardware devices. The addresses visible to programs are translated into the appropriate physical addresses by the processor.

Virtual mode provides greater control over memory protection. Blocks of memory as small as 1 KB can be individually protected from unauthorized access. Protection and relocation enable system software to support multitasking. This capability gives the appearance of simultaneous or near-simultaneous execution of multiple programs.

In MicroBlaze, virtual mode is implemented by the memory-management unit (MMU), available when C_USE_MMU is set to 3 and C_AREA_OPTIMIZED is set to 0. The MMU controls effective-address to physical-address mapping and supports memory protection. Using these capabilities, system software can implement demand-paged virtual memory and other memory management schemes.

The MicroBlaze MMU implementation is based upon PowerPC™ 405. For details, see the *PowerPC Processor Reference Guide* document.

The MMU features are summarized as follows:

- Translates effective addresses into physical addresses
- Controls page-level access during address translation
- Provides additional virtual-mode protection control through the use of zones
- Provides independent control over instruction-address and data-address translation and protection
- Supports eight page sizes: 1 kB, 4 kB, 16 kB, 64 kB, 256 kB, 1 MB, 4 MB, and 16 MB. Any combination of page sizes can be used by system software
- Software controls the page-replacement strategy

Real Mode

The processor references memory when it fetches an instruction and when it accesses data with a load or store instruction. Programs reference memory locations using a 32-bit effective address calculated by the processor. When real mode is enabled, the physical address is identical to the effective address and the processor uses it to access physical memory. After a processor reset, the processor operates in real mode. Real mode can also be enabled by clearing the VM bit in the MSR.

Physical-memory data accesses (loads and stores) are performed in real mode using the effective address. Real mode does not provide system software with virtual address translation, but the full memory access-protection is available, implemented when C_USE_MMU > 1 and C_AREA_OPTIMIZED = 0. Implementation of a real-mode memory manager is more

straightforward than a virtual-mode memory manager. Real mode is often an appropriate solution for memory management in simple embedded environments, when access-protection is necessary, but virtual address translation is not required.

Virtual Mode

In virtual mode, the processor translates an effective address into a physical address using the process shown in [Figure 1-16](#). Virtual mode can be enabled by setting the VM bit in the MSR..

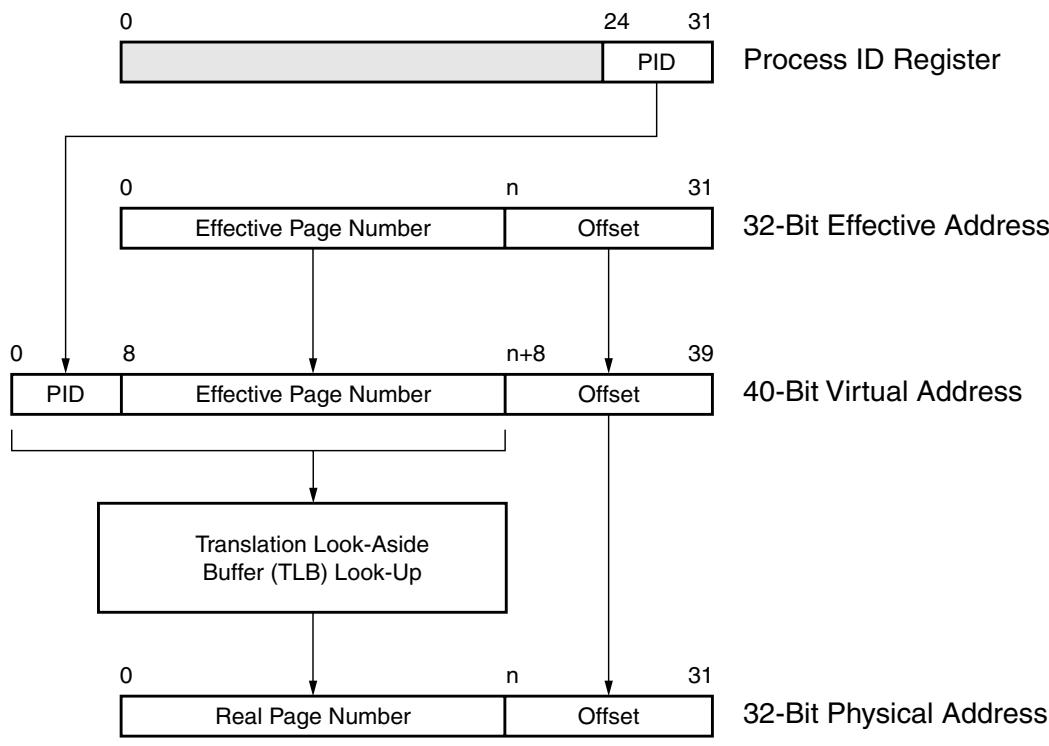


Figure 1-16: Virtual-Mode Address Translation

Each address shown in [Figure 1-16](#) contains a page-number field and an offset field. The page number represents the portion of the address translated by the MMU. The offset represents the byte offset into a page and is not translated by the MMU. The virtual address consists of an additional field, called the process ID (PID), which is taken from the PID register (see Process-ID Register, page 37). The combination of PID and effective page number (EPN) is referred to as the virtual page number (VPN). The value n is determined by the page size, as shown in [Table 1-34](#).

System software maintains a page-translation table that contains entries used to translate each virtual page into a physical page. The page size defined by a page translation entry determines the size of the page number and offset fields. For example, when a 4 kB page size is used, the page-number field is 20 bits and the offset field is 12 bits. The VPN in this case is 28 bits.

Then the most frequently used page translations are stored in the translation look-aside buffer (TLB). When translating a virtual address, the MMU examines the page-translation entries for a matching VPN (PID and EPN). Rather than examining all entries in the table, only entries contained in the processor TLB are examined. When a page-translation entry is found with a matching VPN, the corresponding physical-page number is read from the entry and combined with the offset to form the 32-bit physical address. This physical address is used by the processor to reference memory.

System software can use the PID to uniquely identify software processes (tasks, subroutines, threads) running on the processor. Independently compiled processes can operate in effective-address regions that overlap each other. This overlap must be resolved by system software if multitasking is supported. Assigning a PID to each process enables system software to resolve the overlap by relocating each process into a unique region of virtual-address space. The virtual-address space mappings enable independent translation of each process into the physical-address space.

Page-Translation Table

The page-translation table is a software-defined and software-managed data structure containing page translations. The requirement for software-managed page translation represents an architectural trade-off targeted at embedded-system applications. Embedded systems tend to have a tightly controlled operating environment and a well-defined set of application software. That environment enables virtual-memory management to be optimized for each embedded system in the following ways:

- The page-translation table can be organized to maximize page-table search performance (also called table walking) so that a given page-translation entry is located quickly. Most general-purpose processors implement either an indexed page table (simple search method, large page-table size) or a hashed page table (complex search method, small page-table size). With software table walking, any hybrid organization can be employed that suits the particular embedded system. Both the page-table size and access time can be optimized.
- Independent page sizes can be used for application modules, device drivers, system service routines, and data. Independent page-size selection enables system software to more efficiently use memory by reducing fragmentation (unused memory). For example, a large data structure can be allocated to a 16 MB page and a small I/O device-driver can be allocated to a 1 KB page.
- Page replacement can be tuned to minimize the occurrence of missing page translations. As described in the following section, the most-frequently used page translations are stored in the translation look-aside buffer (TLB). Software is responsible for deciding which translations are stored in the TLB and which translations are replaced when a new translation is required. The replacement strategy can be tuned to avoid thrashing, whereby page-translation entries are constantly being moved in and out of the TLB. The replacement strategy can also be tuned to prevent replacement of critical-page translations, a process sometimes referred to as page locking.

The unified 64-entry TLB, managed by software, caches a subset of instruction and data page-translation entries accessible by the MMU. Software is responsible for reading entries from the page-translation table in system memory and storing them in the TLB. The following section describes the unified TLB in more detail. Internally, the MMU also contains shadow TLBs for instructions and data, with sizes configurable by `C_MMU_ITLB_SIZE` and `C_MMU_DTLB_SIZE` respectively.

These shadow TLBs are managed entirely by the processor (transparent to software) and are used to minimize access conflicts with the unified TLB.

Translation Look-Aside Buffer

The translation look-aside buffer (TLB) is used by the MicroBlaze MMU for address translation when the processor is running in virtual mode, memory protection, and storage control. Each entry within the TLB contains the information necessary to identify a virtual page (PID and effective page number), specify its translation into a physical page, determine the protection characteristics of the page, and specify the storage attributes associated with the page.

The MicroBlaze TLB is physically implemented as three separate TLBs:

- Unified TLB—The UTLB contains 64 entries and is pseudo-associative. Instruction-page and data-page translation can be stored in any UTLB entry. The initialization and management of the UTLB is controlled completely by software.
- Instruction Shadow TLB—The ITLB contains instruction page-translation entries and is fully associative. The page-translation entries stored in the ITLB represent the most-recently accessed instruction-page translations from the UTLB. The ITLB is used to minimize contention between instruction translation and UTLB-update operations. The initialization and management of the ITLB is controlled completely by hardware and is transparent to software.
- Data Shadow TLB—The DTLB contains data page-translation entries and is fully associative. The page-translation entries stored in the DTLB represent the most-recently accessed data-page translations from the UTLB. The DTLB is used to minimize contention between data translation and UTLB-update operations. The initialization and management of the DTLB is controlled completely by hardware and is transparent to software.

Figure 1-17 provides the translation flow for TLB.

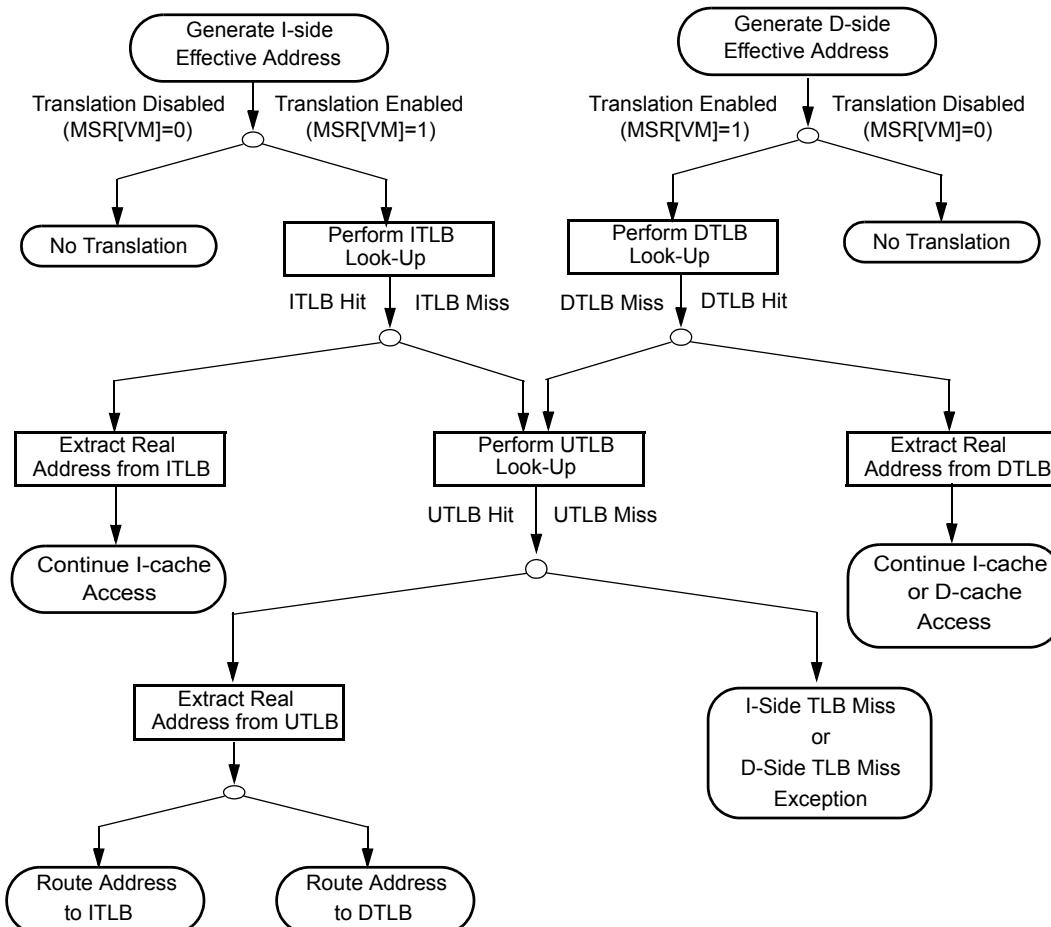
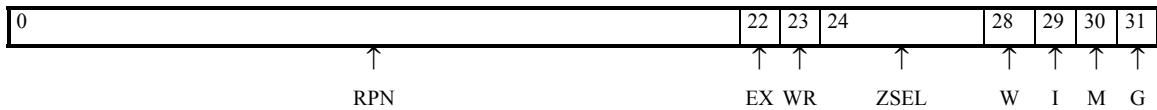


Figure 1-17: TLB Address Translation Flow

TLB Entry Format

Figure 1-18 shows the format of a TLB entry. Each TLB entry is 68 bits and is composed of two portions: TLBLO (also referred to as the data entry), and TLBHI (also referred to as the tag entry).

TLBLO:



TLBHI:

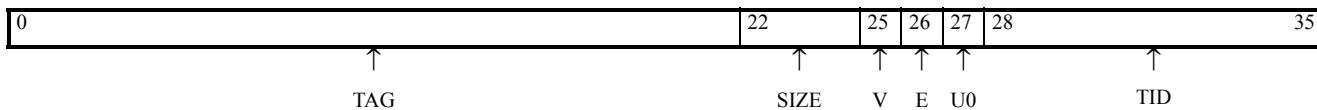


Figure 1-18: **TLB Entry Format**

The TLB entry contents are described in [Table 1-18, page 39](#) and [Table 1-19, page 41](#).

The fields within a TLB entry are categorized as follows:

- Virtual-page identification (TAG, SIZE, V, TID)—These fields identify the page-translation entry. They are compared with the virtual-page number during the translation process.
- Physical-page identification (RPN, SIZE)—These fields identify the translated page in physical memory.
- Access control (EX, WR, ZSEL)—These fields specify the type of access allowed in the page and are used to protect pages from improper accesses.
- Storage attributes (W, I, M, G, E, U0)—These fields specify the storage-control attributes, such as caching policy for the data cache (write-back or write-through), whether a page is cacheable, and how bytes are ordered (endianness).

[Table 1-34](#) shows the relationship between the TLB-entry SIZE field and the translated page size. This table also shows how the page size determines which address bits are involved in a tag comparison, which address bits are used as a page offset, and which bits in the physical page number are used in the physical address.

Table 1-34: Page-Translation Bit Ranges by Page Size

| Page Size | SIZE (TLBHI Field) | Tag Comparison Bit Range | Page Offset | Physical Page Number | RPN Bits Clear to 0 |
|-----------|--------------------|---------------------------|----------------|----------------------|---------------------|
| 1 KB | 000 | TAG[0:21] - Address[0:21] | Address[22:31] | RPN[0:21] | - |
| 4 KB | 001 | TAG[0:19] - Address[0:19] | Address[20:31] | RPN[0:19] | 20:21 |
| 16 KB | 010 | TAG[0:17] - Address[0:17] | Address[18:31] | RPN[0:17] | 18:21 |
| 64 KB | 011 | TAG[0:15] - Address[0:15] | Address[16:31] | RPN[0:15] | 16:21 |
| 256 KB | 100 | TAG[0:13] - Address[0:13] | Address[14:31] | RPN[0:13] | 14:21 |
| 1 MB | 101 | TAG[0:11] - Address[0:11] | Address[12:31] | RPN[0:11] | 12:21 |
| 4 MB | 110 | TAG[0:9] - Address[0:9] | Address[10:31] | RPN[0:9] | 10:21 |
| 16 MB | 111 | TAG[0:7] - Address[0:7] | Address[8:31] | RPN[0:7] | 8:21 |

TLB Access

When the MMU translates a virtual address (the combination of PID and effective address) into a physical address, it first examines the appropriate shadow TLB for the page translation entry. If an entry is found, it is used to access physical memory. If an entry is not found, the MMU examines the UTLB for the entry. A delay occurs each time the UTLB must be accessed due to a shadow TLB miss. The miss latency ranges from 2-32 cycles. The DTLB has priority over the ITLB if both simultaneously access the UTLB.

[Figure 1-19, page 61](#) shows the logical process the MMU follows when examining a page-translation entry in one of the shadow TLBs or the UTLB. All valid entries in the TLB are checked.

A TLB hit occurs when all of the following conditions are met by a TLB entry:

- The entry is valid
- The TAG field in the entry matches the effective address EPN under the control of the SIZE field in the entry
- The TID field in the entry matches the PID

If any of the above conditions are not met, a TLB miss occurs. A TLB miss causes an exception, described as follows:

A TID value of 0x00 causes the MMU to ignore the comparison between the TID and PID. Only the TAG and EA[EPN] are compared. A TLB entry with TID=0x00 represents a process-independent translation. Pages that are accessed globally by all processes should be assigned a TID value of 0x00. A PID value of 0x00 does not identify a process that can access any page. When PID=0x00, a page-translation hit only occurs when TID=0x00. It is possible for software to load the TLB with multiple entries that match an EA[EPN] and PID combination. However, this is considered a programming error and results in undefined behavior.

When a hit occurs, the MMU reads the RPN field from the corresponding TLB entry. Some or all of the bits in this field are used, depending on the value of the SIZE field (see [Table 1-34](#)). For example, if the SIZE field specifies a 256 kB page size, RPN[0:13] represents the physical page number and is used to form the physical address. RPN[14:21] is not used, and software must clear those bits to 0 when initializing the TLB entry. The remainder of the physical address is taken from the page-offset portion of the EA. If the page size is 256 kB, the 32-bit physical address is formed by concatenating RPN[0:13] with bits14:31 of the effective address.

Prior to accessing physical memory, the MMU examines the TLB-entry access-control fields. These fields indicate whether the currently executing program is allowed to perform the requested memory access.

If access is allowed, the MMU checks the storage-attribute fields to determine how to access the page. The storage-attribute fields specify the caching policy for memory accesses.

TLB Access Failures

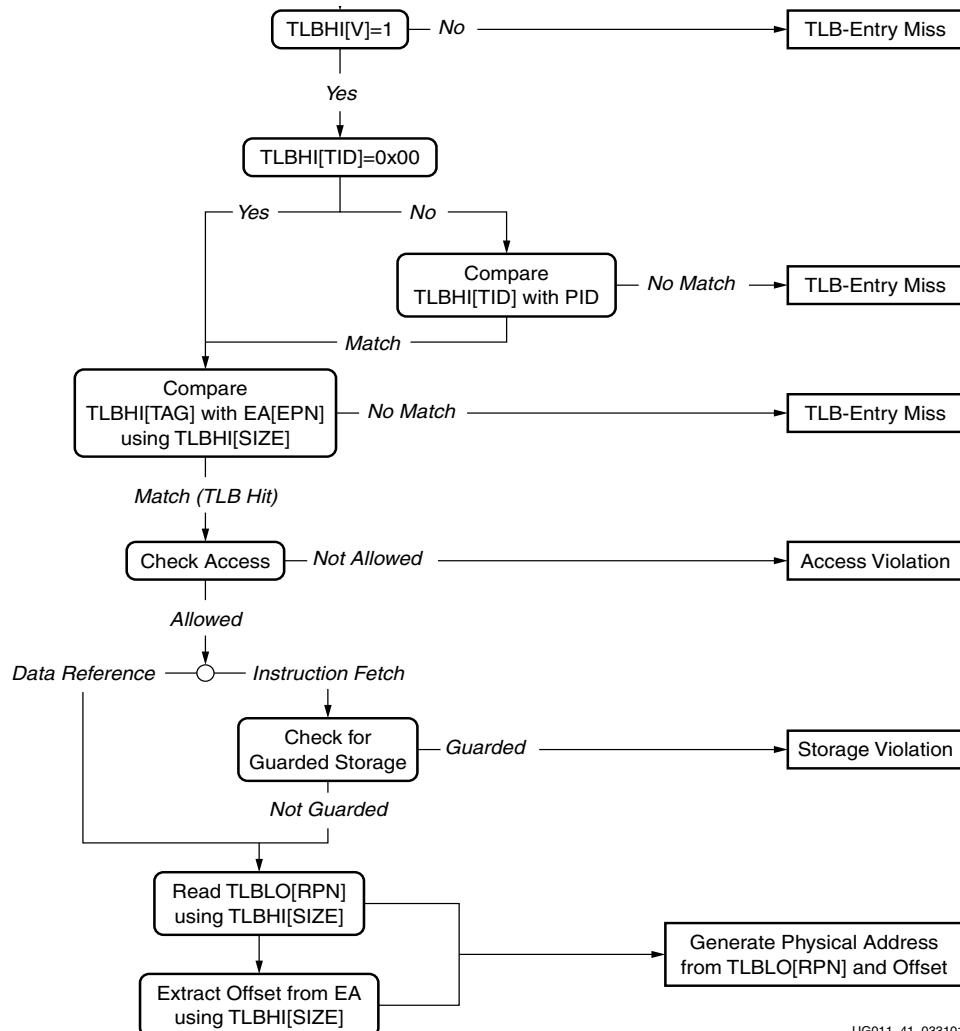
A TLB-access failure causes an exception to occur. This interrupts execution of the instruction that caused the failure and transfers control to an interrupt handler to resolve the failure. A TLB access can fail for two reasons:

- A matching TLB entry was not found, resulting in a TLB miss
- A matching TLB entry was found, but access to the page was prevented by either the storage attributes or zone protection

When an interrupt occurs, the processor enters real mode by clearing MSR[VM] to 0. In real mode, all address translation and memory-protection checks performed by the MMU are disabled. After

system software initializes the UTLB with page-translation entries, management of the MicroBlaze UTLB is usually performed using interrupt handlers running in real mode.

Figure 1-19 diagrams the general process for examining a TLB entry.



UG011_41_033101

Figure 1-19: General Process for Examining a TLB Entry

The following sections describe the conditions under which exceptions occur due to TLB access failures.

Data-Storage Exception

When virtual mode is enabled, ($\text{MSR}[\text{VM}] = 1$), a data-storage exception occurs when access to a page is not permitted for any of the following reasons:

- From user mode:
 - ◆ The TLB entry specifies a zone field that prevents access to the page ($\text{ZPR}[\text{Zn}] = 00$). This applies to load and store instructions.
 - ◆ The TLB entry specifies a read-only page ($\text{TLBLO}[\text{WR}] = 0$) that is not otherwise overridden by the zone field ($\text{ZPR}[\text{Zn}], 11$). This applies to store instructions.

- From privileged mode:
 - ◆ The TLB entry specifies a read-only page (TLBLO[WR]=0) that is not otherwise overridden by the zone field (ZPR[Zn], 10 and ZPR[Zn], 11). This applies to store instructions.

Instruction-Storage Exception

When virtual mode is enabled, (MSR[VM]=1), an instruction-storage exception occurs when access to a page is not permitted for any of the following reasons:

- From user mode:
 - ◆ The TLB entry specifies a zone field that prevents access to the page (ZPR[Zn]=00).
 - ◆ The TLB entry specifies a non-executable page (TLBLO[EX]=0) that is not otherwise overridden by the zone field (ZPR[Zn], 11).
 - ◆ The TLB entry specifies a guarded-storage page (TLBLO[G]=1).
- From privileged mode:
 - ◆ The TLB entry specifies a non-executable page (TLBLO[EX]=0) that is not otherwise overridden by the zone field (ZPR[Zn], 10 and ZPR[Zn], 11).
 - ◆ The TLB entry specifies a guarded-storage page (TLBLO[G]=1).

Data TLB-Miss Exception

When virtual mode is enabled (MSR[VM]=1) a data TLB-miss exception occurs if a valid, matching TLB entry was not found in the TLB (shadow and UTLB). Any load or store instruction can cause a data TLB-miss exception.

Instruction TLB-Miss Exception

When virtual mode is enabled (MSR[VM]=1) an instruction TLB-miss exception occurs if a valid, matching TLB entry was not found in the TLB (shadow and UTLB). Any instruction fetch can cause an instruction TLB-miss exception.

Access Protection

System software uses access protection to protect sensitive memory locations from improper access. System software can restrict memory accesses for both user-mode and privileged-mode software. Restrictions can be placed on reads, writes, and instruction fetches. Access protection is available when virtual protected mode is enabled.

Access control applies to instruction fetches, data loads, and data stores. The TLB entry for a virtual page specifies the type of access allowed to the page. The TLB entry also specifies a zone-protection field in the zone-protection register that is used to override the access controls specified by the TLB entry.

TLB Access-Protection Controls

Each TLB entry controls three types of access:

- Process—Processes are protected from unauthorized access by assigning a unique process ID (PID) to each process. When system software starts a user-mode application, it loads the PID for that application into the PID register. As the application executes, memory addresses are translated using only TLB entries with a TID field in Translation Look-Aside Buffer High (TLBHI) that matches the PID. This enables system software to restrict accesses for an application to a specific area in virtual memory.

A TLB entry with TID=0x00 represents a process-independent translation. Pages that are accessed globally by all processes should be assigned a TID value of 0x00.

- Execution—The processor executes instructions only if they are fetched from a virtual page marked as executable (TLBLO[EX]=1). Clearing TLBLO[EX] to 0 prevents execution of instructions fetched from a page, instead causing an instruction-storage interrupt (ISI) to occur. The ISI does not occur when the instruction is fetched, but instead occurs when the instruction is executed. This prevents speculatively fetched instructions that are later discarded (rather than executed) from causing an ISI.

The zone-protection register can override execution protection.

- Read/Write—Data is written only to virtual pages marked as writable (TLBLO[WR]=1). Clearing TLBLO[WR] to 0 marks a page as read-only. An attempt to write to a read-only page causes a data-storage interrupt (DSI) to occur.

The zone-protection register can override write protection.

TLB entries cannot be used to prevent programs from reading pages. In virtual mode, zone protection is used to read-protect pages. This is done by defining a no-access-allowed zone ($ZPR[Zn] = 00$) and using it to override the TLB-entry access protection. Only programs running in user mode can be prevented from reading a page. Privileged programs always have read access to a page.

Zone Protection

Zone protection is used to override the access protection specified in a TLB entry. Zones are an arbitrary grouping of virtual pages with common access protection. Zones can contain any number of pages specifying any combination of page sizes. There is no requirement for a zone to contain adjacent pages.

The zone-protection register (ZPR) is a 32-bit register used to specify the type of protection override applied to each of 16 possible zones. The protection override for a zone is encoded in the ZPR as a 2-bit field. The 4-bit zone-select field in a TLB entry (TLBLO[ZSEL]) selects one of the 16 zone fields from the ZPR (Z0–Z15). For example, zone Z5 is selected when ZSEL = 0101.

Changing a zone field in the ZPR applies a protection override across all pages in that zone. Without the ZPR, protection changes require individual alterations to each page translation entry within the zone.

UTLB Management

The UTLB serves as the interface between the processor MMU and memory-management software. System software manages the UTLB to tell the MMU how to translate virtual addresses into physical addresses. When a problem occurs due to a missing translation or an access violation, the MMU communicates the problem to system software using the exception mechanism. System software is responsible for providing interrupt handlers to correct these problems so that the MMU can proceed with memory translation.

Software reads and writes UTLB entries using the MFS and MTS instructions, respectively. These instructions use the TLBX register index (numbered 0 to 63) corresponding to one of the 64 entries in the UTLB. The tag and data portions are read and written separately, so software must execute two MFS or MTS instructions to completely access an entry. The UTLB is searched for a specific translation using the TLBSX register. TLBSX locates a translation using an effective address and loads the corresponding UTLB index into the TLBX register.

Individual UTLB entries are invalidated using the MTS instruction to clear the valid bit in the tag portion of a TLB entry (TLBHI[V]).

Recording Page Access and Page Modification

Software management of virtual-memory poses several challenges:

- In a virtual-memory environment, software and data often consume more memory than is physically available. Some of the software and data pages must be stored outside physical memory, such as on a hard drive, when they are not used. Ideally, the most-frequently used pages stay in physical memory and infrequently used pages are stored elsewhere.
- When pages in physical-memory are replaced to make room for new pages, it is important to know whether the replaced (old) pages were modified. If they were modified, they must be saved prior to loading the replacement (new) pages. If the old pages were not modified, the new pages can be loaded without saving the old pages.
- A limited number of page translations are kept in the UTLB. The remaining translations must be stored in the page-translation table. When a translation is not found in the UTLB (due to a miss), system software must decide which UTLB entry to discard so that the missing translation can be loaded. It is desirable for system software to replace infrequently used translations rather than frequently used translations.

Solving the above problems in an efficient manner requires keeping track of page accesses and page modifications. MicroBlaze does not track page access and page modification in hardware. Instead, system software can use the TLB-miss exceptions and the data-storage exception to collect this information. As the information is collected, it can be stored in a data structure associated with the page-translation table.

Page-access information is used to determine which pages should be kept in physical memory and which are replaced when physical-memory space is required. System software can use the valid bit in the TLB entry (TLBHI[V]) to monitor page accesses. This requires page translations be initialized as not valid (TLBHI[V]=0) to indicate they have not been accessed. The first attempt to access a page causes a TLB-miss exception, either because the UTLB entry is marked not valid or because the page translation is not present in the UTLB. The TLB-miss handler updates the UTLB with a valid translation (TLBHI[V]=1). The set valid bit serves as a record that the page and its translation have been accessed. The TLB-miss handler can also record the information in a separate data structure associated with the page-translation entry.

Page-modification information is used to indicate whether an old page can be overwritten with a new page or the old page must first be stored to a hard disk. System software can use the write-protection bit in the TLB entry (TLBLO[WR]) to monitor page modification. This requires page translations be initialized as read-only (TLBLO[WR]=0) to indicate they have not been modified. The first attempt to write data into a page causes a data-storage exception, assuming the page has already been accessed and marked valid as described above. If software has permission to write into the page, the data-storage handler marks the page as writable (TLBLO[WR]=1) and returns. The set write-protection bit serves as a record that a page has been modified. The data-storage handler can also record this information in a separate data structure associated with the page-translation entry.

Tracking page modification is useful when virtual mode is first entered and when a new process is started.

Reset, Interrupts, Exceptions, and Break

MicroBlaze supports reset, interrupt, user exception, break, and hardware exceptions. The following section describes the execution flow associated with each of these events.

The relative priority starting with the highest is:

1. Reset
2. Hardware Exception
3. Non-maskable Break
4. Break
5. Interrupt
6. User Vector (Exception)

Table 1-35 defines the memory address locations of the associated vectors and the hardware enforced register file locations for return addresses. Each vector allocates two addresses to allow full address range branching (requires an IMM followed by a BRAI instruction). The address range 0x28 to 0x4F is reserved for future software support by Xilinx. Allocating these addresses for user applications is likely to conflict with future releases of EDK support software.

Table 1-35: Vectors and Return Address Register File Location

| Event | Vector Address | Register File Return Address |
|-----------------------------------|-------------------------|---------------------------------|
| Reset | 0x00000000 - 0x00000004 | - |
| User Vector (Exception) | 0x00000008 - 0x0000000C | Rx |
| Interrupt | 0x00000010 - 0x00000014 | R14 |
| Break: Non-maskable hardware | 0x00000018 - 0x0000001C | R16 |
| Break: Hardware | | |
| Break: Software | | |
| Hardware Exception | 0x00000020 - 0x00000024 | R17 or BTR |
| Reserved by Xilinx for future use | 0x00000028 - 0x0000004F | - |

All of these events will clear the reservation bit, used together with the LWX and SWX instructions to implement mutual exclusion, such as semaphores and spinlocks.

Reset

When a `Reset` or `Debug_Rst`⁽¹⁾ occurs, MicroBlaze flushes the pipeline and starts fetching instructions from the reset vector (address 0x0). Both external reset signals are active high and should be asserted for a minimum of 16 cycles.

Equivalent Pseudocode

```

PC ← 0x00000000
MSR ← C_RESET_MSR (see "MicroBlaze Core Configurability" in Chapter 2)
EAR ← 0; ESR ← 0; FSR ← 0
PID ← 0; ZPR ← 0; TLBX ← 0
Reservation ← 0

```

1. Reset input controlled by the XMD debugger via MDM.

Hardware Exceptions

MicroBlaze can be configured to trap the following internal error conditions: illegal instruction, instruction and data bus error, and unaligned access. The divide exception can only be enabled if the processor is configured with a hardware divider (`C_USE_DIV=1`). When configured with a hardware floating point unit (`C_USE_FPU>0`), it can also trap the following floating point specific exceptions: underflow, overflow, float division-by-zero, invalid operation, and denormalized operand error.

When configured with a hardware Memory Management Unit, it can also trap the following memory management specific exceptions: Illegal Instruction Exception, Data Storage Exception, Instruction Storage Exception, Data TLB Miss Exception, and Instruction TLB Miss Exception.

A hardware exception causes MicroBlaze to flush the pipeline and branch to the hardware exception vector (address 0x20). The execution stage instruction in the exception cycle is not executed.

The exception also updates the general purpose register R17 in the following manner:

- For the MMU exceptions (Data Storage Exception, Instruction Storage Exception, Data TLB Miss Exception, Instruction TLB Miss Exception) the register R17 is loaded with the appropriate program counter value to re-execute the instruction causing the exception upon return. The value is adjusted to return to a preceding IMM instruction, if any. If the exception is caused by an instruction in a branch delay slot, the value is adjusted to return to the branch instruction, including adjustment for a preceding IMM instruction, if any.
- For all other exceptions the register R17 is loaded with the program counter value of the subsequent instruction, unless the exception is caused by an instruction in a branch delay slot. If the exception is caused by an instruction in a branch delay slot, the ESR[DS] bit is set. In this case the exception handler should resume execution from the branch target address stored in BTR.

The EE and EIP bits in MSR are automatically reverted when executing the RTED instruction.

The VM and UM bits in MSR are automatically reverted from VMS and UMS when executing the RTED, RTBD, and RTID instructions.

Exception Priority

When two or more exceptions occur simultaneously, they are handled in the following order, from the highest priority to the lowest:

- Instruction Bus Exception
- Instruction TLB Miss Exception
- Instruction Storage Exception
- Illegal Opcode Exception
- Privileged Instruction Exception
- Data TLB Miss Exception
- Data Storage Exception
- Unaligned Exception
- Data Bus Exception
- Divide Exception
- FPU Exception
- Fast Simplex Link Exception

Exception Causes

- Fast Simplex Link Exception

The Fast Simplex Link (FSL) exception is caused by executing a `get` or `getd` instruction with the ‘e’ bit set to ‘1’ when there is a control bit mismatch.

- Instruction Bus Exception

The instruction Processor Local Bus (PLB) exception is caused by an active error signal from the slave (`IPLB_MRdErr`) or timeout signal from the arbiter (`IPLB_MTimeout`). The instructions side local memory (ILMB) and CacheLink (IXCL) interfaces cannot cause instruction bus exceptions.

- Illegal Opcode Exception

The illegal opcode exception is caused by an instruction with an invalid major opcode (bits 0 through 5 of instruction). Bits 6 through 31 of the instruction are not checked. Optional processor instructions are detected as illegal if not enabled. If the optional feature `C_OPCODE_0x0_ILLEGAL` is enabled, an illegal opcode exception is also caused if the instruction is equal to `0x00000000`.

- Data Bus Exception

The data Processor Local Bus exception is caused by an active error signal from the slave (`DPLB_MRdErr` or `DPLB_MWrErr`) or timeout signal from the arbiter (`DPLB_MTimeout`). The data side local memory (DLMB) and CacheLink (DXCL) interfaces can not cause data bus exceptions.

- Unaligned Exception

The unaligned exception is caused by a word access where the address to the data bus has bits 30 or 31 set, or a half-word access with bit 31 set.

- Divide Exception

The divide exception is caused by an integer division (`idiv` or `idivu`) where the divisor is zero, or by a signed integer division (`idiv`) where overflow occurs ($-2147483648 / -1$).

- FPU Exception

An FPU exception is caused by an underflow, overflow, divide-by-zero, illegal operation, or denormalized operand occurring with a floating point instruction.

- ◆ Underflow occurs when the result is denormalized.
- ◆ Overflow occurs when the result is not-a-number (NaN).
- ◆ The divide-by-zero FPU exception is caused by the rA operand to `fdiv` being zero when rB is not infinite.
- ◆ Illegal operation is caused by a signaling NaN operand or by illegal infinite or zero operand combinations.

- Privileged Instruction Exception

The Privileged Instruction exception is caused by an attempt to execute a privileged instruction in User Mode.

- Data Storage Exception

The Data Storage exception is caused by an attempt to access data in memory that results in a memory-protection violation.

- Instruction Storage Exception

The Instruction Storage exception is caused by an attempt to access instructions in memory that results in a memory-protection violation.

- Data TLB Miss Exception

The Data TLB Miss exception is caused by an attempt to access data in memory, when a valid Translation Look-Aside Buffer entry is not present, and virtual protected mode is enabled.

- Instruction TLB Miss Exception

The Instruction TLB Miss exception is caused by an attempt to access instructions in memory, when a valid Translation Look-Aside Buffer entry is not present, and virtual protected mode is enabled.

Equivalent Pseudocode

```

ESR[DS] ← exception in delay slot
if ESR[DS] then
    BTR ← branch target PC
    if MMU exception then
        if branch preceded by IMM then
            r17 ← PC - 8
        else
            r17 ← PC - 4
        else
            r17 ← invalid value
    else if MMU exception then
        if instruction preceded by IMM then
            r17 ← PC - 4
        else
            r17 ← PC
    else
        r17 ← PC + 4
PC ← 0x00000020
MSR[EE] ← 0, MSR[EIP]← 1
MSR[UMS] ← MSR[UM] , MSR[UM] ← 0, MSR[VMS] ← MSR[VM] , MSR[VM] ← 0
ESR[EC] ← exception specific value
ESR[ESS] ← exception specific value
EAR ← exception specific value
FSR ← exception specific value
Reservation ← 0

```

Breaks

There are two kinds of breaks:

- Hardware (external) breaks
- Software (internal) breaks

Hardware Breaks

Hardware breaks are performed by asserting the external break signal (that is, the Ext_BRK and Ext_NM_BRK input ports). On a break, the instruction in the execution stage completes while the instruction in the decode stage is replaced by a branch to the break vector (address 0x18). The break return address (the PC associated with the instruction in the decode stage at the time of the break) is automatically loaded into general purpose register R16. MicroBlaze also sets the Break In Progress (BIP) flag in the Machine Status Register (MSR).

A normal hardware break (that is, the Ext_BRK input port) is only handled when MSR[BIP] and MSR[EIP] are set to 0 (that is, there is no break or exception in progress). The Break In Progress flag disables interrupts. A non-maskable break (that is, the Ext_NM_BRK input port) is always handled immediately.

The BIP bit in the MSR is automatically cleared when executing the RTBD instruction.

The Ext_BRK signal must be kept asserted until the break has occurred, and deasserted before the RTBD instruction is executed. The Ext_NM_BRK signal must only be asserted one clock cycle.

Software Breaks

To perform a software break, use the brk and brki instructions. Refer to [Chapter 4, “MicroBlaze Instruction Set Architecture”](#) for detailed information on software breaks.

Latency

The time it takes MicroBlaze to enter a break service routine from the time the break occurs depends on the instruction currently in the execution stage and the latency to the memory storing the break vector.

Equivalent Pseudocode

```
r16 ← PC  
PC ← 0x00000018  
MSR [BIP] ← 1  
MSR [UMS] ← MSR [UM] , MSR [UM] ← 0 , MSR [VMS] ← MSR [VM] , MSR [VM] ← 0  
Reservation ← 0
```

Interrupt

MicroBlaze supports one external interrupt source (connected to the `Interrupt` input port). The processor only reacts to interrupts if the Interrupt Enable (IE) bit in the Machine Status Register (MSR) is set to 1. On an interrupt, the instruction in the execution stage completes while the instruction in the decode stage is replaced by a branch to the interrupt vector (address 0x10). The interrupt return address (the PC associated with the instruction in the decode stage at the time of the interrupt) is automatically loaded into general purpose register R14. In addition, the processor also disables future interrupts by clearing the IE bit in the MSR. The IE bit is automatically set again when executing the RTID instruction.

Interrupts are ignored by the processor if either of the break in progress (BIP) or exception in progress (EIP) bits in the MSR are set to 1.

By using the parameter `C_INTERRUPT_IS_EDGE`, the external interrupt can either be set to level-sensitive or edge-sensitive:

- When using level-sensitive interrupts, the `Interrupt` input must remain set until MicroBlaze has taken the interrupt, and jumped to the interrupt vector. Software must clear the interrupt before returning from the interrupt handler. If not, the interrupt is taken again, as soon as interrupts are enabled when returning from the interrupt handler.
- When using edge-sensitive interrupts, MicroBlaze detects and latches the `Interrupt` input edge, which means that the input only needs to be asserted one clock cycle. The interrupt input can remain asserted, but must be deasserted at least one clock cycle before a new interrupt can be detected.

Latency

The time it takes MicroBlaze to enter an Interrupt Service Routine (ISR) from the time an interrupt occurs, depends on the configuration of the processor and the latency of the memory controller storing the interrupt vectors. If MicroBlaze is configured to have a hardware divider, the largest latency happens when an interrupt occurs during the execution of a division instruction.

Equivalent Pseudocode

```
r14 ← PC  
PC ← 0x00000010  
MSR [IE] ← 0  
MSR [UMS] ← MSR [UM] , MSR [UM] ← 0 , MSR [VMS] ← MSR [VM] , MSR [VM] ← 0  
Reservation ← 0
```

User Vector (Exception)

The user exception vector is located at address 0x8. A user exception is caused by inserting a ‘BRALID Rx, 0x8’ instruction in the software flow. Although Rx could be any general purpose register, Xilinx recommends using R15 for storing the user exception return address, and to use the RTSD instruction to return from the user exception handler.

Pseudocode

```
rx ← PC  
PC ← 0x00000008  
MSR [UMS] ← MSR [UM] , MSR [UM] ← 0 , MSR [VMS] ← MSR [VM] , MSR [VM] ← 0  
Reservation ← 0
```

Instruction Cache

Overview

MicroBlaze can be used with an optional instruction cache for improved performance when executing code that resides outside the LMB address range.

The instruction cache has the following features:

- Direct mapped (1-way associative)
- User selectable cacheable memory address range
- Configurable cache and tag size
- Caching over CacheLink (XCL) interface
- Option to use 4 or 8 word cache-line
- Cache on and off controlled using a bit in the MSR
- Optional WIC instruction to invalidate instruction cache lines
- Optional stream buffers to improve performance by speculatively prefetching instructions
- Optional victim cache to improve performance by saving evicted cache lines

General Instruction Cache Functionality

When the instruction cache is used, the memory address space is split into two segments: a cacheable segment and a non-cacheable segment. The cacheable segment is determined by two parameters: `C_ICACHE_BASEADDR` and `C_ICACHE_HIGHADDR`. All addresses within this range correspond to the cacheable address segment. All other addresses are non-cacheable.

The cacheable segment size must be 2^N , where N is a positive integer. The range specified by `C_ICACHE_BASEADDR` and `C_ICACHE_HIGHADDR` must comprise a complete power-of-two range, such that range = 2^N and the N least significant bits of `C_ICACHE_BASEADDR` must be zero.

The cacheable instruction address consists of two parts: the cache address, and the tag address. The MicroBlaze instruction cache can be configured from 64 bytes to 64 kB. This corresponds to a cache address of between 6 and 16 bits. The tag address together with the cache address should match the full address of cacheable memory. When selecting cache sizes below 2 kB, distributed RAM is used to implement the Tag RAM and Instruction RAM.

For example: in a MicroBlaze configured with `C_ICACHE_BASEADDR= 0x00300000`, `C_ICACHE_HIGHADDR=0x0030ffff`, `C_CACHE_BYTE_SIZE=4096`, and `C_ICACHE_LINE_LEN=8`; the cacheable memory of 64 kB uses 16 bits of byte address, and the 4 kB cache uses 12 bits of byte address, thus the required address tag width is: 16-12=4 bits. The total number of block RAM primitives required in this configuration is: 2 RAMB16 for storing the 1024 instruction words, and 1 RAMB16 for 128 cache line entries, each consisting of: 4 bits of tag, 8 word-valid bits, 1 line-valid bit. In total 3 RAMB16 primitives.

Figure 1-20, page 72 shows the organization of Instruction Cache.

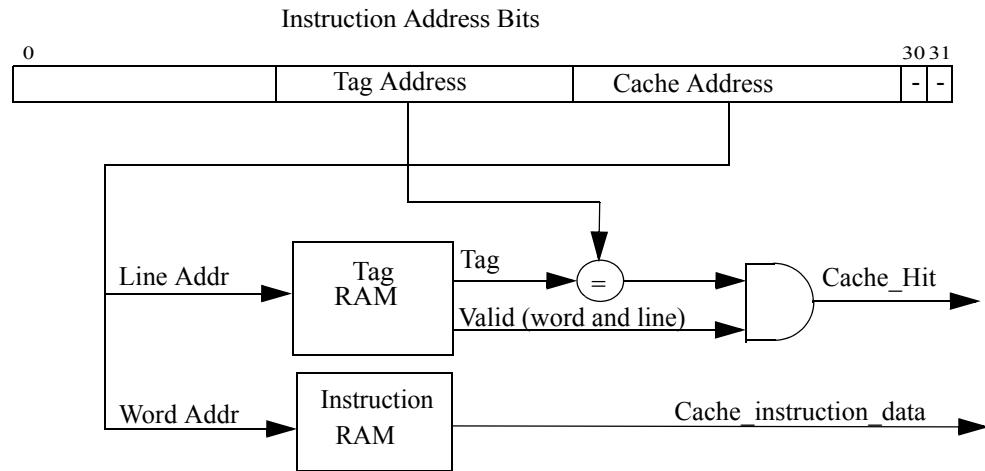


Figure 1-20: Instruction Cache Organization

Instruction Cache Operation

For every instruction fetched, the instruction cache detects if the instruction address belongs to the cacheable segment. If the address is non-cacheable, the cache controller ignores the instruction and lets the PLB or LMB complete the request. If the address is cacheable, a lookup is performed on the tag memory to check if the requested address is currently cached. The lookup is successful if: the word and line valid bits are set, and the tag address matches the instruction address tag segment. On a cache miss, the cache controller requests the new instruction over the instruction CacheLink (IXCL) interface, and waits for the memory controller to return the associated cache line.

Stream Buffers

When stream buffers are enabled, by setting the parameter `C_ICACHE_STREAMS` to 1, the cache will speculatively fetch cache lines in advance in sequence following the last requested address, until the stream buffer is full. The stream buffer can hold up to two cache lines. Should the processor subsequently request instructions from a cache line prefetched by the stream buffer, which occurs in linear code, they are immediately available.

The stream buffer often improves performance, since the processor generally has to spend less time waiting for instructions to be fetched from memory.

To be able to use instruction cache stream buffers, area optimization must not be enabled.

Victim Cache

The victim cache is enabled by setting the parameter `C_ICACHE_VICTIMS` to 2, 4 or 8. This defines the number of cache lines that can be stored in the victim cache. Whenever a cache line is evicted from the cache, it is saved in the victim cache. By saving the most recent lines they can be fetched much faster, should the processor request them, thereby improving performance. If the victim cache is not used, all evicted cache lines must be read from memory again when they are needed.

Note that to be able to use the victim cache, area optimization must not be enabled.

Instruction Cache Software Support

MSR Bit

The ICE bit in the MSR provides software control to enable and disable caches.

The contents of the cache are preserved by default when the cache is disabled. You can invalidate cache lines using the WIC instruction or using the hardware debug logic of MicroBlaze.

WIC Instruction

The optional WIC instruction (`C_ALLOW_ICACHE_WR=1`) is used to invalidate cache lines in the instruction cache from an application. For a detailed description, refer to [Chapter 4, “MicroBlaze Instruction Set Architecture”](#).

Data Cache

Overview

MicroBlaze can be used with an optional data cache for improved performance. The cached memory range must not include addresses in the LMB address range.

The data cache has the following features

- Direct mapped (1-way associative)
- Write-through or write-back
- User selectable cacheable memory address range
- Configurable cache size and tag size
- Caching over CacheLink (XCL) interface
- Option to use 4 or 8 word cache-lines
- Cache on and off controlled using a bit in the MSR
- Optional WDC instruction to invalidate or flush data cache lines
- Optional victim cache with write-back to improve performance by saving evicted cache lines

General Data Cache Functionality

When the data cache is used, the memory address space is split into two segments: a cacheable segment and a non-cacheable segment. The cacheable area is determined by two parameters: `C_DCACHE_BASEADDR` and `C_DCACHE_HIGHADDR`. All addresses within this range correspond to the cacheable address space. All other addresses are non-cacheable.

The cacheable segment size must be 2^N , where N is a positive integer. The range specified by `C_DCACHE_BASEADDR` and `C_DCACHE_HIGHADDR` must comprise a complete power-of-two range, such that range = 2^N and the N least significant bits of `C_DCACHE_BASEADDR` must be zero.

[Figure 1-21](#) shows the Data Cache Organization.

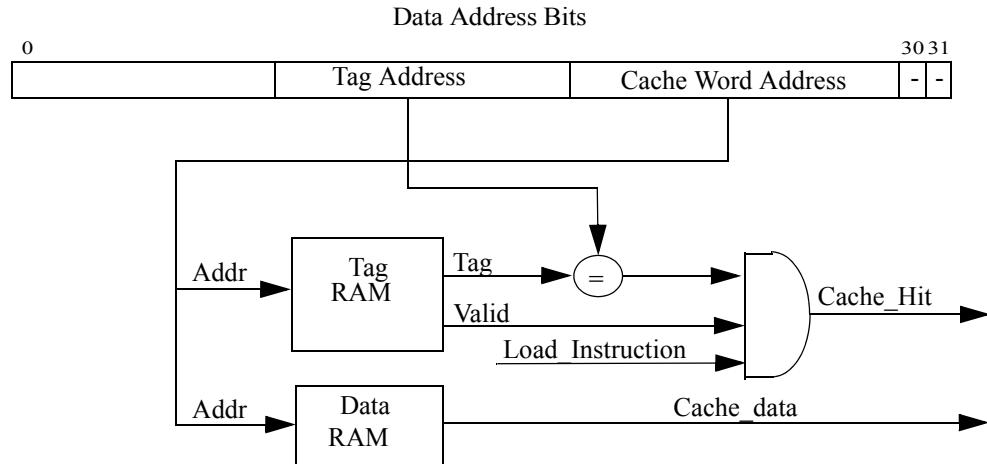


Figure 1-21: Data Cache Organization

The cacheable data address consists of two parts: the cache address, and the tag address. The MicroBlaze data cache can be configured from 64 bytes to 64 kB. This corresponds to a cache address of between 6 and 16 bits. The tag address together with the cache address should match the full address of cacheable memory. When selecting cache sizes below 2 kB, distributed RAM is used to implement the Tag RAM and Data RAM, except that block RAM is always used for the Data RAM when `C_AREA_OPTIMIZED` is set and `C_DCACHE_USE_WRITEBACK` is not set.

For example, in a MicroBlaze configured with `C_DCACHE_BASEADDR=0x00400000`, `C_DCACHE_HIGHADDR=0x00403fff`, `C_DCACHE_BYTE_SIZE=2048`, and `C_DCACHE_LINE_LEN=4`; the cacheable memory of 16 kB uses 14 bits of byte address, and the 2 kB cache uses 11 bits of byte address, thus the required address tag width is $14-11=3$ bits. The total number of block RAM primitives required in this configuration is 1 RAMB16 for storing the 512 data words, and 1 RAMB16 for 128 cache line entries, each consisting of 3 bits of tag, 4 word-valid bits, 1 line-valid bit. In total, 2 RAMB16 primitives.

Data Cache Operation

The caching policy used by the MicroBlaze data cache, write-back or write-through, is determined by the parameter `C_DCACHE_USE_WRITEBACK`. When this parameter is set, a write-back protocol is implemented, otherwise write-through is implemented. However, when configured with an MMU (`C_USE_MMU > 1`, `C_AREA_OPTIMIZED = 0`, `C_DCACHE_USE_WRITEBACK = 1`), the caching policy in virtual mode is determined by the W storage attribute in the TLB entry, whereas write-back is used in real mode.

With the write-back protocol, a store to an address within the cacheable range always updates the cached data. If the target address word is not in the cache (that is, the access is a cache-miss), and the location in the cache contains data that has not yet been written to memory (the cache location is dirty), the old data is written over the data CacheLink (DXCL) to external memory before updating the cache with the new data. If an entire cache line needs to be written, a burst cache line write is used, otherwise single word writes are used. For byte or halfword stores, in case of a cache-miss, the address is first requested over the data CacheLink, while a word store only updates the cache.

With the write-through protocol, a store to an address within the cacheable range generates an equivalent byte, halfword, or word write over the data CacheLink to external memory. The write

also updates the cached data if the target address word is in the cache (that is, the write is a cache-hit). A write cache-miss does not load the associated cache line into the cache.

Provided that the cache is enabled a load from an address within the cacheable range triggers a check to determine if the requested data is currently cached. If it is (that is, on a cache-hit) the requested data is retrieved from the cache. If not (that is, on a cache-miss) the address is requested over the data CacheLink, and the processor pipeline stalls until the cache line associated to the requested address is returned from the external memory controller.

Victim Cache

The victim cache is enabled by setting the parameter `C_DCACHE_VICTIMS` to 2, 4 or 8. This defines the number of cache lines that can be stored in the victim cache. Whenever a complete cache line is evicted from the cache, it is saved in the victim cache. By saving the most recent lines they can be fetched much faster, should the processor request them, thereby improving performance. If the victim cache is not used, all evicted cache lines must be read from memory again when they are needed.

Note that to be able to use the victim cache, write-back must be enabled and area optimization must not be enabled.

Data Cache Software Support

MSR Bit

The DCE bit in the MSR controls whether or not the cache is enabled. When disabling caches the user must ensure that all the prior writes within the cacheable range have been completed in external memory before reading back over PLB. This can be done by writing to a semaphore immediately before turning off caches, and then in a loop poll until it has been written.

The contents of the cache are preserved when the cache is disabled.

WDC Instruction

The optional WDC instruction (`C_ALLOW_DCACHE_WR=1`) is used to invalidate or flush cache lines in the data cache from an application. For a detailed description, please refer to [Chapter 4, “MicroBlaze Instruction Set Architecture”](#).

Floating Point Unit (FPU)

Overview

The MicroBlaze floating point unit is based on the [IEEE 754 standard](#):

- Uses IEEE 754 single precision floating point format, including definitions for infinity, not-a-number (NaN), and zero
- Supports addition, subtraction, multiplication, division, comparison, conversion and square root instructions
- Implements round-to-nearest mode
- Generates sticky status bits for: underflow, overflow, divide-by-zero and invalid operation

For improved performance, the following non-standard simplifications are made:

- Denormalized⁽¹⁾ operands are not supported. A hardware floating point operation on a denormalized number returns a quiet NaN and sets the sticky denormalized operand error bit in FSR; see "Floating Point Status Register (FSR)" on page 36
- A denormalized result is stored as a signed 0 with the underflow bit set in FSR. This method is commonly referred to as Flush-to-Zero (FTZ)
- An operation on a quiet NaN returns the fixed NaN: 0xFFC00000, rather than one of the NaN operands
- Overflow as a result of a floating point operation always returns signed ∞

Format

An IEEE 754 single precision floating point number is composed of the following three fields:

1. 1-bit *sign*
2. 8-bit biased *exponent*
3. 23-bit *fraction* (a.k.a. mantissa or significand)

The fields are stored in a 32 bit word as defined in [Figure 1-22](#):

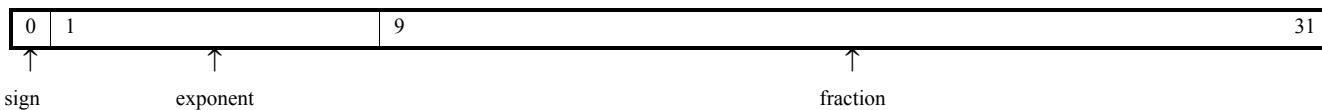


Figure 1-22: IEEE 754 Single Precision Format

The value of a floating point number v in MicroBlaze has the following interpretation:

1. If *exponent* = 255 and *fraction* \neq 0, then $v = \text{NaN}$, regardless of the *sign* bit
2. If *exponent* = 255 and *fraction* = 0, then $v = (-1)^{\text{sign}} * \infty$
3. If $0 < \text{exponent} < 255$, then $v = (-1)^{\text{sign}} * 2^{(\text{exponent}-127)} * (1.\text{fraction})$
4. If *exponent* = 0 and *fraction* \neq 0, then $v = (-1)^{\text{sign}} * 2^{-126} * (0.\text{fraction})$
5. If *exponent* = 0 and *fraction* = 0, then $v = (-1)^{\text{sign}} * 0$

1. Numbers that are so close to 0, that they cannot be represented with full precision, that is, any number n that falls in the following ranges: ($1.17549 \times 10^{-38} > n > 0$), or ($0 > n > -1.17549 \times 10^{-38}$)

For practical purposes only 3 and 5 are useful, while the others all represent either an error or numbers that can no longer be represented with full precision in a 32 bit format.

Rounding

The MicroBlaze FPU only implements the default rounding mode, “Round-to-nearest”, specified in IEEE 754. By definition, the result of any floating point operation should return the nearest single precision value to the infinitely precise result. If the two nearest representable values are equally near, then the one with its least significant bit zero is returned.

Operations

All MicroBlaze FPU operations use the processor's general purpose registers rather than a dedicated floating point register file, see “[General Purpose Registers](#)”.

Arithmetic

The FPU implements the following floating point operations:

- addition, fadd
- subtraction, fsub
- multiplication, fmul
- division, fdiv
- square root, fsqrt (available if C_USE_FPU = 2)

Comparison

The FPU implements the following floating point comparisons:

- compare less-than, fcmp.lt
- compare equal, fcmp.eq
- compare less-or-equal, fcmp.le
- compare greater-than, fcmp.gt
- compare not-equal, fcmp.ne
- compare greater-or-equal, fcmp.ge
- compare unordered, fcmp.un (used for NaN)

Conversion

The FPU implements the following conversions (available if C_USE_FPU = 2):

- convert from signed integer to floating point, flt
- convert from floating point to signed integer, fint

Exceptions

The floating point unit uses the regular hardware exception mechanism in MicroBlaze. When enabled, exceptions are thrown for all the IEEE standard conditions: underflow, overflow, divide-by-zero, and illegal operation, as well as for the MicroBlaze specific exception: denormalized operand error.

A floating point exception inhibits the write to the destination register (Rd). This allows a floating point exception handler to operate on the uncorrupted register file.

Software Support

The EDK compiler system, based on GCC, provides support for the Floating Point Unit compliant with the MicroBlaze API. Compiler flags are automatically added to the GCC command line based on the type of FPU present in the system, when using XPS or SDK.

All double-precision operations are emulated in software. Be aware that the `xil_printf()` function does not support floating-point output. The standard C library `printf()` and related functions do support floating-point output, but will increase the program code size.

Libraries and Binary Compatibility

The EDK compiler system only includes software floating point C runtime libraries. To take advantage of the hardware FPU, the libraries must be recompiled with the appropriate compiler switches.

For all cases where separate compilation is used, it is very important that you ensure the consistency of FPU compiler flags throughout the build.

Operator Latencies

The latencies of the various operations supported by the FPU are listed in [Chapter 4, “MicroBlaze Instruction Set Architecture.”](#) The FPU instructions are not pipelined, so only one operation can be ongoing at any time.

C Language Programming

To gain maximum benefit from the FPU without low-level assembly-language programming, it is important to consider how the C compiler will interpret your source code. Very often the same algorithm can be expressed in many different ways, and some are more efficient than others.

Immediate Constants

Floating-point constants in C are double-precision by default. When using a single-precision FPU, careless coding may result in double-precision software emulation routines being used instead of the native single-precision instructions. To avoid this, explicitly specify (by cast or suffix) that immediate constants in your arithmetic expressions are single-precision values.

For example:

```
float x = 0.0;  
...  
x += (float)1.0; /* float addition */  
x += 1.0F;        /* alternative to above */  
x += 1.0;         /* warning - uses double addition! */
```

Note that the GNU C compiler can be instructed to treat all floating-point constants as single-precision (contrary to the ANSI C standard) by supplying the compiler flag `-fsingle-precision-constants`.

Avoid unnecessary casting

While conversions between floating-point and integer formats are supported in hardware by the FPU, when `C_USE_FPU` is set to 2, it is still best to avoid them when possible.

The following “bad” example calculates the sum of squares of the integers from 1 to 10 using floating-point representation:

```
float sum, t;
int i;
sum = 0.0f;
for (i = 1; i <= 10; i++) {
    t = (float)i;
    sum += t * t;
}
```

The above code requires a cast from an integer to a float on each loop iteration. This can be rewritten as:

```
float sum, t;
int i;
t = sum = 0.0f;
for(i = 1; i <= 10; i++) {
    t += 1.0f;
    sum += t * t;
}
```

Note that the compiler is not at liberty to perform this optimization in general, as the two code fragments above may give different results in some cases (for example, very large t).

Square root runtime library function

The standard C runtime math library functions operate using double-precision arithmetic. When using a single-precision FPU, calls to the square root functions (sqrt()) result in inefficient emulation routines being used instead of FPU instructions:

```
#include <math.h>
...
float x=-1.0F;
...
x = sqrt(x); /* uses double precision */
```

Here the math.h header is included to avoid a warning message from the compiler.

When used with single-precision data types, the result is a cast to double, a runtime library call is made (which does not use the FPU) and then a truncation back to float is performed.

The solution is to use the non-ANSI function sqrtf() instead, which operates using single precision and can be carried out using the FPU. For example:

```
#include <math.h>
...
float x=-1.0F;
...
x = sqrtf(x); /* uses single precision */
```

Note that when compiling this code, the compiler flag -fno-math-errno (in addition to -mhard-float and -mxl-float-sqrt) must be used, to ensure that the compiler does not generate unnecessary code to handle error conditions by updating the errno variable.

Fast Simplex Link (FSL)

MicroBlaze can be configured with up to 16 Fast Simplex Link (FSL) interfaces, each consisting of one input and one output port. The FSL channels are dedicated uni-directional point-to-point data streaming interfaces. For detailed information on the FSL interface, please refer to the *Fast Simplex Link (FSL) Bus* data-sheet, DS449, in the Xilinx EDK IP Documentation.

The FSL interfaces on MicroBlaze are 32 bits wide. A separate bit indicates whether the sent/received word is of control or data type. The get instruction in the MicroBlaze ISA is used to transfer information from an FSL port to a general purpose register. The put instruction is used to transfer data in the opposite direction. Both instructions come in 4 flavors: blocking data, non-blocking data, blocking control, and non-blocking control. For a detailed description of the get and put instructions, please refer to [Chapter 4, “MicroBlaze Instruction Set Architecture”](#).

Hardware Acceleration using FSL

Each FSL provides a low latency dedicated interface to the processor pipeline. Thus they are ideal for extending the processors execution unit with custom hardware accelerators. A simple example is illustrated in [Figure 1-23](#).

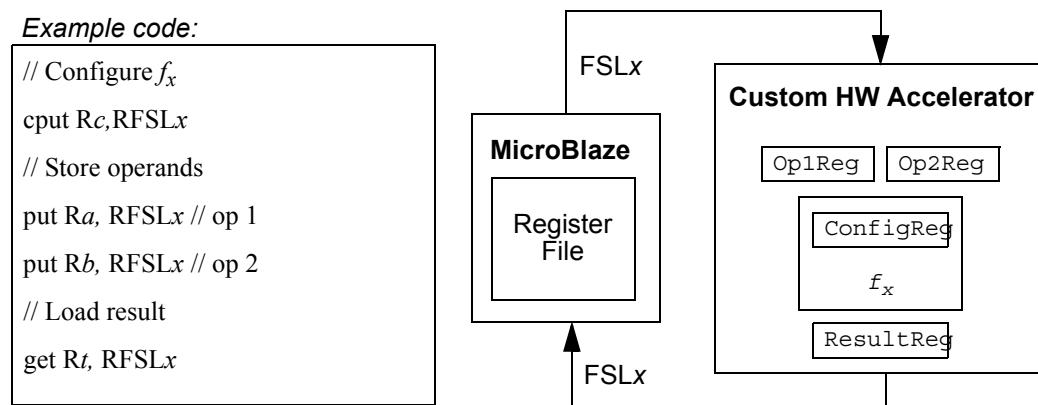


Figure 1-23: FSL Used with HW Accelerated Function f_x

This method is similar to extending the ISA with custom instructions, but has the benefit of not making the overall speed of the processor pipeline dependent on the custom function. Also, there are no additional requirements on the software tool chain associated with this type of functional extension.

Debug and Trace

Debug Overview

MicroBlaze features a debug interface to support JTAG based software debugging tools (commonly known as BDM or Background Debug Mode debuggers) like the Xilinx Microprocessor Debug (XMD) tool. The debug interface is designed to be connected to the Xilinx Microprocessor Debug Module (MDM) core, which interfaces with the JTAG port of Xilinx FPGAs. Multiple MicroBlaze instances can be interfaced with a single MDM to enable multiprocessor debugging. The debugging features include:

- Configurable number of hardware breakpoints and watchpoints and unlimited software breakpoints
- External processor control enables debug tools to stop, reset, and single step MicroBlaze
- Read from and write to: memory, general purpose registers, and special purpose register, except EAR, EDR, ESR, BTR and PVR0 - PVR11, which can only be read
- Support for multiple processors
- Write to instruction and data caches

Trace Overview

The MicroBlaze trace interface exports a number of internal state signals for performance monitoring and analysis. Xilinx recommends that users only use the trace interface through Xilinx developed analysis cores. This interface is not guaranteed to be backward compatible in future releases of MicroBlaze.

MicroBlaze Signal Interface Description

This chapter describes the types of signal interfaces that can be used to connect MicroBlaze™. This chapter contains the following sections:

- “Overview”
- “MicroBlaze I/O Overview”
- “Processor Local Bus (PLB) Interface Description”
- “Local Memory Bus (LMB) Interface Description”
- “Fast Simplex Link (FSL) Interface Description”
- “Xilinx CacheLink (XCL) Interface Description”
- “Debug Interface Description”
- “Trace Interface Description”
- “MicroBlaze Core Configurability”

Overview

The MicroBlaze core is organized as a Harvard architecture with separate bus interface units for data and instruction accesses. The following three memory interfaces are supported: Local Memory Bus (LMB), the IBM Processor Local Bus (PLB), and Xilinx® CacheLink (XCL). The LMB provides single-cycle access to on-chip dual-port block RAM. The PLB interfaces provide a connection to both on-chip and off-chip peripherals and memory. The CacheLink interface is intended for use with specialized external memory controllers. MicroBlaze also supports up to 16 Fast Simplex Link (FSL) ports, each with one master and one slave FSL interface.

Features

MicroBlaze can be configured with the following bus interfaces:

- A 32-bit version of the PLB V4.6 interface (see IBM’s *128-Bit Processor Local Bus Architectural Specifications, Version 4.6*).
- LMB provides simple synchronous protocol for efficient block RAM transfers
- FSL provides a fast non-arbitrated streaming communication mechanism
- XCL provides a fast slave-side arbitrated streaming interface between caches and external memory controllers
- Debug interface for use with the Microprocessor Debug Module (MDM) core
- Trace interface for performance analysis

MicroBlaze I/O Overview

The core interfaces shown in [Figure 2-1](#) and the following [Table 2-1](#) are defined as follows:

- DPLB:** Data interface, Processor Local Bus
- DLMB:** Data interface, Local Memory Bus (BRAM only)
- IPLB:** Instruction interface, Processor Local Bus
- ILMB:** Instruction interface, Local Memory Bus (BRAM only)
- MFSL 0..15:** FSL master interfaces
- DWFSL 0..15:** FSL master direct connection interfaces
- SFSL 0..15:** FSL slave interfaces
- DRFSL 0..15:** FSL slave direct connection interfaces
- DXCL:** Data side Xilinx CacheLink interface (FSL master/slave pair)
- IXCL:** Instruction side Xilinx CacheLink interface (FSL master/slave pair)
- Core:** Miscellaneous signals for: clock, reset, debug, and trace

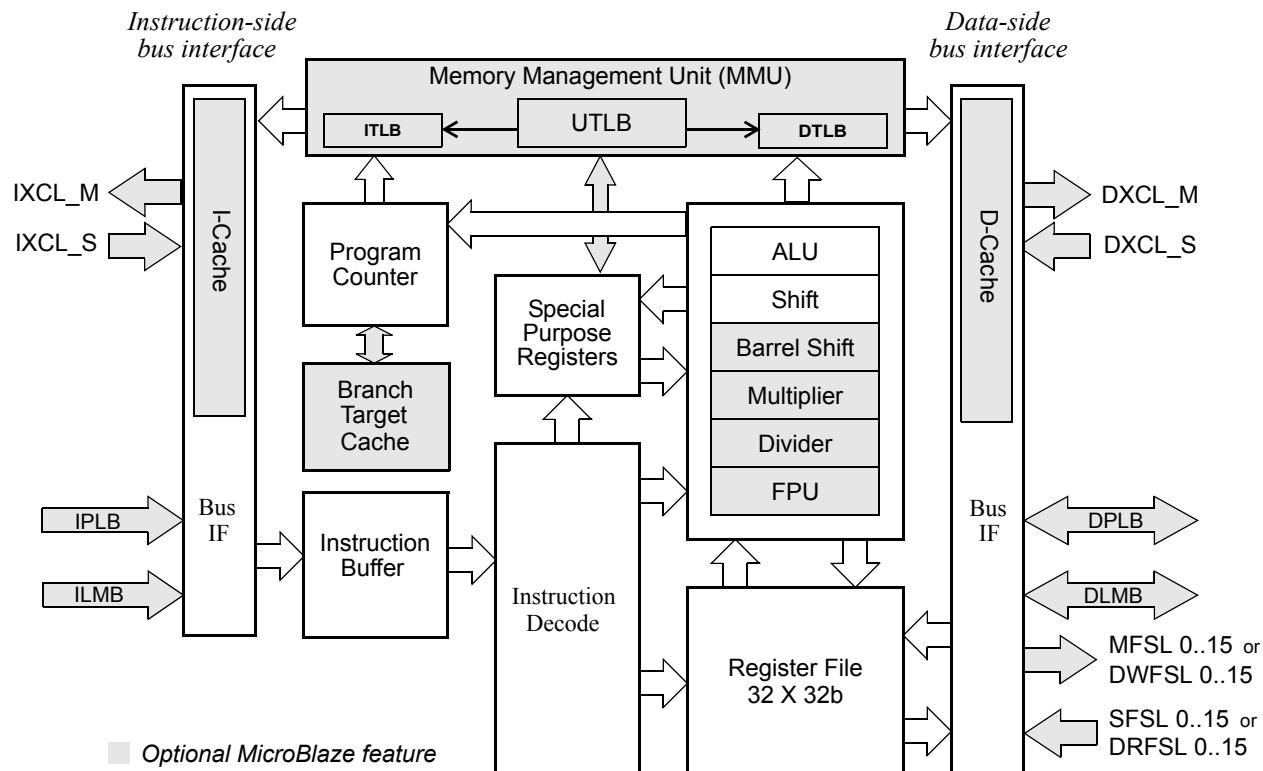


Figure 2-1: MicroBlaze Core Block Diagram

Table 2-1: Summary of MicroBlaze Core I/O

| Signal | Interface | I/O | Description |
|-------------------|------------------|------------|--|
| DPLB_M_AAbort | DPLB | O | Data Interface PLB abort bus request indicator |
| DPLB_M_ABus | DPLB | O | Data Interface PLB address bus |
| DPLB_M_UABus | DPLB | O | Data Interface PLB upper address bus |
| DPLB_M_BE | DPLB | O | Data Interface PLB byte enables |
| DPLB_M_busLock | DPLB | O | Data Interface PLB bus lock |
| DPLB_M_lockErr | DPLB | O | Data Interface PLB lock error indicator |
| DPLB_M_MSize | DPLB | O | Data Interface PLB master data bus size |
| DPLB_M_priority | DPLB | O | Data Interface PLB bus request priority |
| DPLB_M_rdBurst | DPLB | O | Data Interface PLB burst read transfer indicator |
| DPLB_M_request | DPLB | O | Data Interface PLB bus request |
| DPLB_M_RNW | DPLB | O | Data Interface PLB read/not write |
| DPLB_M_size | DPLB | O | Data Interface PLB transfer size |
| DPLB_M_TAttribute | DPLB | O | Data Interface PLB Transfer Attribute bus |
| DPLB_M_type | DPLB | O | Data Interface PLB transfer type |
| DPLB_M_wrBurst | DPLB | O | Data Interface PLB burst write transfer indicator |
| DPLB_M_wrDBus | DPLB | O | Data Interface PLB write data bus |
| DPLB_MBusy | DPLB | I | Data Interface PLB slave busy indicator |
| DPLB_MRdErr | DPLB | I | Data Interface PLB slave read error indicator |
| DPLB_MWrErr | DPLB | I | Data Interface PLB slave write error indicator |
| DPLB_MIRQ | DPLB | I | Data Interface PLB slave interrupt indicator |
| DPLB_MWrBTerm | DPLB | I | Data Interface PLB terminate write burst indicator |
| DPLB_MWrDAck | DPLB | I | Data Interface PLB write data acknowledge |
| DPLB_MAddrAck | DPLB | I | Data Interface PLB address acknowledge |
| DPLB_MRdBTerm | DPLB | I | Data Interface PLB terminate read burst indicator |
| DPLB_MRdDAck | DPLB | I | Data Interface PLB read data acknowledge |
| DPLB_MRdDBus | DPLB | I | Data Interface PLB read data bus |
| DPLB_MRdWdAddr | DPLB | I | Data Interface PLB read word address |

Table 2-1: Summary of MicroBlaze Core I/O (Continued)

| Signal | Interface | I/O | Description |
|-------------------|------------------|------------|---|
| DPLB_MRearbitrate | DPLB | I | Data Interface PLB bus rearbitrate indicator |
| DPLB_MSSize | DPLB | I | Data Interface PLB slave data bus size |
| DPLB_MTImeout | DPLB | I | Data Interface PLB bus timeout |
| IPLB_M_ABort | IPLB | O | Instruction Interface PLB abort bus request indicator |
| IPLB_M_ABus | IPLB | O | Instruction Interface PLB address bus |
| IPLB_M_UABus | IPLB | O | Instruction Interface PLB upper address bus |
| IPLB_M_BE | IPLB | O | Instruction Interface PLB byte enables |
| IPLB_M_busLock | IPLB | O | Instruction Interface PLB bus lock |
| IPLB_M_lockErr | IPLB | O | Instruction Interface PLB lock error indicator |
| IPLB_M_MSize | IPLB | O | Instruction Interface PLB master data bus size |
| IPLB_M_priority | IPLB | O | Instruction Interface PLB bus request priority |
| IPLB_M_rdBurst | IPLB | O | Instruction Interface PLB burst read transfer indicator |
| IPLB_M_request | IPLB | O | Instruction Interface PLB bus request |
| IPLB_M_RNW | IPLB | O | Instruction Interface PLB read/not write |
| IPLB_M_size | IPLB | O | Instruction Interface PLB transfer size |
| IPLB_M_TAttribute | IPLB | O | Instruction Interface PLB Transfer Attribute bus |
| IPLB_M_type | IPLB | O | Instruction Interface PLB transfer type |
| IPLB_M_wrBurst | IPLB | O | Instruction Interface PLB burst write transfer indicator |
| IPLB_M_wrDBus | IPLB | O | Instruction Interface PLB write data bus |
| IPLB_MBusy | IPLB | I | Instruction Interface PLB slave busy indicator |
| IPLB_MRdErr | IPLB | I | Instruction Interface PLB slave read error indicator |
| IPLB_MWrErr | IPLB | I | Instruction Interface PLB slave write error indicator |
| IPLB_MIRQ | IPLB | I | Instruction Interface PLB slave interrupt indicator |
| IPLB_MWrBTerm | IPLB | I | Instruction Interface PLB terminate write burst indicator |

Table 2-1: Summary of MicroBlaze Core I/O (Continued)

| Signal | Interface | I/O | Description |
|--------------------|---------------------|-----|---|
| IPLB_MWrDAck | IPLB | I | Instruction Interface PLB write data acknowledge |
| IPLB_MAddrAck | IPLB | I | Instruction Interface PLB address acknowledge |
| IPLB_MRdBTerm | IPLB | I | Instruction Interface PLB terminate read burst indicator |
| IPLB_MRdDAck | IPLB | I | Instruction Interface PLB read data acknowledge |
| IPLB_MRdDBus | IPLB | I | Instruction Interface PLB read data bus |
| IPLB_MRdWdAddr | IPLB | I | Instruction Interface PLB read word address |
| IPLB_MRearbitrate | IPLB | I | Instruction Interface PLB bus rearbitrate indicator |
| IPLB_MSSize | IPLB | I | Instruction Interface PLB slave data bus size |
| IPLB_MTtimeout | IPLB | I | Instruction Interface PLB bus timeout |
| Data_Addr[0:31] | DLMB | O | Data interface LMB address bus |
| Byte_Enable[0:3] | DLMB | O | Data interface LMB byte enables |
| Data_Write[0:31] | DLMB | O | Data interface LMB write data bus |
| D_AS | DLMB | O | Data interface LMB address strobe |
| Read_Strobe | DLMB | O | Data interface LMB read strobe |
| Write_Strobe | DLMB | O | Data interface LMB write strobe |
| Data_Read[0:31] | DLMB | I | Data interface LMB read data bus |
| DReady | DLMB | I | Data interface LMB data ready |
| Instr_Addr[0:31] | ILMB | O | Instruction interface LMB address bus |
| I_AS | ILMB | O | Instruction interface LMB address strobe |
| IFetch | ILMB | O | Instruction interface LMB instruction fetch |
| Instr[0:31] | ILMB | I | Instruction interface LMB read data bus |
| IReady | ILMB | I | Instruction interface LMB data ready |
| FSL0_M ... FSL15_M | MFSL or DWFSL | O | Master interface to output FSL channels MFSL is used for FSL bus connections, whereas DWFSL is used for direct connections with FSL slaves |
| FSL0_S ... FSL15_S | SFSL or DRFSL | I | Slave interface to input FSL channels SFSL is used for FSL bus connections, whereas DRFSL is used for direct connections with FSL masters |

Table 2-1: Summary of MicroBlaze Core I/O (Continued)

| Signal | Interface | I/O | Description |
|-----------------------|------------------|------------|---|
| ICache_FSL_in... | IXCL_S | IO | Instruction side CacheLink FSL slave interface |
| ICache_FSL_out... | IXCL_M | IO | Instruction side CacheLink FSL master interface |
| DCache_FSL_in... | DXCL_S | IO | Data side CacheLink FSL slave interface |
| DCache_FSL_out... | DXCL_M | IO | Data side CacheLink FSL master interface |
| Interrupt | Core | I | Interrupt |
| Reset ¹ | Core | I | Core reset, active high. Should be held for at least 1 Clk clock cycle. |
| Mb_Reset ¹ | Core | I | Core reset, active high. Should be held for at least 1 Clk clock cycle. |
| Clk | Core | I | Clock ² |
| Ext_BRK | Core | I | Break signal from MDM |
| Ext_NM_BRK | Core | I | Non-maskable break signal from MDM |
| MB_Halted | Core | O | Pipeline is halted, either via the Debug Interface or by setting Dbg_Stop |
| Dbg_Stop | Core | I | Unconditionally force pipeline to halt as soon as possible. Rising-edge detected pulse that should be held for at least 1 Clk clock cycle. The signal only has any effect when C_DEBUG_ENABLED is set to 1. |
| Dbg_... | Core | IO | Debug signals from MDM. See Table 2-9 for details. |
| Trace_... | Core | O | Trace signals for real time HW analysis. See Table 2-10 for details. |

1. The Reset and Mb_Reset signals are functionally equivalent. Mb_Reset is intended for the PLB interfaces.
2. MicroBlaze is a synchronous design clocked with the Clk signal, except for hardware debug logic, which is clocked with the Dbg_Clk signal. If hardware debug logic is not used, there is no minimum frequency limit for Clk. However, if hardware debug logic is used, there are signals transferred between the two clock regions. In this case Clk must have a higher frequency than Dbg_Clk.

Processor Local Bus (PLB) Interface Description

The MicroBlaze PLB interfaces are implemented as byte-enable capable 32-bit masters. Please refer to the *IBM 128-Bit Processor Local Bus Architectural Specification (v4.6)* document for details.

Local Memory Bus (LMB) Interface Description

The LMB is a synchronous bus used primarily to access on-chip block RAM. It uses a minimum number of control signals and a simple protocol to ensure that local block RAM are accessed in a single clock cycle. LMB signals and definitions are shown in the following table. All LMB signals are active high.

LMB Signal Interface

Table 2-2: LMB Bus Signals

| Signal | Data Interface | Instruction Interface | Type | Description |
|---------------------|------------------|-----------------------|------|-------------------------|
| Addr [0 : 31] | Data_Addr[0:31] | Instr_Addr[0:31] | O | Address bus |
| Byte_Enable [0 : 3] | Byte_Enable[0:3] | <i>not used</i> | O | Byte enables |
| Data_Write [0 : 31] | Data_Write[0:31] | <i>not used</i> | O | Write data bus |
| AS | D_AS | I_AS | O | Address strobe |
| Read_Strobe | Read_Strobe | IFetch | O | Read in progress |
| Write_Strobe | Write_Strobe | <i>not used</i> | O | Write in progress |
| Data_Read [0 : 31] | Data_Read[0:31] | Instr[0:31] | I | Read data bus |
| Ready | DReady | IReady | I | Ready for next transfer |
| Clk | Clk | Clk | I | Bus clock |

Addr[0:31]

The address bus is an output from the core and indicates the memory address that is being accessed by the current transfer. It is valid only when AS is high. In multicycle accesses (accesses requiring more than one clock cycle to complete), Addr [0 : 31] is valid only in the first clock cycle of the transfer.

Byte_Enable[0:3]

The byte enable signals are outputs from the core and indicate which byte lanes of the data bus contain valid data. Byte_Enable[0:3] is valid only when AS is high. In multicycle accesses (accesses requiring more than one clock cycle to complete), Byte_Enable [0 : 3] is valid only in the first clock cycle of the transfer. Valid values for Byte_Enable [0 : 31] are shown in the following table:

Table 2-3: Valid Values for Byte_Enable[0:3]

| Byte_Enable[0:3] | Byte Lanes Used | | | | |
|------------------|-----------------|------------|-------------|-------------|---|
| | Data[0:7] | Data[8:15] | Data[16:23] | Data[24:31] | |
| 0000 | | | | | |
| 0001 | | | | | x |
| 0010 | | | x | | |
| 0100 | | x | | | |
| 1000 | x | | | | |
| 0011 | | | x | | x |
| 1100 | x | x | | | |
| 1111 | x | x | x | | x |

Data_Write[0:31]

The write data bus is an output from the core and contains the data that is written to memory. It is valid only when AS is high. Only the byte lanes specified by Byte_Enable [0 : 3] contain valid data.

AS

The address strobe is an output from the core and indicates the start of a transfer and qualifies the address bus and the byte enables. It is high only in the first clock cycle of the transfer, after which it goes low and remains low until the start of the next transfer.

Read_Strobe

The read strobe is an output from the core and indicates that a read transfer is in progress. This signal goes high in the first clock cycle of the transfer, and may remain high until the clock cycle after Ready is sampled high. If a new read transfer is directly started in the next clock cycle, then Read_Strobe remains high.

Write_Strobe

The write strobe is an output from the core and indicates that a write transfer is in progress. This signal goes high in the first clock cycle of the transfer, and may remain high until the clock cycle after Ready is sampled high. If a new write transfer is directly started in the next clock cycle, then Write_Strobe remains high.

Data_Read[0:31]

The read data bus is an input to the core and contains data read from memory. Data_Read [0 : 31] is valid on the rising edge of the clock when Ready is high.

Ready

The Ready signal is an input to the core and indicates completion of the current transfer and that the next transfer can begin in the following clock cycle. It is sampled on the rising edge of the clock. For reads, this signal indicates the Data_Read [0 : 31] bus is valid, and for writes it indicates that the Data_Write [0 : 31] bus has been written to local memory.

Clk

All operations on the LMB are synchronous to the MicroBlaze core clock.

LMB Transactions

The following diagrams provide examples of LMB bus operations.

Generic Write Operation

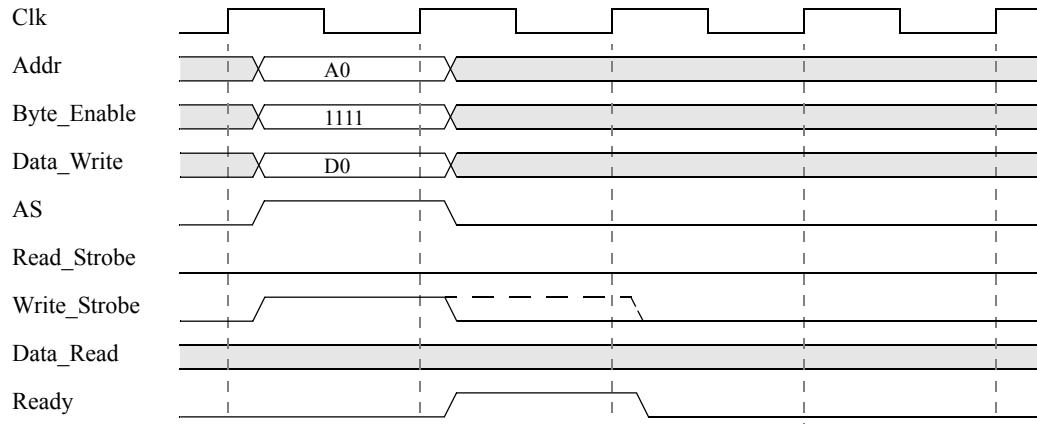


Figure 2-2: LMB Generic Write Operation

Generic Read Operation

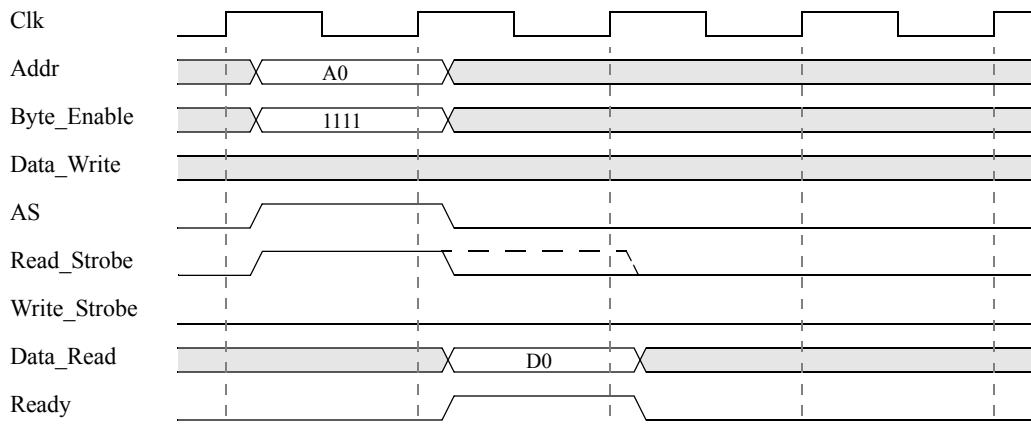


Figure 2-3: LMB Generic Read Operation

Back-to-Back Write Operation

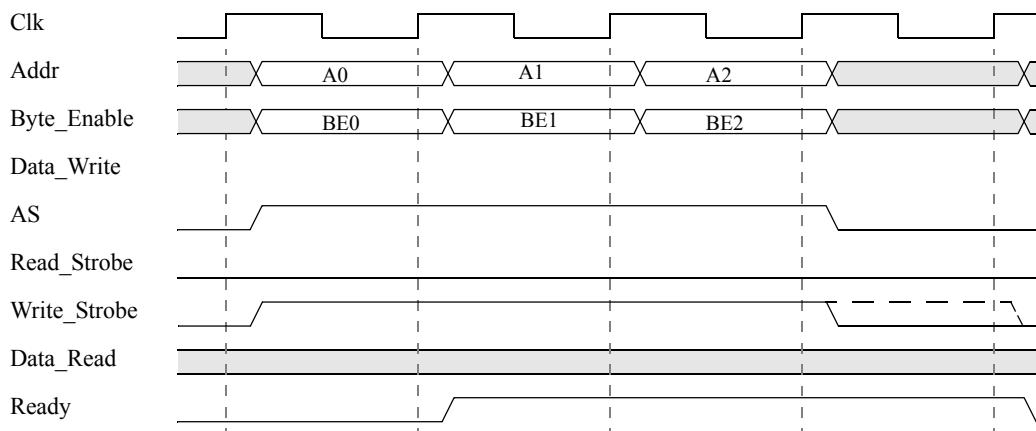


Figure 2-4: LMB Back-to-Back Write Operation

Single Cycle Back-to-Back Read Operation

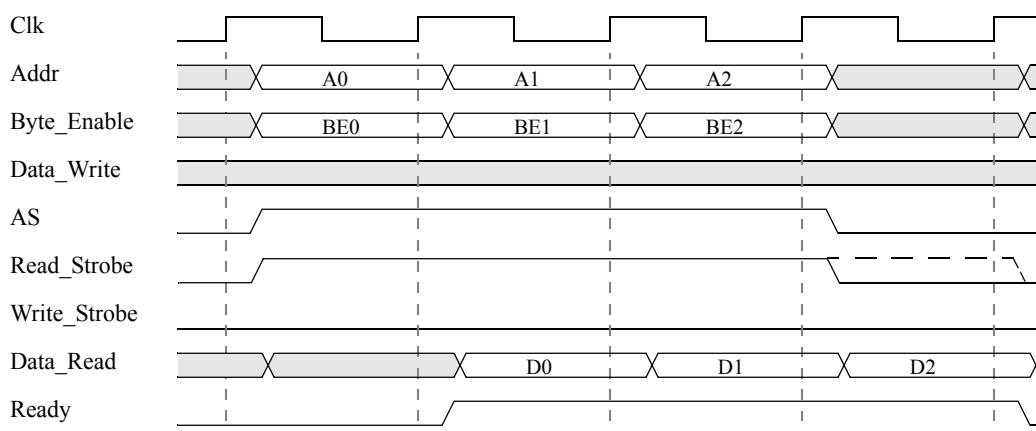


Figure 2-5: LMB Single Cycle Back-to-Back Read Operation

Back-to-Back Mixed Read/Write Operation

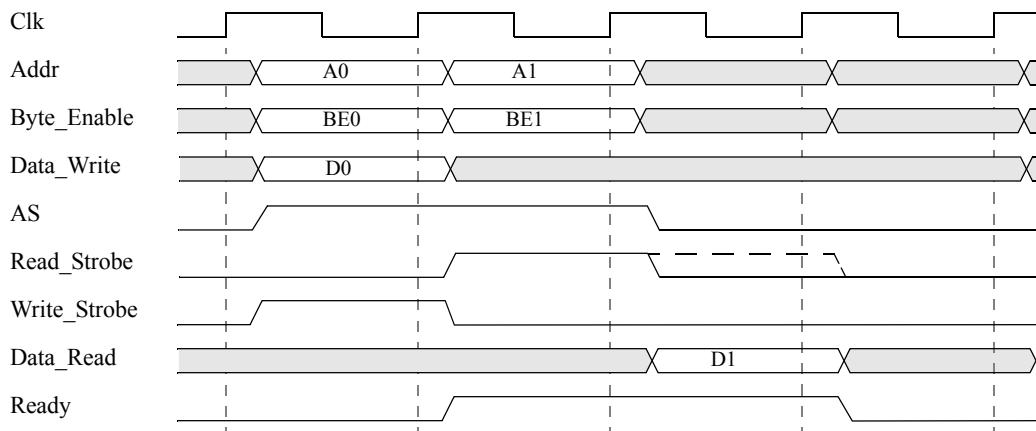


Figure 2-6: Back-to-Back Mixed Read/Write Operation

Read and Write Data Steering

The MicroBlaze data-side bus interface performs the read steering and write steering required to support the following transfers:

- byte, halfword, and word transfers to word devices
- byte and halfword transfers to halfword devices
- byte transfers to byte devices

MicroBlaze does not support transfers that are larger than the addressed device. These types of transfers require dynamic bus sizing and conversion cycles that are not supported by the MicroBlaze bus interface. Data steering for read cycles is shown in [Table 2-4](#), and data steering for write cycles is shown in [Table 2-5](#).

Table 2-4: Read Data Steering (Load to Register rD)

| Register rD Data | | | | | | |
|------------------|-------------------|---------------|---------|----------|-----------|-----------|
| Address [30:31] | Byte_Enable [0:3] | Transfer Size | rD[0:7] | rD[8:15] | rD[16:23] | rD[24:31] |
| 11 | 0001 | byte | | | | Byte3 |
| 10 | 0010 | byte | | | | Byte2 |
| 01 | 0100 | byte | | | | Byte1 |
| 00 | 1000 | byte | | | | Byte0 |
| 10 | 0011 | halfword | | | Byte2 | Byte3 |
| 00 | 1100 | halfword | | | Byte0 | Byte1 |
| 00 | 1111 | word | Byte0 | Byte1 | Byte2 | Byte3 |

Table 2-5: Write Data Steering (Store from Register rD)

| Write Data Bus Bytes | | | | | | |
|----------------------|-------------------|---------------|-----------|-----------|-----------|-----------|
| Address [30:31] | Byte_Enable [0:3] | Transfer Size | Byte0 | Byte1 | Byte2 | Byte3 |
| 11 | 0001 | byte | | | | rD[24:31] |
| 10 | 0010 | byte | | | rD[24:31] | |
| 01 | 0100 | byte | | rD[24:31] | | |
| 00 | 1000 | byte | rD[24:31] | | | |
| 10 | 0011 | halfword | | | rD[16:23] | rD[24:31] |
| 00 | 1100 | halfword | rD[16:23] | rD[24:31] | | |
| 00 | 1111 | word | rD[0:7] | rD[8:15] | rD[16:23] | rD[24:31] |

Note: Other masters may have more restrictive requirements for byte lane placement than those allowed by MicroBlaze. Slave devices are typically attached “left-justified” with byte devices attached to the most-significant byte lane, and halfword devices attached to the most significant halfword lane. The MicroBlaze steering logic fully supports this attachment method.

Fast Simplex Link (FSL) Interface Description

The Fast Simplex Link bus provides a point-to-point communication channel between an output FIFO and an input FIFO. For more information on the generic FSL protocol, see the *Fast Simplex Link (FSL) Bus* data-sheet, DS449, in the Xilinx EDK IP Documentation.

Master FSL Signal Interface

MicroBlaze may contain up to 16 master FSL interfaces. The master signals are depicted in [Table 2-6](#).

Table 2-6: Master FSL Signals

| Signal Name | Description | VHDL Type | Direction |
|----------------|---|------------------|-----------|
| FSLn_M_Clk | Clock | std_logic | input |
| FSLn_M_Write | Write enable signal indicating that data is being written to the output FSL | std_logic | output |
| FSLn_M_Data | Data value written to the output FSL | std_logic_vector | output |
| FSLn_M_Control | Control bit value written to the output FSL | std_logic | output |
| FSLn_M_Full | Full Bit indicating output FSL FIFO is full when set | std_logic | input |

Slave FSL Signal Interface

MicroBlaze may contain up to 16 slave FSL interfaces. The slave FSL interface signals are depicted in [Table 2-7](#).

Table 2-7: Slave FSL Signals

| Signal Name | Description | VHDL Type | Direction |
|----------------|---|------------------|-----------|
| FSLn_S_Clk | Clock | std_logic | input |
| FSLn_S_Read | Read acknowledge signal indicating that data has been read from the input FSL | std_logic | output |
| FSLn_S_Data | Data value currently available at the top of the input FSL | std_logic_vector | input |
| FSLn_S_Control | Control Bit value currently available at the top of the input FSL | std_logic | input |
| FSLn_S_Exists | Flag indicating that data exists in the input FSL | std_logic | input |

FSL Transactions

FSL BUS Write Operation

A write to the FSL bus is performed by MicroBlaze using one of the put or putd instructions. A write operation transfers the register contents to an output FSL bus. The transfer is completed in a single clock cycle for blocking mode writes to the FSL (put and cput instructions) as long as the FSL FIFO does not become full. If the FSL FIFO is full, the processor stalls until the FSL full flag is lowered. The non-blocking instructions (with prefix n), always complete in a single clock cycle even if the FSL was full. If the FSL was full, the write is inhibited and the carry bit is set in the MSR.

FSL BUS Read Operation

A read from the FSL bus is performed by MicroBlaze using one of the get or getd instructions. A read operations transfers the contents of an input FSL to a general purpose register. The transfer is typically completed in 2 clock cycles for blocking mode reads from the FSL as long as data exists in the FSL FIFO. If the FSL FIFO is empty, the processor stalls at this instruction until the FSL exists flag is set. In the non-blocking mode (instructions with prefix n), the transfer is completed in one or two clock cycles irrespective of whether or not the FSL was empty. In the case the FSL was empty, the transfer of data does not take place and the carry bit is set in the MSR.

Direct FSL Connections

A direct FSL connection can be used to avoid the need for the FSL bus. This can be useful in case no buffering is needed between the two connected IP cores, since the FSL bus FIFO is not included with a direct connection. No buffering reduces the communication latency and required implementation resources.

Each of the MicroBlaze FSL interfaces can either use a direct FSL connection or an FSL bus.

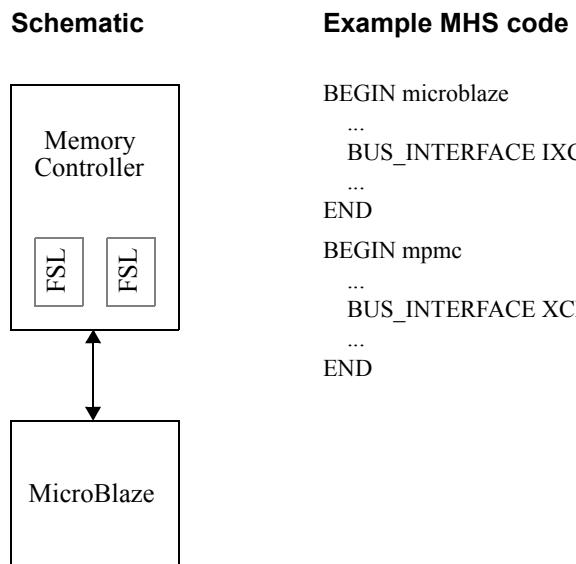
A MicroBlaze DWFSL interface is the initiator on a direct FSL connection, which can only be connected to a DWFSL target. The DWFSL initiator and target have exactly the same signal names, identical to the MFSL signals, depicted in [Table 2-6](#). MicroBlaze uses the DWFSL interface to write data to the target with one of the put or putd instructions.

A MicroBlaze DRFSL interface is the target on a direct FSL connection, which can only be connected to a DRFSL initiator. The DRFSL initiator and target have exactly the same signal names, identical to the SFSL signals, depicted in [Table 2-7](#). MicroBlaze uses the DRFSL interface to read data from the initiator with one of the get or getd instructions.

The Xilinx CacheLink (XCL) interface is implemented with direct FSL connections.

Xilinx CacheLink (XCL) Interface Description

Xilinx CacheLink (XCL) is a high performance solution for external memory accesses. The MicroBlaze CacheLink interface is designed to connect directly to a memory controller with integrated FSL buffers, for example, the MPMC. This method has the lowest latency and minimal number of instantiations (see [Figure 2-7](#)).



**Figure 2-7: CacheLink Connection with Integrated FSL Buffers
(Only Instruction Cache Used in this Example)**

The interface is only available on MicroBlaze when caches are enabled. It is legal to use a CacheLink cache on the instruction side or the data side without caching the other.

How memory locations are accessed depend on the parameter C_ICACHE_ALWAYS_USED for the instruction cache and the parameter C_DCACHE_ALWAYS_USED for the data cache. If the parameter is 1, the cached memory range is always accessed via the CacheLink. If the parameter is 0, the cached memory range is accessed over PLB whenever the caches are software disabled (that is, MSR[DCE]=0 or MSR[ICE]=0).

Memory locations outside the cacheable range are accessed over PLB or LMB.

The CacheLink cache controllers handle 4 or 8-word cache lines, either using critical word first or linear fetch depending on the selected protocol. At the same time the separation from the PLB bus reduces contention for non-cached memory accesses.

CacheLink Signal Interface

The CacheLink signals on MicroBlaze are listed in [Table 2-8](#).

Table 2-8: MicroBlaze Cache Link Signals

| Signal Name | Description | VHDL Type | Direction |
|------------------------|--|----------------------------|-----------|
| ICACHE_FSL_IN_Clk | Clock output to I-side return read data FSL | std_logic | output |
| ICACHE_FSL_IN_Read | Read signal to I-side return read data FSL. | std_logic | output |
| ICACHE_FSL_IN_Data | Read data from I-side return read data FSL | std_logic_vector (0 to 31) | input |
| ICACHE_FSL_IN_Control | FSL control-bit from I-side return read data FSL. Reserved for future use | std_logic | input |
| ICACHE_FSL_IN_Exists | More read data exists in I-side return FSL | std_logic | input |
| ICACHE_FSL_OUT_Clk | Clock output to I-side read access FSL | std_logic | output |
| ICACHE_FSL_OUT_Write | Write new cache miss access request to I-side read access FSL | std_logic | output |
| ICACHE_FSL_OUT_Data | Cache miss access (=address) to I-side read access FSL | std_logic_vector (0 to 31) | output |
| ICACHE_FSL_OUT_Control | FSL control-bit to I-side read access FSL. Reserved for future use | std_logic | output |
| ICACHE_FSL_OUT_Full | FSL access buffer for I-side read accesses is full | std_logic | input |
| DCACHE_FSL_IN_Clk | Clock output to D-side return read data FSL | std_logic | output |
| DCACHE_FSL_IN_Read | Read signal to D-side return read data FSL | std_logic | output |
| DCACHE_FSL_IN_Data | Read data from D-side return read data FSL | std_logic_vector (0 to 31) | input |
| DCACHE_FSL_IN_Control | FSL control bit from D-side return read data FSL | std_logic | input |
| DCACHE_FSL_IN_Exists | More read data exists in D-side return FSL | std_logic | input |
| DCACHE_FSL_OUT_Clk | Clock output to D-side read access FSL | std_logic | output |
| DCACHE_FSL_OUT_Write | Write new cache miss access request to D-side read access FSL | std_logic | output |

Table 2-8: MicroBlaze Cache Link Signals

| Signal Name | Description | VHDL Type | Direction |
|------------------------|---|----------------------------|-----------|
| DCACHE_FSL_OUT_Data | Cache miss access (read address or write address + write data + byte write enable + burst write encoding) to D-side read access FSL | std_logic_vector (0 to 31) | output |
| DCACHE_FSL_OUT_Control | FSL control-bit to D-side read access FSL. Used with address bits [30 to 31] for read/write, byte enable and burst write encoding. | std_logic | output |
| DCACHE_FSL_OUT_Full | FSL access buffer for D-side read accesses is full | std_logic | input |

CacheLink Transactions

All individual CacheLink accesses follow the FSL FIFO based transaction protocol:

- Access information is encoded over the FSL data and control signals (e.g. DCACHE_FSL_OUT_Data, DCACHE_FSL_OUT_Control, ICACHE_FSL_IN_Data, and ICACHE_FSL_IN_Control)
- Information is sent (stored) by raising the write enable signal (e.g. DCACHE_FSL_OUT_Write)
- The sender is only allowed to write if the full signal from the receiver is inactive (e.g. DCACHE_FSL_OUT_Full = 0). The full signal is not used by the instruction cache controller.
- The use of ICACHE_FSL_IN_Read and DCACHE_FSL_IN_Read depends on the selected interface protocol:
 - ◆ With the IXCL and DXCL protocol, information is received (loaded) by raising the read signal. The signal is low, except when the sender signals that new data exists.
 - ◆ With the IXCL2 and DXCL2 protocol, lowering the read signal indicates that the receiver is not able to accept new data. New data is only read when the read signal is high, and the sender signals that data exists.
- The receiver is only allowed to read as long as the sender signals that new data exists (e.g. ICACHE_FSL_IN_Exists = 1)

For details on the generic FSL protocol, please see the *Fast Simplex Link (FSL) Bus* data-sheet, DS449, in the Xilinx EDK IP Documentation.

The CacheLink solution uses one incoming (slave) and one outgoing (master) FSL per cache controller. The outgoing FSL is used to send access requests, while the incoming FSL is used for receiving the requested cache lines. CacheLink also uses a specific encoding of the transaction information over the FSL data and control signals.

The cache lines used for reads in the CacheLink protocol are 4 or 8 words long. Each cache line is either fetched with the critical word first, or in linear order, depending on the selected interface protocol.

- Critical word first is used by the IXCL and DXCL protocol, selected when C_ICACHE_INTERFACE = 0 and C_DCACHE_INTERFACE = 0, respectively. Each cache

line is expected to start with the critical word first (that is, if an access to address 0x348 is a miss with a 4 word cache line, then the returned cache line should have the following address sequence: 0x348, 0x34c, 0x340, 0x344). The cache controller forwards the first word to the execution unit as well as stores it in the cache memory. This allows execution to resume as soon as the first word is back. The cache controller then follows through by filling up the cache line with the remaining 3 or 7 words as they are received.

- Linear fetch is used by the IXCL2 and DXCL2 protocol, selected when `C_ICACHE_INTERFACE = 1` and `C_DCACHE_INTERFACE = 1`, respectively. The address output on the CacheLink is then aligned to the cache line size (that is, if an access to address 0x348 is a miss with a 4 word cache line, then the address output on the CacheLink is 0x340). The cache controller stores data in the cache memory, and forwards the requested word to the execution unit when it is available.

When the parameter `C_DCACHE_USE_WRITEBACK` is set to 1, write operations can store an entire cache line using burst write, as well as single-words. Each cache line is always stored in linear order, and the address output on the CacheLink is aligned to the cache line size. When the parameter `C_DCACHE_USE_WRITEBACK` is cleared to 0, all write operations on the CacheLink are single-word. `C_DCACHE_INTERFACE` must be set to 1 when write-back is used, since burst write is only available with the DXCL2 protocol.

Instruction Cache Read Miss

On a read miss the cache controller performs the following sequence:

1. Write the word aligned⁽¹⁾ or cache line aligned missed address to `ICACHE_FSL_OUT_Data`, with the control bit set low (`ICACHE_FSL_OUT_Control = 0`) to indicate a read access
2. Wait until `ICACHE_FSL_IN_Exists` goes high to indicate that data is available

Note: There must be at least one clock cycle before `ICACHE_FSL_IN_Exists` goes high (that is, at least one wait state must be used).

With the IXCL protocol (critical word first):

3. Store the word from `ICACHE_FSL_IN_Data` to the cache
4. Forward the critical word to the execution unit in order to resume execution
5. Repeat 3 and 4 for the subsequent 3 or 7 words in the cache line

With the IXCL2 protocol (linear fetch):

3. Store words from `ICACHE_FSL_IN_Data` to the cache
4. Forward the relevant word to the execution unit in order to resume execution
5. Store remaining words from `ICACHE_FSL_IN_Data` to the cache

Data Cache Read Miss

On a read miss the cache controller will perform the following sequence:

1. If `DCACHE_FSL_OUT_Full = 1` then stall until it goes low
2. Write the word aligned¹ or cache line aligned missed address to `DCACHE_FSL_OUT_Data`, with the control bit set low (`DCACHE_FSL_OUT_Control = 0`) to indicate a read access
3. Wait until `DCACHE_FSL_IN_Exists` goes high to indicate that data is available

1. Byte and halfword read misses are naturally expected to return complete words, the cache controller then provides the execution unit with the correct bytes.

Note: There must be at least one clock cycle before DCACHE_FSL_IN_Exists goes high (that is, at least one wait state must be used).

With the DXCL protocol (critical word first):

4. Store the word from DCACHE_FSL_IN_Data to the cache
5. Forward the critical word to the execution unit in order to resume execution
6. Repeat 4 and 5 for the subsequent 3 or 7 words in the cache line

With the DXCL2 protocol (linear fetch):

4. Store words from DCACHE_FSL_IN_Data to the cache
5. Forward the requested word to the execution unit in order to resume execution
6. Store remaining words from DCACHE_FSL_IN_Data to the cache

Data Cache Write

When C_DCACHE_INTERFACE is set to 1, the CacheLink can either do burst write or single-word write. A burst write is used when C_DCACHE_USE_WRITEBACK is set to 1 and an entire cache line is valid.

Note that writes to the data cache always are write-through when C_DCACHE_USE_WRITEBACK is cleared to 0, and thus there is a write over the CacheLink regardless of whether there was a hit or miss in the cache.

With the DXCL2 protocol, on a burst cache line write, the cache controller performs the following sequence:

1. If DCACHE_FSL_OUT_Full = 1 then stall until it goes low
2. Write the cache aligned address to DCACHE_FSL_OUT_Data, with the control bit set high (DCACHE_FSL_OUT_Control = 1) to indicate a write access. The two least-significant bits (30:31) of the address are used to encode burst access: 0b10=burst. To separate a burst access from a single byte-write, the control bit for the first data word in step 4 is low for a burst access (DCACHE_FSL_OUT_Control = 0).
3. If DCACHE_FSL_OUT_Full = 1 then stall until it goes low
4. Write the data to be stored to DCACHE_FSL_OUT_Data. The control bit is low (DCACHE_FSL_OUT_Control = 0) for a burst access.
5. Repeat 3 and 4 for the subsequent words in the cache line.

With either the DXCL or DXCL2 protocol, on a single-word write, the cache controller performs the following sequence:

1. If DCACHE_FSL_OUT_Full = 1 then stall until it goes low
2. Write the missed address to DCACHE_FSL_OUT_Data, with the control bit set high (DCACHE_FSL_OUT_Control = 1) to indicate a write access. The two least-significant bits (30:31) of the address are used to encode byte and half-word enables: 0b00=byte0, 0b01=byte1 or halfword0, 0x10=byte2, and 0x11=byte3 or halfword1. The selection of half-word or byte access is based on the control bit for the data word in step 4.
3. If DCACHE_FSL_OUT_Full = 1 then stall until it goes low
4. Write the data to be stored to DCACHE_FSL_OUT_Data. For byte and halfword accesses the data is mirrored onto byte-lanes. Mirroring outputs the byte or halfword to be written on all four byte-lanes or on both halfword-lanes, respectively. The control bit should be low (DCACHE_FSL_OUT_Control = 0) for a word or halfword access, and high for a byte access to separate it from a burst access. Word or halfword accesses can be distinguished by the least significant bit of the address (0=word and 1=halfword).

Debug Interface Description

The debug interface on MicroBlaze is designed to work with the Xilinx Microprocessor Debug Module (MDM) IP core. The MDM is controlled by the Xilinx Microprocessor Debugger (XMD) through the JTAG port of the FPGA. The MDM can control multiple MicroBlaze processors at the same time. The debug signals are grouped in the DEBUG bus. The debug signals on MicroBlaze are listed in [Table 2-9](#).

Table 2-9: MicroBlaze Debug Signals

| Signal Name | Description | VHDL Type | Direction |
|------------------------|--|-----------|-----------|
| Dbg_Clk | JTAG clock from MDM | std_logic | input |
| Dbg_TDI | JTAG TDI from MDM | std_logic | input |
| Dbg_TDO | JTAG TDO to MDM | std_logic | output |
| Dbg_Reg_En | Debug register enable from MDM | std_logic | input |
| Dbg_Shift ¹ | JTAG BSCAN shift signal from MDM | std_logic | input |
| Dbg_Capture | JTAG BSCAN capture signal from MDM | std_logic | input |
| Dbg_Update | JTAG BSCAN update signal from MDM | std_logic | input |
| Debug_Rst ¹ | Reset signal from MDM, active high. Should be held for at least 1 Clk clock cycle. | std_logic | input |

1. Updated for MicroBlaze v7.00: Dbg_Shift added and Debug_Rst included in DEBUG bus

Trace Interface Description

The MicroBlaze core exports a number of internal signals for trace purposes. This signal interface is not standardized and new revisions of the processor may not be backward compatible for signal selection or functionality. It is recommended that you not design custom logic for these signals, but rather to use them via Xilinx provided analysis IP. The trace signals are grouped in the TRACE bus. The current set of trace signals were last updated for MicroBlaze v7.00 and are listed in [Table 2-10](#). The Trace exception types are listed in [Table 2-11](#). All unused Trace exception types are reserved.

Table 2-10: MicroBlaze Trace Signals

| Signal Name | Description | VHDL Type | Direction |
|--------------------------------|---|---|-----------|
| Trace_Valid_Instr | Valid instruction on trace port. | std_logic | output |
| Trace_Instruction ¹ | Instruction code | std_logic_vector (0 to 31) | output |
| Trace_PC ¹ | Program counter | std_logic_vector (0 to 31) | output |
| Trace_Reg_Write ¹ | Instruction writes to the register file | std_logic | output |
| Trace_Reg_Addr ¹ | Destination register address | std_logic_vector (0 to 4) | output |
| Trace_MSR_Reg ¹ | Machine status register | std_logic_vector (0 to 14) ² | output |
| Trace_PID_Reg ^{1,2} | Process identifier register | std_logic_vector (0 to 7) | output |

Table 2-10: MicroBlaze Trace Signals

| Signal Name | Description | VHDL Type | Direction |
|-------------------------------------|---|--|-----------|
| Trace_New_Reg_Value ¹ | Destination register update value | std_logic_vector (0 to 31) | output |
| Trace_Exception_Taken ¹ | Instruction result in taken exception | std_logic | output |
| Trace_Exception_Kind ^{1,3} | Exception type. The description for the exception type is documented below. | std_logic_vector (0 to 4) ² | output |
| Trace_Jump_Taken ¹ | Branch instruction evaluated true, i.e taken | std_logic | output |
| Trace_Jump_Hit ^{1,4,5} | Branch Target Cache hit | std_logic | output |
| Trace_Delay_Slot ¹ | Instruction is in delay slot of a taken branch | std_logic | output |
| Trace_Data_Access ¹ | Valid D-side memory access | std_logic | output |
| Trace_Data_Address ¹ | Address for D-side memory access | std_logic_vector (0 to 31) | output |
| Trace_Data_Write_Value ¹ | Value for D-side memory write access | std_logic_vector (0 to 31) | output |
| Trace_Data_Byte_Enable ¹ | Byte enables for D-side memory access | std_logic_vector (0 to 3) | output |
| Trace_Data_Read ¹ | D-side memory access is a read | std_logic | output |
| Trace_Data_Write ¹ | D-side memory access is a write | std_logic | output |
| Trace_DCache_Req | Data memory address is within D-Cache range | std_logic | output |
| Trace_DCache_Hit | Data memory address is present in D-Cache | std_logic | output |
| Trace_DCache_Rdy ⁴ | Data access completed in D-Cache | std_logic | output |
| Trace_DCache_Read ⁴ | Data read access in D-Cache | std_logic | output |
| Trace_ICache_Req | Instruction memory address is in I-Cache range | std_logic | output |
| Trace_ICache_Hit | Instruction memory address is present in I-Cache | std_logic | output |
| Trace_ICache_Rdy ⁴ | Data access completed in I-Cache | std_logic | output |
| Trace_OF_PipeRun | Pipeline advance for Decode stage | std_logic | output |
| Trace_EX_PipeRun ⁵ | Pipeline advance for Execution stage | std_logic | output |
| Trace_MEM_PipeRun ⁵ | Pipeline advance for Memory stage | std_logic | output |
| Trace_MB_Halted ² | Pipeline is halted by debug | std_logic | output |

1. Valid only when Trace_Valid_Instr = 1
2. Updated for MicroBlaze v7.00: 4 bits added to Trace_MSR_Reg, Trace_PID_Reg added, Trace_MB_Halted added, and 1 bit added to Trace_Exception_Kind
3. Valid only when Trace_Exception_Taken = 1
4. Updated for MicroBlaze v7.30: Trace_Jump_Hit, Trace_DCache_Rdy, Trace_DCache_Read and Trace_ICache_Rdy added
5. Not used with area optimization feature

Table 2-11: Type of Trace Exception

| Trace_Exception_Kind [0:4] | Description |
|----------------------------|---|
| 00000 | Fast Simplex Link exception ¹ |
| 00001 | Unaligned exception |
| 00010 | Illegal Opcode exception |
| 00011 | Instruction Bus exception |
| 00100 | Data Bus exception |
| 00101 | Divide exception |
| 00110 | FPU exception |
| 00111 | Privileged instruction exception ¹ |
| 01010 | Interrupt |
| 01011 | External non maskable break |
| 01100 | External maskable break |
| 10000 | Data storage exception ¹ |
| 10001 | Instruction storage exception ¹ |
| 10010 | Data TLB miss exception ¹ |
| 10011 | Instruction TLB miss exception ¹ |

1. Added for MicroBlaze v7.00

MicroBlaze Core Configurability

The MicroBlaze core has been developed to support a high degree of user configurability. This allows tailoring of the processor to meet specific cost/performance requirements.

Configuration is done via parameters that typically enable, size, or select certain processor features. For example, the instruction cache is enabled by setting the C_USE_ICACHE parameter. The size of the instruction cache, and the cacheable memory range, are all configurable using:

C_CACHE_BYTE_SIZE, C_ICACHE_BASEADDR, and C_ICACHE_HIGHADDR respectively.

Parameters valid for MicroBlaze v7.30 are listed in [Table 2-12](#). Not all of these are recognized by older versions of MicroBlaze; however, the configurability is fully backward compatibility.

Note: Shaded rows indicate that the parameter has a fixed value and *cannot* be modified.

Table 2-12: MPD Parameters

| Parameter Name | Feature/Description | Allowable Values | Default Value | EDK Tool Assigned | VHDL Type |
|----------------------|--|--|---------------|-------------------|----------------------------|
| C_FAMILY | Target Family | aspartan3 aspartan3a aspartan3adsp aspartan3e spartan3 spartan3a spartan3adsp spartan3an spartan3e spartan6 qrvirtex4 qvirtex4 virtex4 virtex5 virtex6 | virtex5 | yes | string |
| C_DATA_SIZE | Data Size | 32 | 32 | NA | integer |
| C_DYNAMIC_BUS_SIZING | Legacy | 1 | 1 | NA | integer |
| C SCO | Xilinx internal | 0 | 0 | NA | integer |
| C_AREA_OPTIMIZED | Select implementation to optimize area with lower instruction throughput | 0, 1 | 0 | | integer |
| C_OPTIMIZATION | Reserved for future use | 0 | 0 | NA | integer |
| C_INTERCONNECT | Select PLB interconnect | 1 | 1 | | integer |
| C_PVR | Processor version register mode selection | 0, 1, 2 | 0 | | integer |
| C_PVR_USER1 | Processor version register USER1 constant | 0x00-0xff | 0x00 | | std_logic_vector (0 to 7) |
| C_PVR_USER2 | Processor version register USER2 constant | 0x00000000-0xffffffff | 0x0000 0000 | | std_logic_vector (0 to 31) |

Table 2-12: MPD Parameters (Continued)

| Parameter Name | Feature/Description | Allowable Values | Default Value | EDK Tool Assigned | VHDL Type |
|------------------------|---|------------------------|---------------|-------------------|------------------|
| C_RESET_MSR | Reset value for MSR register | 0x00, 0x20, 0x80, 0xa0 | 0x00 | | std_logic_vector |
| C_INSTANCE | Instance Name | Any instance name | micro blaze | yes | string |
| C_D_PLB | Data side PLB interface | 0, 1 | 0 | yes | integer |
| C_D_LMB | Data side LMB interface | 0, 1 | 1 | yes | integer |
| C_I_PLB | Instruction side PLB interface | 0, 1 | 0 | yes | integer |
| C_I_LMB | Instruction side LMB interface | 0, 1 | 1 | yes | integer |
| C_USE_BARREL | Include barrel shifter | 0, 1 | 0 | | integer |
| C_USE_DIV | Include hardware divider | 0, 1 | 0 | | integer |
| C_USE_HW_MUL | Include hardware multiplier | 0, 1, 2 | 1 | | integer |
| C_USE_FPU | Include hardware floating point unit | 0, 1, 2 | 0 | | integer |
| C_USE_MSR_INSTR | Enable use of instructions: MSRSET and MSRCLR | 0, 1 | 1 | | integer |
| C_USE_PCMP_INSTR | Enable use of instructions: PCMPBF, PCMPEQ, and PCMPNE | 0, 1 | 1 | | integer |
| C_UNALIGNED_EXCEPTIONS | Enable exception handling for unaligned data accesses | 0, 1 | 0 | | integer |
| C_ILL_OPCODE_EXCEPTION | Enable exception handling for illegal op-code | 0, 1 | 0 | | integer |
| C_IPLB_BUS_EXCEPTION | Enable exception handling for IPLB bus error | 0, 1 | 0 | | integer |
| C_DPLB_BUS_EXCEPTION | Enable exception handling for DPLB bus error | 0, 1 | 0 | | integer |
| C_DIV_ZERO_EXCEPTION | Enable exception handling for division by zero or division overflow | 0, 1 | 0 | | integer |
| C_FPU_EXCEPTION | Enable exception handling for hardware floating point unit exceptions | 0, 1 | 0 | | integer |
| C_OPCODE_0x0_ILLEGAL | Detect opcode 0x0 as an illegal instruction | 0,1 | 0 | | integer |

Table 2-12: MPD Parameters (Continued)

| Parameter Name | Feature/Description | Allowable Values | Default Value | EDK Tool Assigned | VHDL Type |
|--------------------------|---|-------------------------|---------------|-------------------|------------------|
| C_FSL_EXCEPTION | Enable exception handling for Fast Simplex Link | 0,1 | 0 | | integer |
| C_DEBUG_ENABLED | MDM Debug interface | 0,1 | 0 | | integer |
| C_NUMBER_OF_PC_BRK | Number of hardware breakpoints | 0-8 | 1 | | integer |
| C_NUMBER_OF_RD_ADDR_BRK | Number of read address watchpoints | 0-4 | 0 | | integer |
| C_NUMBER_OF_WR_ADDR_BRK | Number of write address watchpoints | 0-4 | 0 | | integer |
| C_INTERRUPT_IS_EDGE | Level/Edge Interrupt | 0, 1 | 0 | yes | integer |
| C_EDGE_IS_POSITIVE | Negative/Positive Edge Interrupt | 0, 1 | 1 | yes | integer |
| C_FSL_LINKS ¹ | Number of FSL interfaces | 0-16 | 0 | yes | integer |
| C_FSL_DATA_SIZE | FSL data bus size | 32 | 32 | NA | integer |
| C_USE_EXTENDED_FSL_INSTR | Enable use of extended FSL instructions | 0, 1 | 0 | | integer |
| C_ICACHE_BASEADDR | Instruction cache base address | 0x00000000 - 0xFFFFFFFF | 0x00000000 | | std_logic_vector |
| C_ICACHE_HIGHADDR | Instruction cache high address | 0x00000000 - 0xFFFFFFFF | 0x3FFFFFFF | | std_logic_vector |
| C_USE_ICACHE | Instruction cache | 0, 1 | 0 | | integer |
| C_ALLOW_ICACHE_WR | Instruction cache write enable | 0, 1 | 1 | | integer |
| C_ICACHE_LINE_LEN | Instruction cache line length | 4, 8 | 4 | | integer |
| C_ICACHE_ALWAYS_USED | Instruction cache CacheLink used for all memory accesses | 0, 1 | 0 | | integer |
| C_ICACHE_INTERFACE | Instruction cache CacheLink interface protocol 0 = IXCL 1 = IXCL2 | 0, 1 | 0 | yes ² | integer |
| C_ICACHE_STREAMS | Instruction cache streams | 0, 1 | 0 | | integer |
| C_ICACHE_VICTIMS | Instruction cache victims | 0, 2, 4, 8 | 0 | | integer |

Table 2-12: MPD Parameters (Continued)

| Parameter Name | Feature/Description | Allowable Values | Default Value | EDK Tool Assigned | VHDL Type |
|------------------------|--|---|---------------|-------------------|------------------|
| C_ADDR_TAG_BITS | Instruction cache address tags | 0-25 | 17 | yes | integer |
| C_CACHE_BYTE_SIZE | Instruction cache size | 64, 128, 256, 512, 1024, 2048, 4096, 8192, 16384, 32768, 65536 ³ | 8192 | | integer |
| C_ICACHE_USE_FSL | Cache over CacheLink instead of peripheral bus for instructions | 1 | 1 | | integer |
| C_DCACHE_BASEADDR | Data cache base address | 0x00000000 - 0xFFFFFFFF | 0x00000000 | | std_logic_vector |
| C_DCACHE_HIGHADDR | Data cache high address | 0x00000000 - 0xFFFFFFFF | 0x3FFF FFFF | | std_logic_vector |
| C_USE_DCACHE | Data cache | 0, 1 | 0 | | integer |
| C_ALLOW_DCACHE_WR | Data cache write enable | 0, 1 | 1 | | integer |
| C_DCACHE_LINE_LEN | Data cache line length | 4, 8 | 4 | | integer |
| C_DCACHE_ALWAYS_USED | Data cache CacheLink used for all accesses | 0, 1 | 0 | | integer |
| C_DCACHE_INTERFACE | Data cache CacheLink interface protocol 0 = DXCL 1 = DXCL2 | 0, 1 | 0 | yes ² | integer |
| C_DCACHE_USE_WRITEBACK | Data cache write-back storage policy used | 0, 1 | 0 | | integer |
| C_DCACHE_VICTIMS | Data cache victims | 0, 2, 4, 8 | 0 | | integer |
| C_DCACHE_ADDR_TAG | Data cache address tags | 0-25 | 17 | yes | integer |
| C_DCACHE_BYTE_SIZE | Data cache size | 64, 128, 256, 512, 1024, 2048, 4096, 8192, 16384, 32768, 65536 ³ | 8192 | | integer |
| C_DCACHE_USE_FSL | Cache over CacheLink instead of peripheral bus for data | 1 | 1 | | integer |
| C_DPLB_DWIDTH | Data side PLB data width | 32 | 32 | | integer |

Table 2-12: MPD Parameters (Continued)

| Parameter Name | Feature/Description | Allowable Values | Default Value | EDK Tool Assigned | VHDL Type |
|-------------------------------|--|------------------|---------------|-------------------|-----------|
| C_DPLB_NATIVE_DWIDTH | Data side PLB native data width | 32 | 32 | | integer |
| C_DPLB_BURST_EN | Data side PLB burst enable | 0 | 0 | | integer |
| C_DPLB_P2P | Data side PLB Point-to-point | 0, 1 | 0 | | integer |
| C_IPLB_DWIDTH | Instruction side PLB data width | 32 | 32 | | integer |
| C_IPLB_NATIVE_DWIDTH | Instruction side PLB native data width | 32 | 32 | | integer |
| C_IPLB_BURST_EN | Instruction side PLB burst enable | 0 | 0 | | integer |
| C_IPLB_P2P | Instruction side PLB Point-to-point | 0, 1 | 0 | | integer |
| C_USE_MMU ⁴ | Memory Management: 0 = None 1 = Usermode 2 = Protection 3 = Virtual | 0, 1, 2, 3 | 0 | | integer |
| C_MMU_DTLB_SIZE ⁴ | Data shadow Translation Look-Aside Buffer size | 1, 2, 4, 8 | 4 | | integer |
| C_MMU_ITLB_SIZE ⁴ | Instruction shadow Translation Look-Aside Buffer size | 1, 2, 4, 8 | 2 | | integer |
| C_MMU_TLB_ACCESS ⁴ | Access to memory management special registers: 0 = Minimal 1 = Read 2 = Write 3 = Full | 0, 1, 2, 3 | 3 | | integer |
| C_MMU_ZONES ⁴ | Number of memory protection zones | 0-16 | 16 | | integer |
| C_USE_INTERRUPT | Enable interrupt handling | 0,1 | 0 | yes | integer |
| C_USE_EXT_BRK | Enable external break handling | 0,1 | 0 | yes | integer |
| C_USE_EXT_NM_BRK | Enable external non-maskable break handling | 0,1 | 0 | yes | integer |

Table 2-12: MPD Parameters (Continued)

| Parameter Name | Feature/Description | Allowable Values | Default Value | EDK Tool Assigned | VHDL Type |
|---|--|------------------|---------------|-------------------|-----------|
| C_USE_BRANCH_TARGET_CACHE ⁴ | Enable Branch Target Cache | 0,1 | 0 | | integer |
| C_BRANCH_TARGET_CACHE_SIZE ⁴ | Branch Target Cache size: 0 = Default 1 = 8 entries 2 = 16 entries 3 = 32 entries 4 = 64 entries 5 = 512 entries 6 = 1024 entries 7 = 2048 entries | 0-7 | 0 | | integer |

1. The number of FSL Links is assigned by the tool itself if you are using the co-processor wizard. If you add the IP manually, you must update the parameter manually.
2. EDK tool assigned value can be overridden by explicit assignment.
3. Not all sizes are permitted in all architectures. The cache uses between 0 and 32 RAMB primitives (0 if cache size is less than 2048).
4. Not available when C_AREA_OPTIMIZED is set to 1.

MicroBlaze Application Binary Interface

This chapter describes MicroBlaze™ Application Binary Interface (ABI), which is important for developing software in assembly language for the soft processor. The MicroBlaze GNU compiler follows the conventions described in this document. Any code written by assembly programmers should also follow the same conventions to be compatible with the compiler generated code. Interrupt and Exception handling is also explained briefly.

This chapter contains the following sections:

- “Data Types”
- “Register Usage Conventions”
- “Stack Convention”
- “Memory Model”
- “Interrupt and Exception Handling”

Data Types

The data types used by MicroBlaze assembly programs are shown in [Table 3-1](#). Data types such as data8, data16, and data32 are used in place of the usual byte, half-word, and word.register

Table 3-1: Data Types in MicroBlaze Assembly Programs

| MicroBlaze data types (for assembly programs) | Corresponding ANSI C data types | Size (bytes) |
|--|------------------------------------|--------------|
| data8 | char | 1 |
| data16 | short | 2 |
| data32 | int | 4 |
| data32 | long int | 4 |
| data32 | float | 4 |
| data32 | enum | 4 |
| data16/data32 | pointer ^a | 2/4 |

a. Pointers to small data areas, which can be accessed by global pointers are data16.

Register Usage Conventions

The register usage convention for MicroBlaze is given in [Table 3-2](#).

Table 3-2: Register Usage Conventions

| Register | Type | Enforcement | Purpose |
|--------------|---------------------------|---|--|
| R0 | Dedicated | HW | Value 0 |
| R1 | Dedicated | SW | Stack Pointer |
| R2 | Dedicated | SW | Read-only small data area anchor |
| R3-R4 | Volatile | SW | Return Values/Temporaries |
| R5-R10 | Volatile | SW | Passing parameters/Temporaries |
| R11-R12 | Volatile | SW | Temporaries |
| R13 | Dedicated | SW | Read-write small data area anchor |
| R14 | Dedicated | HW | Return address for Interrupt |
| R15 | Dedicated | SW | Return address for Sub-routine |
| R16 | Dedicated | HW | Return address for Trap (Debugger) |
| R17 | Dedicated | HW, if configured to support HW exceptions, else SW | Return address for Exceptions |
| R18 | Dedicated | SW | Reserved for Assembler/Compiler Temporaries |
| R19 | Non-volatile | SW | Must be saved across function calls. Callee-save |
| R20 | Dedicated or Non-volatile | SW | Reserved for storing a pointer to the Global Offset Table (GOT) in Position Independent Code (PIC). Non-volatile in non-PIC code. Must be saved across function calls. Callee-save |
| R21-R31 | Non-volatile | SW | Must be saved across function calls. Callee-save |
| RPC | Special | HW | Program counter |
| RMSR | Special | HW | Machine Status Register |
| REAR | Special | HW | Exception Address Register |
| RESR | Special | HW | Exception Status Register |
| RFSR | Special | HW | Floating Point Status Register |
| RBTR | Special | HW | Branch Target Register |
| REDR | Special | HW | Exception Data Register |
| RPID | Special | HW | Process Identifier Register |
| RZPR | Special | HW | Zone Protection Register |
| RTLBLO | Special | HW | Translation Look-Aside Buffer Low Register |
| RTLBHI | Special | HW | Translation Look-Aside Buffer High Register |
| RTLBX | Special | HW | Translation Look-Aside Buffer Index Register |
| RTLBSX | Special | HW | Translation Look-Aside Buffer Search Index |
| RPVR0-RPVR11 | Special | HW | Processor Version Register 0 through 11 |

The architecture for MicroBlaze defines 32 general purpose registers (GPRs). These registers are classified as volatile, non-volatile, and dedicated.

- The volatile registers (also known as caller-save) are used as temporaries and do not retain values across the function calls. Registers R3 through R12 are volatile, of which R3 and R4 are used for returning values to the caller function, if any. Registers R5 through R10 are used for passing parameters between subroutines.
- Registers R19 through R31 retain their contents across function calls and are hence termed as non-volatile registers (a.k.a callee-save). The callee function is expected to save those non-volatile registers, which are being used. These are typically saved to the stack during the prologue and then reloaded during the epilogue.
- Certain registers are used as dedicated registers and programmers are not expected to use them for any other purpose.
 - ◆ Registers R14 through R17 are used for storing the return address from interrupts, subroutines, traps, and exceptions in that order. Subroutines are called using the branch and link instruction, which saves the current Program Counter (PC) onto register R15.
 - ◆ Small data area pointers are used for accessing certain memory locations with 16-bit immediate value. These areas are discussed in the memory model section of this document. The read only small data area (SDA) anchor R2 (Read-Only) is used to access the constants such as literals. The other SDA anchor R13 (Read-Write) is used for accessing the values in the small data read-write section.
 - ◆ Register R1 stores the value of the stack pointer and is updated on entry and exit from functions.
 - ◆ Register R18 is used as a temporary register for assembler operations.
- MicroBlaze includes special purpose registers such as: program counter (rpc), machine status register (rmsr), exception status register (resr), exception address register (rear), floating point status register (rfsr), branch target register (rbtr), exception data register (redr), memory management registers (rpid, rzpr, rtlblo, rtlbhi, rtlbx, rtlbsx), and processor version registers (rpvr0-rpvr11). These registers are not mapped directly to the register file and hence the usage of these registers is different from the general purpose registers. The value of a special purpose registers can be transferred to or from a general purpose register by using **mts** and **mfs** instructions respectively.

Stack Convention

The stack conventions used by MicroBlaze are detailed in [Table 3-3](#).

The shaded area in [Table 3-3](#) denotes a part of the stack frame for a caller function, while the unshaded area indicates the callee frame function. The ABI conventions of the stack frame define the protocol for passing parameters, preserving non-volatile register values, and allocating space for the local variables in a function.

Functions that contain calls to other subroutines are called as non-leaf functions. These non-leaf functions have to create a new stack frame area for its own use. When the program starts executing, the stack pointer has the maximum value. As functions are called, the stack pointer is decremented by the number of words required by every function for its stack frame. The stack pointer of a caller function always has a higher value as compared to the callee function.

Table 3-3: Stack Convention

| | |
|-------------------|---|
| High Address | |
| | Function Parameters for called sub-routine (Arg n .. Arg1) (Optional: Maximum number of arguments required for any called procedure from the current procedure). |
| Old Stack Pointer | Link Register (R15) |
| | Callee Saved Register (R31....R19) (Optional: Only those registers which are used by the current procedure are saved) |
| | Local Variables for Current Procedure (Optional: Present only if Locals defined in the procedure) |
| | Functional Parameters (Arg n .. Arg 1) (Optional: Maximum number of arguments required for any called procedure from the current procedure) |
| New Stack Pointer | Link Register |
| Low Address | |

Consider an example where Func1 calls Func2, which in turn calls Func3. The stack representation at different instances is depicted in [Figure 3-1](#). After the call from Func 1 to Func 2, the value of the stack pointer (SP) is decremented. This value of SP is again decremented to accommodate the stack frame for Func3. On return from Func 3 the value of the stack pointer is increased to its original value in the function, Func 2.

Details of how the stack is maintained are shown in [Figure 3-1](#).

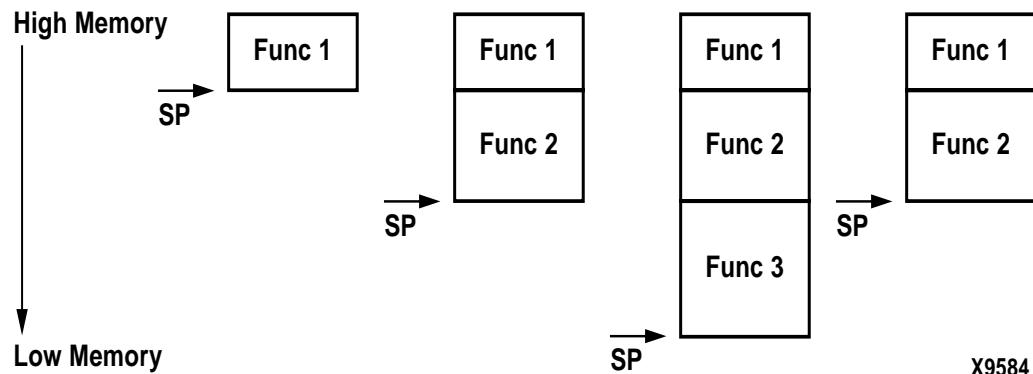


Figure 3-1: Stack Frame

Calling Convention

The caller function passes parameters to the callee function using either the registers (R5 through R10) or on its own stack frame. The callee uses the stack area of the caller to store the parameters passed to the callee.

Refer to [Figure 3-1](#). The parameters for Func 2 are stored either in the registers R5 through R10 or on the stack frame allocated for Func 1.

Memory Model

The memory model for MicroBlaze classifies the data into four different parts: Small Data Area, Data Area, Common Un-Initialized Area, and Literals or Constants.

Small Data Area

Global initialized variables which are small in size are stored in this area. The threshold for deciding the size of the variable to be stored in the small data area is set to 8 bytes in the MicroBlaze C compiler (mb-gcc), but this can be changed by giving a command line option to the compiler. Details about this option are discussed in the *GNU Compiler Tools* chapter. 64 kilobytes of memory is allocated for the small data areas. The small data area is accessed using the read-write small data area anchor (R13) and a 16-bit offset. Allocating small variables to this area reduces the requirement of adding **IMM** instructions to the code for accessing global variables. Any variable in the small data area can also be accessed using an absolute address.

Data Area

Comparatively large initialized variables are allocated to the data area, which can either be accessed using the read-write SDA anchor R13 or using the absolute address, depending on the command line option given to the compiler.

Common Un-Initialized Area

Un-initialized global variables are allocated in the common area and can be accessed either using the absolute address or using the read-write small data area anchor R13.

Literals or Constants

Constants are placed into the read-only small data area and are accessed using the read-only small data area anchor R2.

The compiler generates appropriate global pointers to act as base pointers. The actual values of the SDA anchors are decided by the linker, in the final linking stages. For more information on the various sections of the memory please refer to *MicroBlaze Linker Script Sections* in the *Embedded System Tools Reference Manual*. The compiler generates appropriate sections, depending on the command line options. Please refer to the *GNU Compiler Tools* chapter in the *Embedded System Tools Reference Manual* for more information about these options.

Interrupt and Exception Handling

MicroBlaze assumes certain address locations for handling interrupts and exceptions as indicated in [Table 3-4](#). At these locations, code is written to jump to the appropriate handlers.

Table 3-4: Interrupt and Exception Handling

| On | Hardware jumps to | Software Labels |
|-----------------------------------|-------------------|-----------------------|
| Start / Reset | 0x0 | _start |
| User exception | 0x8 | _exception_handler |
| Interrupt | 0x10 | _interrupt_handler |
| Break (HW/SW) | 0x18 | - |
| Hardware exception | 0x20 | _hw_exception_handler |
| Reserved by Xilinx for future use | 0x28 - 0x4F | - |

The code expected at these locations is as shown below. For programs compiled without the **-xl-mode-xmdstub** compiler option, the `crt0.o` initialization file is passed by the **mb-gcc** compiler to the **mb-ld** linker for linking. This file sets the appropriate addresses of the exception handlers.

For programs compiled with the **-xl-mode-xmdstub** compiler option, the `crt1.o` initialization file is linked to the output program. This program has to be run with the xmdstub already loaded in the memory at address location 0x0. Hence at run-time, the initialization code in `crt1.o` writes the appropriate instructions to location 0x8 through 0x14 depending on the address of the exception and interrupt handlers.

The following is code for passing control to Exception and Interrupt handlers:

```

0x00: bri _start1
0x04: nop
0x08: imm high bits of address (user exception handler)
0x0c: bri _exception_handler
0x10: imm high bits of address (interrupt handler)
0x14: bri _interrupt_handler
0x20: imm high bits of address (HW exception handler)
0x24: bri _hw_exception_handler

```

MicroBlaze allows exception and interrupt handler routines to be located at any address location addressable using 32 bits. The user exception handler code starts with the label `_exception_handler`, the hardware exception handler starts with `_hw_exception_handler`, while the interrupt handler code starts with the label `_interrupt_handler`.

In the current MicroBlaze system, there are dummy routines for interrupt and exception handling, which you can change. In order to override these routines and link your interrupt and exception handlers, you must define the interrupt handler code with an attribute `interrupt_handler`. For more details about the use and syntax of the interrupt handler attribute, please refer to the *GNU Compiler Tools* chapter in the *Embedded System Tools Reference Guide*.

When software breakpoints are used in the Xilinx Microprocessor Debug (XMD) tool, the Break (HW/SW) address location is reserved for handling the software breakpoint.

MicroBlaze Instruction Set Architecture

This chapter provides a detailed guide to the Instruction Set Architecture of MicroBlaze™ and contains the following sections:

- “Notation”
- “Formats”
- “Instructions”

Notation

The symbols used throughout this chapter are defined in [Table 4-1](#).

Table 4-1: Symbol Notation

| Symbol | Meaning |
|----------|--------------------------------------|
| + | Add |
| - | Subtract |
| × | Multiply |
| / | Divide |
| ^ | Bitwise logical AND |
| ∨ | Bitwise logical OR |
| ⊕ | Bitwise logical XOR |
| x̄ | Bitwise logical complement of x |
| ← | Assignment |
| >> | Right shift |
| << | Left shift |
| rx | Register x |
| $x[i]$ | Bit i in register x |
| $x[i:j]$ | Bits i through j in register x |
| = | Equal comparison |
| ≠ | Not equal comparison |
| > | Greater than comparison |
| ≥ | Greater than or equal comparison |

Table 4-1: Symbol Notation (Continued)

| Symbol | Meaning |
|---------------------------|---|
| < | Less than comparison |
| \leq | Less than or equal comparison |
| $\text{sext}(x)$ | Sign-extend x |
| $\text{Mem}(x)$ | Memory location at address x |
| FSLx | FSL interface x |
| LSW(x) | Least Significant Word of x |
| $\text{isDnz}(x)$ | Floating point: true if x is denormalized |
| $\text{isInfinite}(x)$ | Floating point: true if x is $+\infty$ or $-\infty$ |
| $\text{isPosInfinite}(x)$ | Floating point: true if x is $+\infty$ |
| $\text{isNegInfinite}(x)$ | Floating point: true if x $-\infty$ |
| $\text{isNaN}(x)$ | Floating point: true if x is a quiet or signalling NaN |
| $\text{isZero}(x)$ | Floating point: true if x is $+0$ or -0 |
| $\text{isQuietNaN}(x)$ | Floating point: true if x is a quiet NaN |
| $\text{isSigNaN}(x)$ | Floating point: true if x is a signaling NaN |
| $\text{signZero}(x)$ | Floating point: return $+0$ for $x > 0$, and -0 if $x < 0$ |
| $\text{signInfinite}(x)$ | Floating point: return $+\infty$ for $x > 0$, and $-\infty$ if $x < 0$ |

Formats

MicroBlaze uses two instruction formats: Type A and Type B.

Type A

Type A is used for register-register instructions. It contains the opcode, one destination and two source registers.

Type B

Type B is used for register-immediate instructions. It contains the opcode, one destination and one source registers, and a source 16-bit immediate value.

| Opcode | Destination Reg | Source Reg A | | Immediate Value |
|--------|-----------------|--------------|----|-----------------|
| 0 | 6 | 11 | 16 | 31 |

Instructions

This section provides descriptions of MicroBlaze instructions. Instructions are listed in alphabetical order. For each instruction Xilinx® provides the mnemonic, encoding, a description, pseudocode of its semantics, and a list of registers that it modifies.

add

Arithmetic Add

| | | |
|-------|------------|-------------------------------|
| add | rD, rA, rB | Add |
| addc | rD, rA, rB | Add with Carry |
| addk | rD, rA, rB | Add and Keep Carry |
| addkc | rD, rA, rB | Add with Carry and Keep Carry |

| | | | | | | | | | | | | | | | | | |
|----------|----------|----------|----------|----------|----------|----|----|----|----------|----------|----------|----------|----------|----------|----------|----------|----------|
| 0 | 0 | 0 | K | C | 0 | rD | rA | rB | 0 |
| 0 | | | 6 | | | 1 | | 1 | | | 2 | | | | | | 3 |

Description

The sum of the contents of registers rA and rB, is placed into register rD.

Bit 3 of the instruction (labeled as K in the figure) is set to one for the mnemonic addk. Bit 4 of the instruction (labeled as C in the figure) is set to one for the mnemonic addc. Both bits are set to one for the mnemonic addkc.

When an add instruction has bit 3 set (addk, addkc), the carry flag will keep its previous value regardless of the outcome of the execution of the instruction. If bit 3 is cleared (add, addc), then the carry flag will be affected by the execution of the instruction.

When bit 4 of the instruction is set to one (addc, addkc), the content of the carry flag (MSR[C]) affects the execution of the instruction. When bit 4 is cleared (add, addk), the content of the carry flag does not affect the execution of the instruction (providing a normal addition).

Pseudocode

```

if C = 0 then
    (rD) ← (rA) + (rB)
else
    (rD) ← (rA) + (rB) + MSR [C]
if K = 0 then
    MSR [C] ← CarryOut

```

Registers Altered

- rD
 - MSR[C]

Latency

1 cycle

Note

The C bit in the instruction opcode is not the same as the carry bit in the MSR.

The “add r0, r0, r0” (= 0x00000000) instruction is never used by the compiler and usually indicates uninitialized memory. If you are using illegal instruction exceptions you can trap these instructions by setting the MicroBlaze parameter C_OPCODE_0x0_ILLEGAL=1.

addi

Arithmetic Add Immediate

| | | |
|--------|-------------|---|
| addi | rD, rA, IMM | Add Immediate |
| addic | rD, rA, IMM | Add Immediate with Carry |
| addik | rD, rA, IMM | Add Immediate and Keep Carry |
| addikc | rD, rA, IMM | Add Immediate with Carry and Keep Carry |

| | | | | | | | | | |
|----------|----------|----------|----------|----------|----------|-----------|-----------|------------|---|
| 0 | 0 | 1 | K | C | 0 | rD | rA | IMM | |
| 0 | | 6 | | | 1 | | 1 | | 3 |

Description

The sum of the contents of registers rA and the value in the IMM field, sign-extended to 32 bits, is placed into register rD. Bit 3 of the instruction (labeled as K in the figure) is set to one for the mnemonic addik. Bit 4 of the instruction (labeled as C in the figure) is set to one for the mnemonic addic. Both bits are set to one for the mnemonic addikc.

When an addi instruction has bit 3 set (addik, addikc), the carry flag will keep its previous value regardless of the outcome of the execution of the instruction. If bit 3 is cleared (addi, addic), then the carry flag will be affected by the execution of the instruction.

When bit 4 of the instruction is set to one (addic, addikc), the content of the carry flag (MSR[C]) affects the execution of the instruction. When bit 4 is cleared (addi, addik), the content of the carry flag does not affect the execution of the instruction (providing a normal addition).

Pseudocode

```

if C = 0 then
    (rD) ← (rA) + sext(IMM)
else
    (rD) ← (rA) + sext(IMM) + MSR[C]
if K = 0 then
    MSR[C] ← CarryOut

```

Registers Altered

- rD
 - MSR[C]

Latency

1 cycle

Notes

The C bit in the instruction opcode is not the same as the carry bit in the MSR.

By default, Type B Instructions take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “[imm](#),” page 161 for details on using 32-bit immediate values.

and**Logical AND**

and rD, rA, rB

| 1 | 0 | 0 | 0 | 0 | 1 | rD | rA | rB | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
|---|---|---|---|---|---|----|----|----|---|---|---|---|---|---|---|---|---|---|---|
| 0 | | | | | | 6 | 1 | 1 | 2 | | | | | | | | | 3 | 1 |

Description

The contents of register rA are ANDed with the contents of register rB; the result is placed into register rD.

Pseudocode

$(rD) \leftarrow (rA) \wedge (rB)$

Registers Altered

- rD

Latency

1 cycle

andi**Logical AND with Immediate****andi** rD, rA, IMM

| 1 | 0 | 1 | 0 | 0 | 1 | rD | rA | IMM |
|---|---|---|---|---|---|----|----|-----|
| 0 | | 6 | | 1 | | 1 | 1 | 6 |

Description

The contents of register rA are ANDed with the value of the IMM field, sign-extended to 32 bits; the result is placed into register rD.

Pseudocode
$$(rD) \leftarrow (rA) \wedge \text{sext(IMM)}$$
Registers Altered

- rD

Latency

1 cycle

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

and n

Logical AND NOT

andn rD, rA, rB

| | | | | |
|--------------------|-----------|-----------|-----------|----------------------------------|
| 1 0 0 0 1 1 | rD | rA | rB | 0 0 0 0 0 0 0 0 0 0 0 0 0 |
| 0 | 6 | 1 1 | 1 6 | 2 1 1 |

Description

The contents of register rA are ANDed with the logical complement of the contents of register rB; the result is placed into register rD.

Pseudocode

$$(rD) \leftarrow (rA) \wedge (\overline{r}B)$$

Registers Altered

- rD

Latency

1 cycle

andni**Logical AND NOT with Immediate****andni** rD, rA, IMM

| 1 | 0 | 1 | 0 | 1 | 1 | rD | rA | IMM |
|---|---|---|---|---|---|----|----|-----|
| 0 | | | 6 | | | 1 | 1 | 1 |

1
6
1
1
6
3
1

Description

The IMM field is sign-extended to 32 bits. The contents of register rA are ANDed with the logical complement of the extended IMM field; the result is placed into register rD.

Pseudocode

$$(rD) \leftarrow (rA) \wedge (\text{sext(IMM)})$$
Registers Altered

- rD

Latency

1 cycle

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

beq**Branch if Equal**

| | | |
|-------------|--------|----------------------------|
| beq | rA, rB | Branch if Equal |
| beqd | rA, rB | Branch if Equal with Delay |

| | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|----|----|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 1 | 1 | 1 | D | 0 | 0 | 0 | 0 | rA | rB | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | | | | | | 6 | | | | | 1 | 1 | 2 | | | | | | | 3 |

Description

Branch if rA is equal to 0, to the instruction located in the offset value of rB. The target of the branch will be the instruction at address PC + rB.

The mnemonic beqd will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```
If rA = 0 then
    PC ← PC + rB
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set)

Note

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

beqi**Branch Immediate if Equal**

| | | |
|--------------|---------|--------------------------------------|
| beqi | rA, IMM | Branch Immediate if Equal |
| beqid | rA, IMM | Branch Immediate if Equal with Delay |

| 1 | 0 | 1 | 1 | 1 | D | 0 | 0 | 0 | 0 | rA | | IMM |
|---|---|---|---|---|---|---|---|---|---|----|--|-----|
| 0 | | | | | 6 | | 1 | | | 1 | | 3 |

1
6
1
1
6
1

Description

Branch if rA is equal to 0, to the instruction located in the offset value of IMM. The target of the branch will be the instruction at address PC + IMM.

The mnemonic beqid will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```
If rA = 0 then
    PC ← PC + sext(IMM)
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken, or successful branch prediction occurs)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set, or a branch prediction mispredict occurs)

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

bge**Branch if Greater or Equal**

| | | |
|-------------|--------|---------------------------------------|
| bge | rA, rB | Branch if Greater or Equal |
| bged | rA, rB | Branch if Greater or Equal with Delay |

| | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|----|----|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 1 | 1 | 1 | D | 0 | 1 | 0 | 1 | rA | rB | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 0 | | | | | | 6 | | | | | 1 | 1 | 1 | 2 | | | | | | 3 | | 1 |

Description

Branch if rA is greater or equal to 0, to the instruction located in the offset value of rB. The target of the branch will be the instruction at address PC + rB.

The mnemonic bged will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```
If rA >= 0 then
    PC ← PC + rB
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set)

Note

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

bgei**Branch Immediate if Greater or Equal**

| | | |
|--------------|---------|---|
| bgei | rA, IMM | Branch Immediate if Greater or Equal |
| bgeid | rA, IMM | Branch Immediate if Greater or Equal with Delay |

| 1 | 0 | 1 | 1 | 1 | 1 | D | 0 | 1 | 0 | 1 | rA | | | IMM |
|---|---|---|---|---|---|---|---|---|---|---|----|--|--|-----|
| 0 | | | | | | 6 | | | | | 1 | | | 3 |

1 6 1 1 6 1

Description

Branch if rA is greater or equal to 0, to the instruction located in the offset value of IMM. The target of the branch will be the instruction at address PC + IMM.

The mnemonic bgeid will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```
If rA >= 0 then
    PC ← PC + sext(IMM)
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken, or successful branch prediction occurs)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set, or a branch prediction mispredict occurs)

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

bgt**Branch if Greater Than****bgt**

rA, rB

Branch if Greater Than

bgtd

rA, rB

Branch if Greater Than with Delay

| | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|----|----|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 1 | 1 | 1 | D | 0 | 1 | 0 | 0 | rA | rB | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | | | | | | 6 | | | | | 1 | 1 | 2 | | | | | | | 3 | | 1 |

Description

Branch if rA is greater than 0, to the instruction located in the offset value of rB. The target of the branch will be the instruction at address PC + rB.

The mnemonic bgtd will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```

If rA > 0 then
    PC ← PC + rB
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution

```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set)

Note

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

bgti**Branch Immediate if Greater Than**

| | | |
|--------------|---------|---|
| bgti | rA, IMM | Branch Immediate if Greater Than |
| bgtid | rA, IMM | Branch Immediate if Greater Than with Delay |

| 1 | 0 | 1 | 1 | 1 | 1 | D | 0 | 1 | 0 | 0 | rA | | | IMM |
|---|---|---|---|---|---|---|---|---|---|---|----|--|--|-----|
| 0 | | | | | | 6 | | | | | 1 | | | 3 |

1 6 1 1 6 1

Description

Branch if rA is greater than 0, to the instruction located in the offset value of IMM. The target of the branch will be the instruction at address PC + IMM.

The mnemonic bgtid will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```
If rA > 0 then
    PC ← PC + sext(IMM)
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken, or successful branch prediction occurs)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set, or a branch prediction mispredict occurs)

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

ble**Branch if Less or Equal**

| | | |
|-------------|--------|------------------------------------|
| ble | rA, rB | Branch if Less or Equal |
| bled | rA, rB | Branch if Less or Equal with Delay |

| | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|----|----|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 1 | 1 | 1 | D | 0 | 0 | 1 | 1 | rA | rB | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 0 | | | | | | 6 | | | | | 1 | 1 | 1 | 2 | | | | | | 3 | | 1 |

Description

Branch if rA is less or equal to 0, to the instruction located in the offset value of rB. The target of the branch will be the instruction at address PC + rB.

The mnemonic bled will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```
If rA <= 0 then
    PC ← PC + rB
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set)

Note

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

blei**Branch Immediate if Less or Equal**

| | | |
|--------------|---------|--|
| blei | rA, IMM | Branch Immediate if Less or Equal |
| bleid | rA, IMM | Branch Immediate if Less or Equal with Delay |

| 1 | 0 | 1 | 1 | 1 | D | 0 | 0 | 1 | 1 | rA | | IMM |
|---|---|---|---|---|---|---|---|---|---|----|--|-----|
| 0 | | | 6 | | | 1 | | 1 | | 1 | | 3 |

1
1

Description

Branch if rA is less or equal to 0, to the instruction located in the offset value of IMM. The target of the branch will be the instruction at address PC + IMM.

The mnemonic bleid will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```
If rA <= 0 then
    PC ← PC + sext(IMM)
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken, or successful branch prediction occurs)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set, or a branch prediction mispredict occurs)

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

blt**Branch if Less Than**

| | | |
|-------------|--------|--------------------------------|
| blt | rA, rB | Branch if Less Than |
| bltd | rA, rB | Branch if Less Than with Delay |

| | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|----|----|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 1 | 1 | 1 | D | 0 | 0 | 1 | 0 | rA | rB | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | | | | | | 6 | | | | | 1 | 1 | 2 | | | | | | | 3 | |

Description

Branch if rA is less than 0, to the instruction located in the offset value of rB. The target of the branch will be the instruction at address PC + rB.

The mnemonic bltd will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```
If rA < 0 then
    PC ← PC + rB
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set)

Note

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

blti**Branch Immediate if Less Than**

| | | |
|--------------|---------|--|
| blti | rA, IMM | Branch Immediate if Less Than |
| bltid | rA, IMM | Branch Immediate if Less Than with Delay |

| 1 | 0 | 1 | 1 | 1 | 1 | D | 0 | 0 | 1 | 0 | rA | | | IMM |
|---|---|---|---|---|---|---|---|---|---|---|----|--|--|-----|
| 0 | | | | | | 6 | | | | | 1 | | | 3 |

1 6 1 1 6 1

Description

Branch if rA is less than 0, to the instruction located in the offset value of IMM. The target of the branch will be the instruction at address PC + IMM.

The mnemonic bltid will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```
If rA < 0 then
    PC ← PC + sext(IMM)
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken, or successful branch prediction occurs)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set, or a branch prediction mispredict occurs)

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

bne**Branch if Not Equal**

| | | |
|-------------|--------|--------------------------------|
| bne | rA, rB | Branch if Not Equal |
| bned | rA, rB | Branch if Not Equal with Delay |

| | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|--|----|--|----|--|---|---|---|---|---|---|---|---|---|---|---|--|
| 1 | 0 | 0 | 1 | 1 | 1 | D | 0 | 0 | 0 | 1 | | rA | | rB | | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 0 | | | | | | 6 | | | | | | 1 | | 1 | | 2 | | | | | | | | | | 3 | |

Description

Branch if rA not equal to 0, to the instruction located in the offset value of rB. The target of the branch will be the instruction at address PC + rB.

The mnemonic bned will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```
If rA ≠ 0 then
    PC ← PC + rB
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set)

Note

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

bnei**Branch Immediate if Not Equal**

| | | |
|--------------|---------|--|
| bnei | rA, IMM | Branch Immediate if Not Equal |
| bneid | rA, IMM | Branch Immediate if Not Equal with Delay |

| 1 | 0 | 1 | 1 | 1 | D | 0 | 0 | 0 | 1 | rA | | | IMM |
|---|---|---|---|---|---|---|---|---|---|----|--|--|-----|
| 0 | | | 6 | | | 1 | | | 1 | | | | 3 |

1 6 1 1 6 1 3
 1 1 1 6 1 1

Description

Branch if rA not equal to 0, to the instruction located in the offset value of IMM. The target of the branch will be the instruction at address PC + IMM.

The mnemonic bneid will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```
If rA ≠ 0 then
    PC ← PC + sext(IMM)
else
    PC ← PC + 4
if D = 1 then
    allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if branch is not taken, or successful branch prediction occurs)
- 2 cycles (if branch is taken and the D bit is set)
- 3 cycles (if branch is taken and the D bit is not set, or a branch prediction mispredict occurs)

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

br**Unconditional Branch**

| | | |
|--------------|--------|-------------------------------------|
| br | rB | Branch |
| bra | rB | Branch Absolute |
| brd | rB | Branch with Delay |
| brad | rB | Branch Absolute with Delay |
| brld | rD, rB | Branch and Link with Delay |
| brald | rD, rB | Branch Absolute and Link with Delay |

| 1 | 0 | 0 | 1 | 1 | 0 | rD | D | A | L | 0 | 0 | rB | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
|---|---|---|---|---|---|----|---|---|---|---|---|----|---|---|---|---|---|---|---|---|---|---|---|
| 0 | | | 6 | | | 1 | | | 1 | | | 1 | 2 | | | | | | | | 3 | | 1 |

Description

Branch to the instruction located at address determined by rB.

The mnemonics brld and brald will set the L bit. If the L bit is set, linking will be performed. The current value of PC will be stored in rD.

The mnemonics bra, brad and brald will set the A bit. If the A bit is set, it means that the branch is to an absolute value and the target is the value in rB, otherwise, it is a relative branch and the target will be PC + rB.

The mnemonics brd, brad, brld and brald will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction.

If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

Pseudocode

```

if L = 1 then
    (rD) ← PC
if A = 1 then
    PC ← (rB)
else
    PC ← PC + (rB)
    if D = 1 then
        allow following instruction to complete execution

```

Registers Altered

- rD
- PC

Latency

- 2 cycles (if the D bit is set)
- 3 cycles (if the D bit is not set)

Note

The instructions brl and bral are not available. A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

bri**Unconditional Branch Immediate**

| | | |
|---------------|---------|---|
| bri | IMM | Branch Immediate |
| brai | IMM | Branch Absolute Immediate |
| brid | IMM | Branch Immediate with Delay |
| braid | IMM | Branch Absolute Immediate with Delay |
| brlid | rD, IMM | Branch and Link Immediate with Delay |
| bralid | rD, IMM | Branch Absolute and Link Immediate with Delay |

| 1 0 1 1 1 0 | rD | D A L 0 0 | IMM | |
|-------------|----|-----------|--------|--------|
| 0 | 6 | 1 1 | 1 6 | 3 1 |

Description

Branch to the instruction located at address determined by IMM, sign-extended to 32 bits.

The mnemonics brlid and bralid will set the L bit. If the L bit is set, linking will be performed. The current value of PC will be stored in rD.

The mnemonics brai, braid and bralid will set the A bit. If the A bit is set, it means that the branch is to an absolute value and the target is the value in IMM, otherwise, it is a relative branch and the target will be PC + IMM.

The mnemonics brid, braid, brlid and bralid will set the D bit. The D bit determines whether there is a branch delay slot or not. If the D bit is set, it means that there is a delay slot and the instruction following the branch (that is, in the branch delay slot) is allowed to complete execution before executing the target instruction. If the D bit is not set, it means that there is no delay slot, so the instruction to be executed after the branch is the target instruction.

As a special case, when MicroBlaze is configured to use an MMU (`C_USE_MMU >= 1`) and “bralid rd, 0x8“ is used to perform a User Vector Exception, the Machine Status Register bits User Mode and Virtual Mode are cleared.

Pseudocode

```

if L = 1 then
    (rD) ← PC
if A = 1 then
    PC ← sext(IMM)
else
    PC ← PC + sext(IMM)
if D = 1 then
    allow following instruction to complete execution
if D = 1 and A = 1 and L = 1 and IMM = 0x8 then
    MSR [UMS] ← MSR [UM]
    MSR [VMS] ← MSR [VM]
    MSR [UM] ← 0
    MSR [VM] ← 0

```

Registers Altered

- rD
- PC
- MSR[UM], MSR[VM]

Latency

- 1 cycle (if successful branch prediction occurs)
- 2 cycles (if the D bit is set)
- 3 cycles (if the D bit is not set, or a branch prediction mispredict occurs)

Notes

The instructions brli and brali are not available.

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

brk

Break

brk rD, rB

| | | | | |
|--------------------|-----------|------------------|-----------|--|
| 1 0 0 1 1 0 | rD | 0 1 1 0 0 | rB | 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 |
| 0 | 6 | 11 | 16 | 21 31 |

Description

Branch and link to the instruction located at address value in rB. The current value of PC will be stored in rD. The BIP flag in the MSR will be set, and the reservation bit will be cleared.

When MicroBlaze is configured to use an MMU (`C_USE_MMU >= 1`) this instruction is privileged. This means that if the instruction is attempted in User Mode (`MSR[UM] = 1`) a Privileged Instruction exception occurs.

Pseudocode

```

if MSR[UM] = 1 then
    ESR[EC] ← 00111
else
    (rD) ← PC
    PC ← (rB)
    MSR[BIP] ← 1
    Reservation ← 0

```

Registers Altered

- rD
- PC
- MSR[BIP]
- ESR[EC], in case a privileged instruction exception is generated

Latency

- 3 cycles

brki**Break Immediate**

brki rD, IMM

| | | | |
|--------------------|-----------|------------------|------------|
| 1 0 1 1 1 0 | rD | 0 1 1 0 0 | IMM |
| 0 | 6 | 11 | 16 |

31

Description

Branch and link to the instruction located at address value in IMM, sign-extended to 32 bits. The current value of PC will be stored in rD. The BIP flag in the MSR will be set, and the reservation bit will be cleared.

When MicroBlaze is configured to use an MMU ($C_USE_MMU \geq 1$) this instruction is privileged, except as a special case when “brki rD, 0x8” or “brki rD, 0x18” is used to perform a Software Break. This means that, apart from the special case, if the instruction is attempted in User Mode (MSR [UM] = 1) a Privileged Instruction exception occurs.

As a special case, when MicroBlaze is configured to use an MMU ($C_USE_MMU \geq 1$) and “brki rD, 0x8” or “brki rD, 0x18” is used to perform a Software Break, the Machine Status Register bits User Mode and Virtual Mode are cleared.

Pseudocode

```

if MSR[UM] = 1 and IMM ≠ 0x8 and IMM ≠ 0x18 then
    ESR[EC] ← 00111
else
    (rD) ← PC
    PC ← sext(IMM)
    MSR[BIP] ← 1
    Reservation ← 0
    if IMM = 0x8 or IMM = 0x18 then
        MSR[UMS] ← MSR[UM]   MSR[VMS] ← MSR[VM]
        MSR[UM] ← 0           MSR[VM] ← 0

```

Registers Altered

- rD, unless an exception is generated, in which case the register is unchanged
- PC
- MSR[BIP], MSR[UM], MSR[VM]
- ESR[EC], in case a privileged instruction exception is generated

Latency

- 3 cycles

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

bs**Barrel Shift**

| | | |
|-------------|------------|---------------------------------|
| bsrl | rD, rA, rB | Barrel Shift Right Logical |
| bsra | rD, rA, rB | Barrel Shift Right Arithmetical |
| bsll | rD, rA, rB | Barrel Shift Left Logical |

| 0 | 1 | 0 | 0 | 0 | 1 | rD | rA | rB | S | T | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
|---|---|---|---|---|---|----|----|----|---|---|---|---|---|---|---|---|---|---|---|
| 0 | | 6 | | 1 | | 1 | | 1 | 2 | | 3 | | | | | | | | 1 |

Description

Shifts the contents of register rA by the amount specified in register rB and puts the result in register rD.

The mnemonic bsll sets the S bit (Side bit). If the S bit is set, the barrel shift is done to the left. The mnemonics bsrl and bsra clear the S bit and the shift is done to the right.

The mnemonic bsra will set the T bit (Type bit). If the T bit is set, the barrel shift performed is Arithmetical. The mnemonics bsrl and bsll clear the T bit and the shift performed is Logical.

Pseudocode

```

if S = 1 then
    (rD) ← (rA) << (rB) [27:31]
else
    if T = 1 then
        if ((rB) [27:31]) ≠ 0 then
            (rD) [0:(rB) [27:31]-1] ← (rA)[0]
            (rD) [(rB) [27:31]:31] ← (rA) >> (rB) [27:31]
        else
            (rD) ← (rA)
        else
            (rD) ← (rA) >> (rB) [27:31]
    end
end

```

Registers Altered

- rD

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

Note

These instructions are optional. To use them, MicroBlaze has to be configured to use barrel shift instructions (C_USE_BARREL=1).

bsi**Barrel Shift Immediate**

| | | |
|--------------|-------------|---|
| bsrli | rD, rA, IMM | Barrel Shift Right Logical Immediate |
| bsrai | rD, rA, IMM | Barrel Shift Right Arithmetical Immediate |
| bslli | rD, rA, IMM | Barrel Shift Left Logical Immediate |

| 0 | 1 | 1 | 0 | 0 | 1 | rD | rA | 0 | 0 | 0 | 0 | 0 | S | T | 0 | 0 | 0 | 0 | IMM |
|---|---|---|---|---|---|----|----|---|---|---|---|---|---|---|---|---|---|---|-----|
| 0 | | | 6 | | | 1 | | 1 | | 6 | | | 2 | | 2 | | 7 | | 3 |

Description

Shifts the contents of register rA by the amount specified by IMM and puts the result in register rD.

The mnemonic bsll sets the S bit (Side bit). If the S bit is set, the barrel shift is done to the left. The mnemonics bsrli and bsra clear the S bit and the shift is done to the right.

The mnemonic bsra will set the T bit (Type bit). If the T bit is set, the barrel shift performed is Arithmetical. The mnemonics bsrli and bsll clear the T bit and the shift performed is Logical.

Pseudocode

```

if S = 1 then
    (rD) ← (rA) << IMM
else
    if T = 1 then
        if IMM ≠ 0 then
            (rD) [0:IMM-1] ← (rA)[0]
            (rD) [IMM:31] ← (rA) >> IMM
        else
            (rD) ← (rA)
    else
        (rD) ← (rA) >> IMM

```

Registers Altered

- rD

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

Notes

These are not Type B Instructions. There is no effect from a preceding imm instruction.

These instructions are optional. To use them, MicroBlaze has to be configured to use barrel shift instructions (C_USE_BARREL=1).

cmp

Integer Compare

| | | |
|------|------------|-------------------------------|
| cmp | rD, rA, rB | compare rB with rA (signed) |
| cmpl | rD, rA, rB | compare rB with rA (unsigned) |

| | | | | |
|--------------------|-----------|-----------|-----------|------------------------------|
| 0 0 0 1 0 1 | rD | rA | rB | 0 0 0 0 0 0 0 0 0 U 1 |
| 0 | 6 | 1 1 | 1 6 | 2 1 1 |

Description

The contents of register rA is subtracted from the contents of register rB and the result is placed into register rD.

The MSB bit of rD is adjusted to show true relation between rA and rB. If the U bit is set, rA and rB are considered unsigned values. If the U bit is clear, rA and rB are considered signed values.

Pseudocode

$$(rD) \leftarrow (rB) + (\overline{rA}) + 1$$

$$(rD)_{(MSB)} \leftarrow (rA) > (rB)$$

Registers Altered

- rD

Latency

- 1 cycle

fadd**Floating Point Arithmetic Add**

fadd rD, rA, rB Add

| | | | | |
|--------------------|-----------|-----------|-----------|--|
| 0 1 0 1 1 0 | rD | rA | rB | 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 |
| 0 | 6 | 11 | 16 | 21 |

31
Description

The floating point sum of registers rA and rB, is placed into register rD.

Pseudocode

```

if isDnz(rA) or isDnz(rB) then
    (rD) ← 0xFFC00000
    FSR[DO] ← 1
    ESR[EC] ← 00110
else if isSigNaN(rA) or isSigNaN(rB) or
    (isPosInfinite(rA) and isNegInfinite(rB)) or
    (isNegInfinite(rA) and isPosInfinite(rB)) then
    (rD) ← 0xFFC00000
    FSR[IO] ← 1
    ESR[EC] ← 00110
else if isQuietNaN(rA) or isQuietNaN(rB) then
    (rD) ← 0xFFC00000
else if isDnz((rA)+(rB)) then
    (rD) ← signZero((rA)+(rB))
    FSR[UF] ← 1
    ESR[EC] ← 00110
else if isNaN((rA)+(rB)) then
    (rD) ← signInfinite((rA)+(rB))
    FSR[OF] ← 1
    ESR[EC] ← 00110
else
    (rD) ← (rA) + (rB)

```

Registers Altered

- rD, unless an FP exception is generated, in which case the register is unchanged
- ESR[EC], if an FP exception is generated
- FSR[IO,UF,OF,DO]

Latency

- 4 cycles with C_AREA_OPTIMIZED=0
- 6 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is only available when the MicroBlaze parameter C_USE_FPU is greater than 0.

frsub

Reverse Floating Point Arithmetic Subtraction

frsub rD, rA, rB Reverse subtract

| | | | | | |
|--------------------|-----------|-----------|-----------|--------------------------------|----|
| 0 1 0 1 1 0 | rD | rA | rB | 0 0 0 1 0 0 0 0 0 0 0 0 | |
| 0 | 6 | 11 | 16 | 21 | 31 |

Description

The floating point value in rA is subtracted from the floating point value in rB and the result is placed into register rD.

Pseudocode

```

if isDnz(rA) or isDnz(rB) then
    (rD) ← 0xFFC00000
    FSR[DO] ← 1
    ESR[EC] ← 00110
else if (isSigNaN(rA) or isSigNaN(rB) or
        (isPosInfinite(rA) and isPosInfinite(rB)) or
        (isNegInfinite(rA) and isNegInfinite(rB))) then
    (rD) ← 0xFFC00000
    FSR[IO] ← 1
    ESR[EC] ← 00110
else if isQuietNaN(rA) or isQuietNaN(rB) then
    (rD) ← 0xFFC00000
else if isDnz((rB) - (rA)) then
    (rD) ← signZero((rB) - (rA))
    FSR[UF] ← 1
    ESR[EC] ← 00110
else if isNaN((rB) - (rA)) then
    (rD) ← signInfinite((rB) - (rA))
    FSR[OF] ← 1
    ESR[EC] ← 00110
else
    (rD) ← (rB) - (rA)

```

Registers Altered

- rD, unless an FP exception is generated, in which case the register is unchanged
- ESR[EC], if an FP exception is generated
- FSR[IO,UF,OF,DO]

Latency

- 4 cycles with C_AREA_OPTIMIZED=0
- 6 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is only available when the MicroBlaze parameter C_USE_FPU is greater than 0.

fmul**Floating Point Arithmetic Multiplication**

fmul rD, rA, rB Multiply

| | | | | |
|--------------------|-----------|-----------|-----------|--|
| 0 1 0 1 1 0 | rD | rA | rB | 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 |
| 0 | 6 | 11 | 16 | 21 |

31
Description

The floating point value in rA is multiplied with the floating point value in rB and the result is placed into register rD.

Pseudocode

```

if isDnz(rA) or isDnz(rB) then
    (rD) ← 0xFFC00000
    FSR[DO] ← 1
    ESR[EC] ← 00110
else
    if isSigNaN(rA) or isSigNaN(rB) or (isZero(rA) and isInfinite(rB)) or
        (isZero(rB) and isInfinite(rA)) then
        (rD) ← 0xFFC00000
        FSR[IO] ← 1
        ESR[EC] ← 00110
    else if isQuietNaN(rA) or isQuietNaN(rB) then
        (rD) ← 0xFFC00000
    else if isDnz((rB)*(rA)) then
        (rD) ← signZero((rA)*(rB))
        FSR[UF] ← 1
        ESR[EC] ← 00110
    else if isNaN((rB)*(rA)) then
        (rD) ← signInfinite((rB)*(rA))
        FSR[OF] ← 1
        ESR[EC] ← 00110
    else
        (rD) ← (rB) * (rA)

```

Registers Altered

- rD, unless an FP exception is generated, in which case the register is unchanged
- ESR[EC], if an FP exception is generated
- FSR[IO,UF,OF,DO]

Latency

- 4 cycles with C_AREA_OPTIMIZED=0
- 6 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is only available when the MicroBlaze parameter C_USE_FPU is greater than 0.

fdiv**Floating Point Arithmetic Division**

fdiv rD, rA, rB Divide

| | | | | | |
|--------------------|-----------|-----------|-----------|--------------------------------|----|
| 0 1 0 1 1 0 | rD | rA | rB | 0 0 1 1 0 0 0 0 0 0 0 0 | |
| 0 | 6 | 11 | 16 | 21 | 31 |

Description

The floating point value in rB is divided by the floating point value in rA and the result is placed into register rD.

Pseudocode

```

if isDnz(rA) or isDnz(rB) then
    (rD) ← 0xFFC00000
    FSR[DO] ← 1
    ESR[EC] ← 00110
else
    if isSigNaN(rA) or isSigNaN(rB) or (isZero(rA) and isZero(rB)) or
        (isInfinite(rA) and isInfinite(rB)) then
        (rD) ← 0xFFC00000
        FSR[IO] ← 1
        ESR[EC] ← 00110
    else if isQuietNaN(rA) or isQuietNaN(rB) then
        (rD) ← 0xFFC00000
    else if isZero(rA) and not isInfinite(rB) then
        (rD) ← signInfinite((rB)/(rA))
        FSR[DZ] ← 1
        ESR[EC] ← 00110
    else if isDnz((rB)/(rA)) then
        (rD) ← signZero((rB)/(rA))
        FSR[UF] ← 1
        ESR[EC] ← 00110
    else if isNaN((rB)/(rA)) then
        (rD) ← signInfinite((rB)/(rA))
        FSR[OF] ← 1
        ESR[EC] ← 00110
    else
        (rD) ← (rB) / (rA)

```

Registers Altered

- rD, unless an FP exception is generated, in which case the register is unchanged
- ESR[EC], if an FP exception is generated
- FSR[IO,UF,OF,DO,DZ]

Latency

- 28 cycles with C_AREA_OPTIMIZED=0, 30 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is only available when the MicroBlaze parameter C_USE_FPU is greater than 0.

fcmp**Floating Point Number Comparison**

| | | |
|---------|------------|--|
| fcmp.un | rD, rA, rB | Unordered floating point comparison |
| fcmp.lt | rD, rA, rB | Less-than floating point comparison |
| fcmp.eq | rD, rA, rB | Equal floating point comparison |
| fcmp.le | rD, rA, rB | Less-or-Equal floating point comparison |
| fcmp.gt | rD, rA, rB | Greater-than floating point comparison |
| fcmp.ne | rD, rA, rB | Not-Equal floating point comparison |
| fcmp.ge | rD, rA, rB | Greater-or-Equal floating point comparison |

| | | | | | | |
|-------------|----|----|----|---------|-------|---------|
| 0 1 0 1 1 0 | rD | rA | rB | 0 1 0 0 | OpSel | 0 0 0 0 |
| 0 | 6 | 11 | 16 | 21 | 25 | 28 31 |

Description

The floating point value in rB is compared with the floating point value in rA and the comparison result is placed into register rD. The OpSel field in the instruction code determines the type of comparison performed.

Pseudocode

```

if isDnz(rA) or isDnz(rB) then
    (rD) ← 0
    FSR[DO] ← 1
    ESR[EC] ← 00110
else
    {read out behavior from Table 4-2}

```

Registers Altered

- rD, unless an FP exception is generated, in which case the register is unchanged
- ESR[EC], if an FP exception is generated
- FSR[IO,DO]

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 3 cycles with C_AREA_OPTIMIZED=1

Note

These instructions are only available when the MicroBlaze parameter C_USE_FPU is greater than 0.

[Table 4-2, page 152](#) lists the floating point comparison operations.

Table 4-2: Floating Point Comparison Operation

| Comparison Type | | Operand Relationship | | | | |
|------------------|-------|----------------------|-------------|-------------|--|--|
| Description | OpSel | (rB) > (rA) | (rB) < (rA) | (rB) = (rA) | isSigNaN(rA) or isSigNaN(rB) | isQuietNaN(rA) or isQuietNaN(rB) |
| Unordered | 000 | (rD) ← 0 | (rD) ← 0 | (rD) ← 0 | (rD) ← 1 FSR[IO] ← 1 ESR[EC] ← 00110 | (rD) ← 1 |
| Less-than | 001 | (rD) ← 0 | (rD) ← 1 | (rD) ← 0 | (rD) ← 0 FSR[IO] ← 1 ESR[EC] ← 00110 | (rD) ← 0 FSR[IO] ← 1 ESR[EC] ← 00110 |
| Equal | 010 | (rD) ← 0 | (rD) ← 0 | (rD) ← 1 | (rD) ← 0 FSR[IO] ← 1 ESR[EC] ← 00110 | (rD) ← 0 |
| Less-or-equal | 011 | (rD) ← 0 | (rD) ← 1 | (rD) ← 1 | (rD) ← 0 FSR[IO] ← 1 ESR[EC] ← 00110 | (rD) ← 0 FSR[IO] ← 1 ESR[EC] ← 00110 |
| Greater-than | 100 | (rD) ← 1 | (rD) ← 0 | (rD) ← 0 | (rD) ← 0 FSR[IO] ← 1 ESR[EC] ← 00110 | (rD) ← 0 FSR[IO] ← 1 ESR[EC] ← 00110 |
| Not-equal | 101 | (rD) ← 1 | (rD) ← 1 | (rD) ← 0 | (rD) ← 1 FSR[IO] ← 1 ESR[EC] ← 00110 | (rD) ← 1 |
| Greater-or-equal | 110 | (rD) ← 1 | (rD) ← 0 | (rD) ← 1 | (rD) ← 0 FSR[IO] ← 1 ESR[EC] ← 00110 | (rD) ← 0 FSR[IO] ← 1 ESR[EC] ← 00110 |

flt

Floating Point Convert Integer to Float

flt rD, rA

| | | | | |
|--------------------|-----------|-----------|----------|------------------------------------|
| 0 1 0 1 1 0 | rD | rA | 0 | 0 1 0 1 0 0 0 0 0 0 0 0 0 0 |
| 0 | 6 | 11 | 16 | 21 |

Description

Converts the signed integer in register rA to floating point and puts the result in register rD. This is a 32-bit rounding signed conversion that will produce a 32-bit floating point result.

Pseudocode

(rD) \leftarrow float ((rA))

Registers Altered

- rD

Latency

- 4 cycles with C_AREA_OPTIMIZED=0
 - 6 cycles with C AREA OPTIMIZED=1

Note

This instruction is only available when the MicroBlaze parameter C_USE_FPU is set to 2.

fint**Floating Point Convert Float to Integer**

fint rD, rA

| | | | | |
|--------------------|-----------|-----------|----------|--------------------------------------|
| 0 1 0 1 1 0 | rD | rA | 0 | 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 |
| 0 | 6 | 11 | 16 | 21 |

31
Description

Converts the floating point number in register rA to a signed integer and puts the result in register rD. This is a 32-bit signed conversion that will produce a 32-bit integer result.

Pseudocode

```

if isDnz(rA) then
    (rD) ← 0xFFC00000
    FSR[DO] ← 1
    ESR[EC] ← 00110
else if isNaN(rA) then
    (rD) ← 0xFFC00000
    FSR[IO] ← 1
    ESR[EC] ← 00110
else if isInf(rA) or (rA) < -231 or (rA) > 231 - 1 then
    (rD) ← 0xFFC00000
    FSR[IO] ← 1
    ESR[EC] ← 00110
else
    (rD) ← int ((rA))

```

Registers Altered

- rD, unless an FP exception is generated, in which case the register is unchanged
- ESR[EC], if an FP exception is generated
- FSR[IO,UF,OF,DO]

Latency

- 5 cycles with C_AREA_OPTIMIZED=0
- 7 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is only available when the MicroBlaze parameter C_USE_FPU is set to 2.

fsqrt**Floating Point Arithmetic Square Root**

fsqrt rD, rA Square Root

| | | | | |
|--------------------|-----------|-----------|----------|--------------------------------|
| 0 1 0 1 1 0 | rD | rA | 0 | 0 1 1 1 0 0 0 0 0 0 0 0 |
| 0 | 6 | 11 | 16 | 21 |

31
Description

Performs a floating point square root on the value in rA and puts the result in register rD.

Pseudocode

```

if isDnz(rA) then
    (rD) ← 0xFFC00000
    FSR[DO] ← 1
    ESR[EC] ← 00110
else if isSigNaN(rA) then
    (rD) ← 0xFFC00000
    FSR[IO] ← 1
    ESR[EC] ← 00110
else if isQuietNaN(rA) then
    (rD) ← 0xFFC00000
else if (rA) < 0 then
    (rD) ← 0xFFC00000
    FSR[IO] ← 1
    ESR[EC] ← 00110
else if (rA) = -0 then
    (rD) ← -0
else
    (rD) ← sqrt ((rA))

```

Registers Altered

- rD, unless an FP exception is generated, in which case the register is unchanged
- ESR[EC], if an FP exception is generated
- FSR[IO,UF,OF,DO]

Latency

- 27 cycles with C_AREA_OPTIMIZED=0
- 29 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is only available when the MicroBlaze parameter C_USE_FPU is set to 2.

get**get from fsl interface**

| | | |
|----------------|----------|---|
| <i>tmeaget</i> | rD, FSLx | get data from FSL x <i>t</i> = test-only <i>n</i> = non-blocking <i>e</i> = exception if control bit set <i>a</i> = atomic |
| <i>tmcaget</i> | rD, FSLx | get control from FSL x <i>t</i> = test-only <i>n</i> = non-blocking <i>e</i> = exception if control bit not set <i>a</i> = atomic |

| | | | | |
|--------------------|-----------|------------------|--------------------------------|-------------|
| 0 1 1 0 1 1 | rD | 0 0 0 0 0 | 0 n c t a e 0 0 0 0 0 0 | FSLx |
| 0 | 6 | 11 | 16 | 28 31 |

Description

MicroBlaze will read from the FSLx interface and place the result in register rD.

The get instruction has 32 variants.

The blocking versions (when ‘n’ bit is ‘0’) will stall MicroBlaze until the data from the FSL interface is valid. The non-blocking versions will not stall MicroBlaze and will set carry to ‘0’ if the data was valid and to ‘1’ if the data was invalid. In case of an invalid access the destination register contents is undefined.

All data get instructions (when ‘c’ bit is ‘0’) expect the control bit from the FSL interface to be ‘0’. If this is not the case, the instruction will set MSR[FSL_Error] to ‘1’. All control get instructions (when ‘c’ bit is ‘1’) expect the control bit from the FSL interface to be ‘1’. If this is not the case, the instruction will set MSR[FSL_Error] to ‘1’.

The exception versions (when ‘e’ bit is ‘1’) will generate an exception if there is a control bit mismatch. In this case ESR is updated with EC set to the exception cause and ESS set to the FSL index. The target register, rD, is not updated when an exception is generated, instead the FSL data is stored in EDR.

The test versions (when ‘t’ bit is ‘1’) will be handled as the normal case, except that the read signal to the FSL link is not asserted.

Atomic versions (when ‘a’ bit is ‘1’) are not interruptible. This means that a sequence of atomic FSL instructions can be grouped together without an interrupt breaking the program flow. However, note that exceptions may still occur.

When MicroBlaze is configured to use an MMU ($C_USE_MMU \geq 1$) these instructions are privileged. This means that if these instructions are attempted in User Mode (MSR [UM]=1) a Privileged Instruction exception occurs.

Pseudocode

```

if MSR[UM] = 1 then
    ESR[EC] ← 00111
else
    (rD) ← FSLx
    if (n = 1) then
        MSR[Carry] ← not (FSLx Exists bit)
    if (FSLx Control bit ≠ c) and (FSLx Exists bit) then
        MSR[FSL_Error] ← 1
    if (e = 1) then
        ESR[EC] ← 00000
        ESR[ESS] ← instruction bits [28:31]
        EDR ← FSLx

```

Registers Altered

- rD, unless an exception is generated, in which case the register is unchanged
- MSR[FSL_Error]
- MSR[Carry]
- ESR[EC], in case an FSL exception or a privileged instruction exception is generated
- ESR[ESS], in case an FSL exception is generated
- EDR, in case an FSL exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

The blocking versions of this instruction will stall the pipeline of MicroBlaze until the instruction can be completed. Interrupts are served when the parameter C_USE_EXTENDED_FSL_INSTR is set to 1, and the instruction is not atomic.

Note

To refer to an FSLx interface in assembly language, use rfsl0, rfsl1, ... rfsl15.

The blocking versions of this instruction should not be placed in a delay slot when the parameter C_USE_EXTENDED_FSL_INSTR is set to 1, since this prevents interrupts from being served.

For non-blocking versions, an rsubc instruction can be used to decrement an index variable.

The ‘e’ bit does not have any effect unless C_FSL_EXCEPTION is set to 1.

These instructions are only available when the MicroBlaze parameter C_FSL_LINKS is greater than 0.

The extended instructions (exception, test and atomic versions) are only available when the MicroBlaze parameter C_USE_EXTENDED_FSL_INSTR is set to 1.

getd**get from fsl interface dynamic**

| | | |
|-----------------|--------|---|
| <i>tmeagetd</i> | rD, rB | get data from FSL rB[28:31] <i>t</i> = test-only <i>n</i> = non-blocking <i>e</i> = exception if control bit set <i>a</i> = atomic |
| <i>tmeagetd</i> | rD, rB | get control from FSL rB[28:31] <i>t</i> = test-only <i>n</i> = non-blocking <i>e</i> = exception if control bit not set <i>a</i> = atomic |

| | | | | | |
|--------------------|-----------|------------------|-----------|--------------------------------|----|
| 0 1 0 0 1 1 | rD | 0 0 0 0 0 | rB | 0 n c t a e 0 0 0 0 0 0 | |
| 0 | 6 | 11 | 16 | 21 | 31 |

Description

MicroBlaze will read from the FSL interface defined by the four least significant bits in rB and place the result in register rD.

The getd instruction has 32 variants.

The blocking versions (when ‘n’ bit is ‘0’) will stall MicroBlaze until the data from the FSL interface is valid. The non-blocking versions will not stall MicroBlaze and will set carry to ‘0’ if the data was valid and to ‘1’ if the data was invalid. In case of an invalid access the destination register contents is undefined.

All data get instructions (when ‘c’ bit is ‘0’) expect the control bit from the FSL interface to be ‘0’. If this is not the case, the instruction will set MSR[FSL_Error] to ‘1’. All control get instructions (when ‘c’ bit is ‘1’) expect the control bit from the FSL interface to be ‘1’. If this is not the case, the instruction will set MSR[FSL_Error] to ‘1’.

The exception versions (when ‘e’ bit is ‘1’) will generate an exception if there is a control bit mismatch. In this case ESR is updated with EC set to the exception cause and ESS set to the FSL index. The target register, rD, is not updated when an exception is generated, instead the FSL data is stored in EDR.

The test versions (when ‘t’ bit is ‘1’) will be handled as the normal case, except that the read signal to the FSL link is not asserted.

Atomic versions (when ‘a’ bit is ‘1’) are not interruptible. This means that a sequence of atomic FSL instructions can be grouped together without an interrupt breaking the program flow. However, note that exceptions may still occur.

When MicroBlaze is configured to use an MMU (`C_USE_MMU >= 1`) these instructions are privileged. This means that if these instructions are attempted in User Mode (MSR[UM] = 1) a Privileged Instruction exception occurs.

Pseudocode

```
if MSR[UM] = 1 then
    ESR[EC] ← 00111
else
    (rD) ← (FSL rB[28:31])
    if (n = 1) then
        MSR[Carry] ← not (FSL rB[28:31] Exists bit)
    if (FSL rB[28:31] Control bit ≠ c) and (FSL rB[28:31] Exists bit) then
        MSR[FSL_Error] ← 1
    if (e = 1) then
        ESR[EC] ← 00000
        ESR[ESS] ← rB[28:31]
        EDR ← (FSL rB[28:31])
```

Registers Altered

- rD, unless an exception is generated, in which case the register is unchanged
- MSR[FSL_Error]
- MSR[Carry]
- ESR[EC], in case an FSL exception or a privileged instruction exception is generated
- ESR[ESS], in case an FSL exception is generated
- EDR, in case an FSL exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

The blocking versions of this instruction will stall the pipeline of MicroBlaze until the instruction can be completed. Interrupts are served unless the instruction is atomic, which ensures that the instruction cannot be interrupted.

Note

The blocking versions of this instruction should not be placed in a delay slot, since this prevents interrupts from being served.

For non-blocking versions, an rsubc instruction can be used to decrement an index variable.

The ‘e’ bit does not have any effect unless C_FSL_EXCEPTION is set to 1.

These instructions are only available when the MicroBlaze parameter C_FSL_LINKS is greater than 0 and the parameter C_USE_EXTENDED_FSL_INSTR is set to 1.

idiv

Integer Divide

| | | |
|-------|------------|----------------------------|
| idiv | rD, rA, rB | divide rB by rA (signed) |
| idivu | rD, rA, rB | divide rB by rA (unsigned) |

Description

The contents of register rB is divided by the contents of register rA and the result is placed into register rD.

If the U bit is set, rA and rB are considered unsigned values. If the U bit is clear, rA and rB are considered signed values.

If the value of rA is 0, the DZO bit in MSR will be set and the value in rD will be 0, unless an exception is generated.

If the U bit is clear, the value of rA is -1, and the value of rB is -2147483648, the DZO bit in MSR will be set and the value in rD will be -2147483648, unless an exception is generated.

Pseudocode

```

if (rA) = 0 then
  (rD) <- 0
  MSR[DZO] <- 1
  ESR[EC] <- 00101
  ESR[DEC] <- 0
else if U = 0 and (rA) = -1 and (rB) = -2147483648 then
  (rD) <- -2147483648
  MSR[DZO] <- 1
  ESR[EC] <- 00101
  ESR[DEC] <- 1
else
  (rD) ← (rB) / (rA)

```

Registers Altered

- rD, unless a divide exception is generated, in which case the register is unchanged
 - MSR[DZO], if the value in rA is zero
 - ESR[EC], if the value in rA is zero

Latency

- 1 cycle if $(rA) = 0$, otherwise 32 cycles with $C_AREA_OPTIMIZED=0$
 - 1 cycle if $(rA) = 0$, otherwise 34 cycles with $C_AREA_OPTIMIZED=1$

Note

This instruction is only valid if MicroBlaze is configured to use a hardware divider (C_USE_DIV = 1).

imm**Immediate**

imm IMM

| | | | | IMM |
|---|---|---|---|------------|
| 1 | 0 | 1 | 1 | 0 0 |
| 0 | 6 | 1 | 1 | 3 |

0 6 1 1 3
 1 6 1

Description

The instruction imm loads the IMM value into a temporary register. It also locks this value so it can be used by the following instruction and form a 32-bit immediate value.

The instruction imm is used in conjunction with Type B instructions. Since Type B instructions have only a 16-bit immediate value field, a 32-bit immediate value cannot be used directly. However, 32-bit immediate values can be used in MicroBlaze. By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. The imm instruction locks the 16-bit IMM value temporarily for the next instruction. A Type B instruction that immediately follows the imm instruction will then form a 32-bit immediate value from the 16-bit IMM value of the imm instruction (upper 16 bits) and its own 16-bit immediate value field (lower 16 bits). If no Type B instruction follows the imm instruction, the locked value gets unlocked and becomes useless.

Latency

- 1 cycle

Notes

The imm instruction and the Type B instruction following it are atomic; consequently, no interrupts are allowed between them.

The assembler provided by Xilinx automatically detects the need for imm instructions. When a 32-bit IMM value is specified in a Type B instruction, the assembler converts the IMM value to a 16-bit one to assemble the instruction and inserts an imm instruction before it in the executable file.

Ibu**Load Byte Unsigned**

Ibu rD, rA, rB

| | | | | |
|--------------------|-----------|-----------|-----------|--|
| 1 1 0 0 0 0 | rD | rA | rB | 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 |
| 0 | 6 | 11 | 16 | 21 31 |

Description

Loads a byte (8 bits) from the memory location that results from adding the contents of registers rA and rB. The data is placed in the least significant byte of register rD and the other three bytes in rD are cleared.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB.

A data storage exception occurs if access is prevented by a no-access-allowed zone protection. This only applies to accesses with user mode and virtual protected mode enabled.

Pseudocode

```

Addr ← (rA) + (rB)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 0
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[UM] = 1 and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 0; ESR[DIZ] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else
    (rD)[24:31] ← Mem(Addr)
    (rD)[0:23] ← 0

```

Registers Altered

- rD, unless an exception is generated, in which case the register is unchanged
- MSR[UM], MSR[VM], MSR[UMS], MSR[VMS], if an exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

Ibui**Load Byte Unsigned Immediate**

Ibui rD, rA, IMM

| | | | | |
|--------------------|-----------|-----------|------------|----|
| 1 1 1 0 0 0 | rD | rA | IMM | |
| 0 | 6 | 11 | 16 | 31 |

Description

Loads a byte (8 bits) from the memory location that results from adding the contents of register rA with the value in IMM, sign-extended to 32 bits. The data is placed in the least significant byte of register rD and the other three bytes in rD are cleared.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB.

A data storage exception occurs if access is prevented by a no-access-allowed zone protection. This only applies to accesses with user mode and virtual protected mode enabled.

Pseudocode

```

Addr ← (rA) + sext(IMM)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 0
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[UM] = 1 and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 0; ESR[DIZ] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else
    (rD)[24:31] ← Mem(Addr)
    (rD)[0:23] ← 0

```

Registers Altered

- rD, unless an exception is generated, in which case the register is unchanged
- MSR[UM], MSR[VM], MSR[UMS], MSR[VMS], if an exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

Ihu

Load Halfword Unsigned

Ihu rD, rA, rB

| | | | | |
|--------------------|-----------|-----------|-----------|--|
| 1 1 0 0 0 1 | rD | rA | rB | 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 |
| 0 | 6 | 11 | 16 | 21 |

Description

Loads a halfword (16 bits) from the halfword aligned memory location that results from adding the contents of registers rA and rB. The data is placed in the least significant halfword of register rD and the most significant halfword in rD is cleared.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB.

A data storage exception occurs if access is prevented by a no-access-allowed zone protection. This only applies to accesses with user mode and virtual protected mode enabled.

An unaligned data access exception occurs if the least significant bit in the address is not zero.

Pseudocode

```

Addr ← (rA) + (rB)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 0
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[UM] = 1 and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 0; ESR[DIZ] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Addr[31] ≠ 0 then
    ESR[EC] ← 00001; ESR[W] ← 0; ESR[S] ← 0; ESR[Rx] ← rD
else
    (rD)[16:31] ← Mem(Addr)
    (rD)[0:15] ← 0

```

Registers Altered

- rD, unless an exception is generated, in which case the register is unchanged
 - MSR[UM], MSR[VM], MSR[UMS], MSR[VMS], if a TLB miss exception or a data storage exception is generated
 - ESR[EC], ESR[S], if an exception is generated
 - ESR[DIZ], if a data storage exception is generated
 - ESR[W], ESR[Rx], if an unaligned data access exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
 - 2 cycles with C_AREA_OPTIMIZED=1

Ihui**Load Halfword Unsigned Immediate**

Ihui rD, rA, IMM

| 1 | 1 | 1 | 0 | 0 | 1 | rD | rA | IMM | |
|---|---|---|---|---|----|----|----|-----|----|
| 0 | | | 6 | | 11 | | 16 | | 31 |

Description

Loads a halfword (16 bits) from the halfword aligned memory location that results from adding the contents of register rA and the value in IMM, sign-extended to 32 bits. The data is placed in the least significant halfword of register rD and the most significant halfword in rD is cleared.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB. A data storage exception occurs if access is prevented by a no-access-allowed zone protection. This only applies to accesses with user mode and virtual protected mode enabled. An unaligned data access exception occurs if the least significant bit in the address is not zero.

Pseudocode

```

Addr ← (rA) + sext(IMM)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 0
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[UM] = 1 and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 0; ESR[DIZ] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Addr[31] ≠ 0 then
    ESR[EC] ← 00001; ESR[W] ← 0; ESR[S] ← 0; ESR[Rx] ← rD
else
    (rD)[16:31] ← Mem(Addr)
    (rD)[0:15] ← 0

```

Registers Altered

- rD, unless an exception is generated, in which case the register is unchanged
- MSR[UM], MSR[VM], MSR[Ums], MSR[Vms], if a TLB miss exception or a data storage exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated
- ESR[W], ESR[Rx], if an unaligned data access exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

Iw**Load Word****Iw** rD, rA, rB

| | | | | |
|--------------------|-----------|-----------|-----------|--|
| 1 1 0 0 1 0 | rD | rA | rB | 0 0 |
| 0 | 6 | 11 | 16 | 21 |

Description

Loads a word (32 bits) from the word aligned memory location that results from adding the contents of registers rA and rB. The data is placed in register rD.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB.

A data storage exception occurs if access is prevented by a no-access-allowed zone protection. This only applies to accesses with user mode and virtual protected mode enabled.

An unaligned data access exception occurs if the two least significant bits in the address are not zero.

Pseudocode

```

Addr ← (rA) + (rB)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 0
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[UM] = 1 and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 0; ESR[DIZ] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Addr[30:31] ≠ 0 then
    ESR[EC] ← 00001; ESR[W] ← 1; ESR[S] ← 0; ESR[Rx] ← rD
else
    (rD) ← Mem(Addr)

```

Registers Altered

- rD, unless an exception is generated, in which case the register is unchanged
- MSR[UM], MSR[VM], MSR[Ums], MSR[Vms], if a TLB miss exception or a data storage exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated
- ESR[W], ESR[Rx], if an unaligned data access exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

lwi**Load Word Immediate**

lwi rD, rA, IMM

| | | | | |
|--------------------|-----------|-----------|------------|----|
| 1 1 1 0 1 0 | rD | rA | IMM | |
| 0 | 6 | 11 | 16 | 31 |

Description

Loads a word (32 bits) from the word aligned memory location that results from adding the contents of register rA and the value IMM, sign-extended to 32 bits. The data is placed in register rD. A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB. A data storage exception occurs if access is prevented by a no-access-allowed zone protection. This only applies to accesses with user mode and virtual protected mode enabled. An unaligned data access exception occurs if the two least significant bits in the address are not zero.

Pseudocode

```

Addr ← (rA) + sext(IMM)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 0
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[UM] = 1 and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 0; ESR[DIZ] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Addr[30:31] ≠ 0 then
    ESR[EC] ← 00001; ESR[W] ← 1; ESR[S] ← 0; ESR[Rx] ← rD
else
    (rD) ← Mem(Addr)

```

Registers Altered

- rD, unless an exception is generated, in which case the register is unchanged
- MSR[UM], MSR[VM], MSR[Ums], MSR[Vms], if a TLB miss exception or a data storage exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated
- ESR[W], ESR[Rx], if an unaligned data access exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

lwx

Load Word Exclusive

Iwx rD, rA, rB

Description

Loads a word (32 bits) from the word aligned memory location that results from adding the contents of registers rA and rB. The data is placed in register rD, the reservation bit is set, and the carry flag (MSR[C]) is cleared.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB.

A data storage exception occurs if access is prevented by a no-access-allowed zone protection. This only applies to accesses with user mode and virtual protected mode enabled.

An unaligned data access exception will *not* occur, even if the two least significant bits in the address are not zero.

Pseudocode

```

Addr ← (rA) + (rB)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 0
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[UM] = 1 and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 0; ESR[DIZ] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else
    (rD) ← Mem(Addr); Reservation ← 1; MSR[C] ← 0

```

Registers Altered

- rD and MSR[C], unless an exception is generated, in which case they are unchanged
 - MSR[UM], MSR[VM], MSR[UMS], MSR[VMS], if a TLB miss exception or a data storage exception is generated
 - ESR[EC], ESR[S], if an exception is generated
 - ESR[DIZ], if a data storage exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
 - 2 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is used together with SWX to implement exclusive access, such as semaphores and spinlocks.

mfs**Move From Special Purpose Register**

mfs rD, rS

| | | | | | | | | | | | | | | |
|---|---|---|---|---|---|----|---|---|----|---|----|---|---|----|
| 1 | 0 | 0 | 1 | 0 | 1 | rD | 0 | 0 | 0 | 0 | 0 | 1 | 0 | rS |
| 0 | | | 6 | | | 11 | | | 16 | | 18 | | | 31 |

Description

Copies the contents of the special purpose register rS into register rD. The special purpose registers TLBLO and TLBHI are used to copy the contents of the Unified TLB entry indexed by TLBX.

Pseudocode

```
switch (rS):
    case 0x0000 : (rD) ← PC
    case 0x0001 : (rD) ← MSR
    case 0x0003 : (rD) ← EAR
    case 0x0005 : (rD) ← ESR
    case 0x0007 : (rD) ← FSR
    case 0x000B : (rD) ← BTR
    case 0x000D : (rD) ← EDR
    case 0x1000 : (rD) ← PID
    case 0x1001 : (rD) ← ZPR
    case 0x1002 : (rD) ← TLBX
    case 0x1003 : (rD) ← TLBLO
    case 0x1004 : (rD) ← TLBHI
    case 0x200x : (rD) ← PVR[x] (where x = 0 to 11)
    default : (rD) ← Undefined
```

Registers Altered

- rD

Latency

- 1 cycle

Notes

To refer to special purpose registers in assembly language, use rpc for PC, rmsr for MSR, rear for EAR, resr for ESR, rfsr for FSR, rbtr for BTR, redr for EDR, rpid for PID, rzpr for ZPR, rtlblo for TLBLO, rtlbhi for TLBHI, rtlbx for TLBX, and rpvr0 - rpvr11 for PVR0 - PVR11.

The value read from MSR may not include effects of the immediately preceding instruction (dependent on pipeline stall behavior). An instruction that does not affect MSR must precede the MFS instruction to guarantee correct MSR value.

The value read from FSR may not include effects of the immediately preceding instruction (dependent on pipeline stall behavior). An instruction that does not affect FSR must precede the MFS instruction to guarantee correct FSR value.

EAR, ESR and BTR are only valid as operands when at least one of the MicroBlaze C_*_EXCEPTION parameters are set to 1.

EDR is only valid as operand when the parameter C_FSL_EXCEPTION is set to 1 and the parameter C_FSL_LINKS is greater than 0.

FSR is only valid as an operand when the C_USE_FPU parameter is greater than 0.

PID, ZPR, TLBLO and TLBHI are only valid as operands when the parameter C_USE_MMU > 1 and the parameter C_MMU_TLB_ACCESS = 1 or 3.

TLBX is only valid as operand when the parameter C_USE_MMU > 1 and the parameter C_MMU_TLB_ACCESS > 0.

PVR0 is only valid as an operand when C_PVR is 1 or 2, and PVR1 - PVR11 are only valid as operands when C_PVR is set to 2.

msrclr**Read MSR and clear bits in MSR****msrclr** rD, Imm

| | | | |
|--------------------|-----------|--------------------|--------------|
| 1 0 0 1 0 1 | rD | 1 0 0 0 1 0 | Imm15 |
| 0 | 6 | 11 | 16 17 |

31

Description

Copies the contents of the special purpose register MSR into register rD.

Bit positions in the IMM value that are 1 are cleared in the MSR. Bit positions that are 0 in the IMM value are left untouched.

When MicroBlaze is configured to use an MMU (`C_USE_MMU >= 1`) this instruction is privileged for all IMM values except those only affecting C. This means that if the instruction is attempted in User Mode (MSR [UM] = 1) in this case a Privileged Instruction exception occurs.

Pseudocode

```

if MSR[UM] = 1 and IMM ≠ 0x4 then
    ESR[EC] ← 00111
else
    (rD) ← (MSR)
    (MSR) ← (MSR) ∧ (IMM)

```

Registers Altered

- rD
- MSR
- ESR[EC], in case a privileged instruction exception is generated

Latency

- 1 cycle

Note

MSRCLR will affect the Carry bit immediately while the remaining bits will take effect one cycle after the instruction has been executed. When clearing the IE bit, it is guaranteed that the processor will not react to any interrupt for the subsequent instructions.

The value read from MSR may not include effects of the immediately preceding instruction (dependent on pipeline stall behavior). An instruction that does not affect MSR must precede the MSRCLR instruction to guarantee correct MSR value.

The immediate values has to be less than 2^{15} when `C_USE_MMU >= 1`, and less than 2^{14} otherwise. Only bits 17 to 31 of the MSR can be cleared when `C_USE_MMU >= 1`, and bits 18 to 31 otherwise.

This instruction is only available when the parameter `C_USE_MSR_INSTR` is set to 1.

When clearing MSR[VM] the instruction must always be followed by a synchronizing branch instruction, for example BRI 4.

msrset**Read MSR and set bits in MSR****msrset rD, Imm**

| | | | | |
|--------------------|-----------|--------------------|--------------|----|
| 1 0 0 1 0 1 | rD | 1 0 0 0 0 0 | Imm15 | |
| 0 | 6 | 11 | 16 17 | 31 |

Description

Copies the contents of the special purpose register MSR into register rD.

Bit positions in the IMM value that are 1 are set in the MSR. Bit positions that are 0 in the IMM value are left untouched.

When MicroBlaze is configured to use an MMU ($C_USE_MMU \geq 1$) this instruction is privileged for all IMM values except those only affecting C. This means that if the instruction is attempted in User Mode (MSR [UM] = 1) in this case a Privileged Instruction exception occurs.

Pseudocode

```

if MSR[UM] = 1 and IMM ≠ 0x4 then
    ESR[EC] ← 00111
else
    (rD) ← (MSR)
    (MSR) ← (MSR) ∨ (IMM)

```

Registers Altered

- rD
- MSR
- ESR[EC], in case a privileged instruction exception is generated

Latency

- 1 cycle

Note

MSRSET will affect the Carry bit immediately while the remaining bits will take effect one cycle after the instruction has been executed. When setting the EIP or BIP bit, it is guaranteed that the processor will not react to any interrupt or normal hardware break for the subsequent instructions.

The value read from MSR may not include effects of the immediately preceding instruction (dependent on pipeline stall behavior). An instruction that does not affect MSR must precede the MSRSET instruction to guarantee correct MSR value.

The immediate values has to be less than 2^{15} when $C_USE_MMU \geq 1$, and less than 2^{14} otherwise. Only bits 17 to 31 of the MSR can be set when $C_USE_MMU \geq 1$, and bits 18 to 31 otherwise.

This instruction is only available when the parameter $C_USE_MSR_INSTR$ is set to 1.

When setting MSR[VM] the instruction must always be followed by a synchronizing branch instruction, for example BRI 4.

mts**Move To Special Purpose Register**

mts rS, rA

| | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|----|---|---|----|----|----|---|---|--|--|----|--|
| 1 | 0 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | rA | 1 | 1 | | | rS | |
| 0 | | | | 6 | | 11 | | | 16 | 18 | | | | | | 31 | |

Description

Copies the contents of register rD into the special purpose register rS. The special purpose registers TLBLO and TLBHI are used to copy to the Unified TLB entry indexed by TLBX.

When MicroBlaze is configured to use an MMU ($C_USE_MMU \geq 1$) this instruction is privileged. This means that if the instruction is attempted in User Mode (MSR [UM] = 1) a Privileged Instruction exception occurs.

Pseudocode

```
if MSR[UM] = 1 then
    ESR[EC] ← 00111
else
    switch (rS)
        case 0x0001 : MSR ← (rA)
        case 0x0007 : FSR ← (rA)
        case 0x1000 : PID ← (rA)
        case 0x1001 : ZPR ← (rA)
        case 0x1002 : TLBX ← (rA)
        case 0x1003 : TLBLO ← (rA)
        case 0x1004 : TLBHI ← (rA)
        case 0x1005 : TLBSX ← (rA)
```

Registers Altered

- rS
- ESR[EC], in case a privileged instruction exception is generated

Latency

- 1 cycle

Notes

When writing MSR using MTS, all bits take effect one cycle after the instruction has been executed. An MTS instruction writing MSR should never be followed back-to-back by an instruction that uses the MSR content. When clearing the IE bit, it is guaranteed that the processor will not react to any interrupt for the subsequent instructions. When setting the EIP or BIP bit, it is guaranteed that the processor will not react to any interrupt or normal hardware break for the subsequent instructions.

To refer to special purpose registers in assembly language, use rmsr for MSR, rfsr for FSR, rpid for PID, rzpr for ZPR, rtlblo for TLBLO, rtlbhi for TLBHI, rtlbx for TLBX, and rtlbsx for TLBSX.

The PC, ESR, EAR, BTR, EDR and PVR0 - PVR11 cannot be written by the MTS instruction.

The FSR is only valid as a destination if the MicroBlaze parameter C_USE_FPU is greater than 0.

PID, ZPR and TLBSX are only valid as destinations when the parameter $C_USE_MMU > 1$ and the parameter $C_MMU_TLB_ACCESS > 1$. TLBLO, TLBHI and TLBX are only valid as destinations when the parameter $C_USE_MMU > 1$.

When changing MSR[VM] or PID the instruction must always be followed by a synchronizing branch instruction, for example BRI 4.

mul

Multiply

mul rD, rA, rB

Description

Multiplies the contents of registers rA and rB and puts the result in register rD. This is a 32-bit by 32-bit multiplication that will produce a 64-bit result. The least significant word of this value is placed in rD. The most significant word is discarded.

Pseudocode

(rD) \leftarrow LSW((rA) \times (rB))

Registers Altered

- rD

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
 - 3 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is only valid if the target architecture has multiplier primitives, and if present, the MicroBlaze parameter C USE HW MUL is greater than 0.

mulh**Multiply High****mulh rD, rA, rB**

| 0 1 0 0 0 0 | rD | rA | rB | 0 0 0 0 0 0 0 0 0 0 1 |
|--------------------|-----------|-----------|-----------|------------------------------|
| 0 | 6 | 1 | 1 | 2 1 3 1 |

Description

Multiplies the contents of registers rA and rB and puts the result in register rD. This is a 32-bit by 32-bit signed multiplication that will produce a 64-bit result. The most significant word of this value is placed in rD. The least significant word is discarded.

Pseudocode

```
(rD) ← MSW( (rA) × (rB) ), signed
```

Registers Altered

- rD

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 3 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is only valid if the target architecture has multiplier primitives, and if present, the MicroBlaze parameter C_USE_HW_MUL is set to 2.

When MULH is used, bit 30 and 31 in the MUL instruction must be zero to distinguish between the two instructions. In previous versions of MicroBlaze, these bits were defined as zero, but the actual values were not relevant.

mulhu

Multiply High Unsigned

mulhu rD, rA, rB

| 0 | 1 | 0 | 0 | 0 | 0 | rD | rA | rB | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 |
|---|---|---|---|---|---|----|----|----|---|---|---|---|---|---|---|---|---|---|---|
| 0 | | 6 | | 1 | | 1 | | 1 | 2 | | 3 | | 1 | | | | | | |

Description

Multiplies the contents of registers rA and rB and puts the result in register rD. This is a 32-bit by 32-bit unsigned multiplication that will produce a 64-bit unsigned result. The most significant word of this value is placed in rD. The least significant word is discarded.

Pseudocode

```
(rD) ← MSW( (rA) × (rB) ), unsigned
```

Registers Altered

- rD

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 3 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is only valid if the target architecture has multiplier primitives, and if present, the MicroBlaze parameter C_USE_HW_MUL is set to 2.

When MULHU is used, bit 30 and 31 in the MUL instruction must be zero to distinguish between the two instructions. In previous versions of MicroBlaze, these bits were defined as zero, but the actual values were not relevant.

mulhsu

Multiply High Signed Unsigned

mulhsu rD, rA, rB

| | | | | |
|--------------------|-----------|-----------|-----------|--------------------------------|
| 0 1 0 0 0 0 | rD | rA | rB | 0 0 0 0 0 0 0 0 0 0 1 0 |
| 0 | 6 | 1 | 1 | 2 |

Description

Multiplies the contents of registers rA and rB and puts the result in register rD. This is a 32-bit signed by 32-bit unsigned multiplication that will produce a 64-bit signed result. The most significant word of this value is placed in rD. The least significant word is discarded.

Pseudocode

(rD) \leftarrow MSW((rA), signed \times (rB), unsigned), signed

Registers Altered

- rD

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
 - 3 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is only valid if the target architecture has multiplier primitives, and if present, the MicroBlaze parameter C_USE_HW_MUL is set to 2.

When MULHSU is used, bit 30 and 31 in the MUL instruction must be zero to distinguish between the two instructions. In previous versions of MicroBlaze, these bits were defined as zero, but the actual values were not relevant.

muli**Multiply Immediate**

muli rD, rA, IMM

| 0 | 1 | 1 | 0 | 0 | 0 | rD | rA | IMM |
|---|---|---|---|---|---|----|----|-----|
| 0 | | | 6 | | | 1 | 1 | 3 |

1 6 1 1 6 1

Description

Multiplies the contents of registers rA and the value IMM, sign-extended to 32 bits; and puts the result in register rD. This is a 32-bit by 32-bit multiplication that will produce a 64-bit result. The least significant word of this value is placed in rD. The most significant word is discarded.

Pseudocode

```
(rD) ← LSW( (rA) × sext(IMM) )
```

Registers Altered

- rD

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 3 cycles with C_AREA_OPTIMIZED=1

Notes

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

This instruction is only valid if the target architecture has multiplier primitives, and if present, the MicroBlaze parameter C_USE_HW_MUL is greater than 0.

Or**Logical OR****or** rD, rA, rB

| 1 0 0 0 0 0 | rD | rA | rB | 0 0 0 0 0 0 0 0 0 0 0 0 | | | | | | | |
|-------------|----|----|----|-------------------------|---|---|---|---|---|---|---|
| 0 | 6 | 1 | 1 | 2 | 3 | 1 | 6 | 1 | 1 | 1 | 1 |

Description

The contents of register rA are ORed with the contents of register rB; the result is placed into register rD.

Pseudocode
$$(rD) \leftarrow (rA) \vee (rB)$$
Registers Altered

- rD

Latency

- 1 cycle

ori**Logical OR with Immediate**

ori rD, rA, IMM

| 1 | 0 | 1 | 0 | 0 | 0 | rD | rA | IMM |
|---|---|---|---|---|---|----|----|-----|
| 0 | | 6 | | 1 | | 1 | 1 | 6 |

Description

The contents of register rA are ORed with the extended IMM field, sign-extended to 32 bits; the result is placed into register rD.

Pseudocode

$(rD) \leftarrow (rA) \vee \text{sext(IMM)}$

Registers Altered

- rD

Latency

- 1 cycle

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

pcmpbf

Pattern Compare Byte Find

pcmpbf

rD, rA, rB

bitwise comparison returning position of first match

| | | | | |
|--------------------|-----------|-----------|-----------|----------------------------------|
| 1 0 0 0 0 0 | rD | rA | rB | 1 0 0 0 0 0 0 0 0 0 0 0 0 |
| 0 | 6 | 1 | 1 | 2 |

Description

The contents of register rA is bytewise compared with the contents in register rB.

- rD is loaded with the position of the first matching byte pair, starting with MSB as position 1, and comparing until LSB as position 4
 - If none of the byte pairs match, rD is set to 0

Pseudocode

```

if rB[0:7] = rA[0:7] then
  (rD) ← 1
else
  if rB[8:15] = rA[8:15] then
    (rD) ← 2
  else
    if rB[16:23] = rA[16:23] then
      (rD) ← 3
    else
      if rB[24:31] = rA[24:31] then
        (rD) ← 4
      else
        (rD) ← 0

```

Registers Altered

- rD

Latency

- 1 cycle

Note

This instruction is only available when the parameter C USE PCMP INSTR is set to 1.

pcmpeq

Pattern Compare Equal

pcmpaq

rD, rA, rB

equality comparison with a positive boolean result

| | | | | |
|--------------------|-----------|-----------|-----------|----------------------------------|
| 1 0 0 0 1 0 | rD | rA | rB | 1 0 0 0 0 0 0 0 0 0 0 0 0 |
| 0 | 6 | 1 | 1 | 2 |

Description

The contents of register rA is compared with the contents in register rB.

- rD is loaded with 1 if they match, and 0 if not

Pseudocode

```

if (rB) = (rA) then
  (rD) ← 1
else
  (rD) ← 0

```

Registers Altered

- rD

Latency

- 1 cycle

Note

This instruction is only available when the parameter C_USE_PCMP_INSTR is set to 1.

pcmpne

Pattern Compare Not Equal

pcmpne

rD, rA, rB

equality comparison with a negative boolean result

| | | | | |
|--------------------|-----------|-----------|-----------|----------------------------------|
| 1 0 0 0 1 1 | rD | rA | rB | 1 0 0 0 0 0 0 0 0 0 0 0 0 |
| 0 | 6 | 1 | 1 | 2 |

Description

The contents of register rA is compared with the contents in register rB.

- rD is loaded with 0 if they match, and 1 if not

Pseudocode

```

if (rB) = (rA) then
    (rD) ← 0
else
    (rD) ← 1

```

Registers Altered

- rD

Latency

- 1 cycle

Note

This instruction is only available when the parameter C USE PCMP INSTR is set to 1.

put

Put to fsl interface

| | | |
|-----------------|----------|--|
| <i>n</i> aput | rA, FSLx | put data to FSL x <i>n</i> = non-blocking <i>a</i> = atomic |
| <i>t</i> naput | FSLx | put data to FSL x test-only <i>n</i> = non-blocking <i>a</i> = atomic |
| <i>n</i> caput | rA, FSLx | put control to FSL x <i>n</i> = non-blocking <i>a</i> = atomic |
| <i>t</i> ncaput | FSLx | put control to FSL x test-only <i>n</i> = non-blocking <i>a</i> = atomic |

| | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|----|---|---|---|---|----|---|---|---|---|---|----|---|------|
| 0 | 1 | 1 | 0 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | rA | 1 | n | c | t | a | 0 | 0 | 0 | 0 | 0 | 0 | 0 | FSLx |
| 0 | | | | | | 6 | | | | | 11 | | | | | 16 | | | | | | 28 | | 31 |

Description

MicroBlaze will write the value from register rA to the FSLx interface.

The put instruction has 16 variants.

The blocking versions (when ‘n’ is ‘0’) will stall MicroBlaze until there is space available in the FSL interface. The non-blocking versions will not stall MicroBlaze and will set carry to ‘0’ if space was available and to ‘1’ if no space was available.

All data put instructions (when ‘c’ is ‘0’) will set the control bit to the FSL interface to ‘0’ and all control put instructions (when ‘c’ is ‘1’) will set the control bit to ‘1’.

The test versions (when ‘t’ bit is ‘1’) will be handled as the normal case, except that the write signal to the FSL link is not asserted (thus no source register is required).

Atomic versions (when ‘a’ bit is ‘1’) are not interruptible. This means that a sequence of atomic FSL instructions can be grouped together without an interrupt breaking the program flow. However, note that exceptions may still occur.

When MicroBlaze is configured to use an MMU (`C_USE_MMU` ≥ 1) these instructions are privileged. This means that if these instructions are attempted in User Mode (`MSR[UM]` = 1) a Privileged Instruction exception occurs.

Pseudocode

```

if MSR[UM] = 1 then
    ESR[EC] ← 00111
else
    (FSLx) ← (rA)
    if (n = 1) then
        MSR[Carry] ← (FSLx Full bit)
        (FSLx Control bit) ← C

```

Registers Altered

- MSR[Carry]
- ESR[EC], in case a privileged instruction exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

The blocking versions of this instruction will stall the pipeline of MicroBlaze until the instruction can be completed. Interrupts are served when the parameter C_USE_EXTENDED_FSL_INSTR is set to 1, and the instruction is not atomic.

Note

To refer to an FSLx interface in assembly language, use rfsl0, rfsl1, ... rfsl15.

The blocking versions of this instruction should not be placed in a delay slot when the parameter C_USE_EXTENDED_FSL_INSTR is set to 1, since this prevents interrupts from being served.

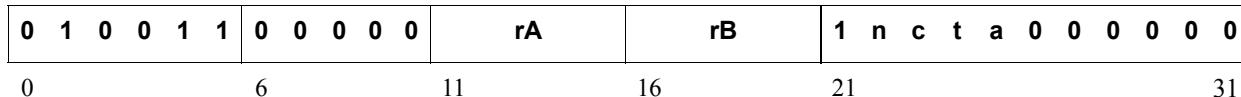
These instructions are only available when the MicroBlaze parameter C_FSL_LINKS is greater than 0.

The extended instructions (atomic versions) are only available when the MicroBlaze parameter C_USE_EXTENDED_FSL_INSTR is set to 1.

putd

Put to fsl interface dynamic

| | | |
|-----------------|--------|--|
| <i>naputd</i> | rA, rB | put data to FSL rB[28:31] <i>n</i> = non-blocking <i>a</i> = atomic |
| <i>tnaputd</i> | rB | put data to FSL rB[28:31] test-only <i>n</i> = non-blocking <i>a</i> = atomic |
| <i>ncaputd</i> | rA, rB | put control to FSL rB[28:31] <i>n</i> = non-blocking <i>a</i> = atomic |
| <i>tncaputd</i> | rB | put control to FSL rB[28:31] test-only <i>n</i> = non-blocking <i>a</i> = atomic |



Description

MicroBlaze will write the value from register rA to the FSL interface defined by the four least significant bits in rB.

The putd instruction has 16 variants.

The blocking versions (when ‘n’ is ‘0’) will stall MicroBlaze until there is space available in the FSL interface. The non-blocking versions will not stall MicroBlaze and will set carry to ‘0’ if space was available and to ‘1’ if no space was available.

All data putd instructions (when ‘c’ is ‘0’) will set the control bit to the FSL interface to ‘0’ and all control putd instructions (when ‘c’ is ‘1’) will set the control bit to ‘1’.

The test versions (when ‘t’ bit is ‘1’) will be handled as the normal case, except that the write signal to the FSL link is not asserted (thus no source register is required).

Atomic versions (when ‘a’ bit is ‘1’) are not interruptible. This means that a sequence of atomic FSL instructions can be grouped together without an interrupt breaking the program flow. However, note that exceptions may still occur.

When MicroBlaze is configured to use an MMU ($C_USE_MMU \geq 1$) these instructions are privileged. This means that if these instructions are attempted in User Mode (MSR[UM] = 1) a Privileged Instruction exception occurs.

Pseudocode

```

if MSR[UM] = 1 then
    ESR[EC] ← 00111
else
    (FSL rB[28:31]) ← (rA)
    if (n = 1) then
        MSR[Carry] ← (FSL rB[28:31] Full bit)
        (FSL rB[28:31] Control bit) ← C

```

Registers Altered

- MSR[Carry]
- ESR[EC], in case a privileged instruction exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

The blocking versions of this instruction will stall the pipeline of MicroBlaze until the instruction can be completed. Interrupts are served unless the instruction is atomic, which ensures that the instruction cannot be interrupted.

Note

The blocking versions of this instruction should not be placed in a delay slot, since this prevents interrupts from being served.

These instructions are only available when the MicroBlaze parameter C_FSL_LINKS is greater than 0 and the parameter C_USE_EXTENDED_FSL_INSTR is set to 1.

rsub

Arithmetic Reverse Subtract

| | | |
|--------|------------|------------------------------------|
| rsub | rD, rA, rB | Subtract |
| rsubc | rD, rA, rB | Subtract with Carry |
| rsubk | rD, rA, rB | Subtract and Keep Carry |
| rsubkc | rD, rA, rB | Subtract with Carry and Keep Carry |

| | | | | | | | | | | | | | | | | | |
|----------|----------|----------|----------|----------|----------|-----------|-----------|-----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|
| 0 | 0 | 0 | K | C | 1 | rD | rA | rB | 0 |
| 0 | | | | | | 6 | 1 | 1 | 2 | | | | | | | | 3 |

Description

The contents of register rA is subtracted from the contents of register rB and the result is placed into register rD. Bit 3 of the instruction (labeled as K in the figure) is set to one for the mnemonic rsubk. Bit 4 of the instruction (labeled as C in the figure) is set to one for the mnemonic rsubc. Both bits are set to one for the mnemonic rsubkc.

When an rsub instruction has bit 3 set (rsubk, rsubkc), the carry flag will keep its previous value regardless of the outcome of the execution of the instruction. If bit 3 is cleared (rsub, rsubc), then the carry flag will be affected by the execution of the instruction.

When bit 4 of the instruction is set to one (rsubc, rsubkc), the content of the carry flag (MSR[C]) affects the execution of the instruction. When bit 4 is cleared (rsub, rsubk), the content of the carry flag does not affect the execution of the instruction (providing a normal subtraction).

Pseudocode

```

if C = 0 then
    (rD) ← (rB) + (rA) + 1
else
    (rD) ← (rB) + (rA) + MSR [C]
if K = 0 then
    MSR [C] ← CarryOut

```

Registers Altered

- rD
 - MSR[C]

Latency

- 1 cycle

Notes

In subtractions, Carry = (Borrow). When the Carry is set by a subtraction, it means that there is no Borrow, and when the Carry is cleared, it means that there is a Borrow.

rsubi

Arithmetic Reverse Subtract Immediate

| | | |
|---------|-------------|--|
| rsubi | rD, rA, IMM | Subtract Immediate |
| rsubic | rD, rA, IMM | Subtract Immediate with Carry |
| rsubik | rD, rA, IMM | Subtract Immediate and Keep Carry |
| rsubikc | rD, rA, IMM | Subtract Immediate with Carry and Keep Carry |

| | | | | | | | | |
|----------|----------|----------|----------|----------|----------|----|----|-----|
| 0 | 0 | 1 | K | C | 1 | rD | rA | IMM |
| 0 | | | 6 | | | 1 | 1 | 3 |

Description

The contents of register rA is subtracted from the value of IMM, sign-extended to 32 bits, and the result is placed into register rD. Bit 3 of the instruction (labeled as K in the figure) is set to one for the mnemonic rsubik. Bit 4 of the instruction (labeled as C in the figure) is set to one for the mnemonic rsubic. Both bits are set to one for the mnemonic rsubikc.

When an rsubi instruction has bit 3 set (rsubik, rsubikc), the carry flag will keep its previous value regardless of the outcome of the execution of the instruction. If bit 3 is cleared (rsubi, rsubic), then the carry flag will be affected by the execution of the instruction. When bit 4 of the instruction is set to one (rsubic, rsubikc), the content of the carry flag (MSR[C]) affects the execution of the instruction. When bit 4 is cleared (rsubi, rsubik), the content of the carry flag does not affect the execution of the instruction (providing a normal subtraction).

Pseudocode

```

if C = 0 then
  (rD) ← sext( IMM ) + ( rA ) + 1
else
  (rD) ← sext( IMM ) + ( rA ) + MSR [C]
if K = 0 then
  MSR [C] ← CarryOut

```

Registers Altered

- rD
 - MSR[C]

Latency

- 1 cycle

Notes

In subtractions, Carry = (Borrow). When the Carry is set by a subtraction, it means that there is no Borrow, and when the Carry is cleared, it means that there is a Borrow. By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm.” page 161 for details on using 32-bit immediate values.

rtbd**Return from Break**

rtbd rA, IMM

| | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|----|--|-----|--|----|
| 1 | 0 | 1 | 1 | 0 | 1 | 1 | 0 | 0 | 1 | 0 | rA | | IMM | | |
| 0 | | | | | | 6 | | | | | 11 | | 16 | | 31 |

Description

Return from break will branch to the location specified by the contents of rA plus the IMM field, sign-extended to 32 bits. It will also enable breaks after execution by clearing the BIP flag in the MSR.

This instruction always has a delay slot. The instruction following the RTBD is always executed before the branch target. That delay slot instruction has breaks disabled.

When MicroBlaze is configured to use an MMU ($C_USE_MMU \geq 1$) this instruction is privileged. This means that if the instruction is attempted in User Mode (MSR [UM] = 1) a Privileged Instruction exception occurs.

Pseudocode

```

if MSR[UM] = 1 then
    ESR[EC] ← 00111
else
    PC ← (rA) + sext(IMM)
    allow following instruction to complete execution
    MSR[BIP] ← 0
    MSR[UM] ← MSR[UMS]
    MSR[VM] ← MSR[VMS]

```

Registers Altered

- PC
- MSR[BIP], MSR[UM], MSR[VM]
- ESR[EC], in case a privileged instruction exception is generated

Latency

- 2 cycles

Note

Convention is to use general purpose register r16 as rA.

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

rtid**Return from Interrupt**

rtid rA, IMM

| 1 | 0 | 1 | 1 | 0 | 1 | 1 | 0 | 0 | 0 | 1 | rA | IMM |
|---|---|---|---|---|---|---|---|---|---|----|----|-----|
| 0 | | | | | | 6 | | | | 11 | 16 | 31 |

Description

Return from interrupt will branch to the location specified by the contents of rA plus the IMM field, sign-extended to 32 bits. It will also enable interrupts after execution.

This instruction always has a delay slot. The instruction following the RTID is always executed before the branch target. That delay slot instruction has interrupts disabled.

When MicroBlaze is configured to use an MMU ($C_USE_MMU \geq 1$) this instruction is privileged. This means that if the instruction is attempted in User Mode (MSR [UM] = 1) a Privileged Instruction exception occurs.

Pseudocode

```

if MSR[UM] = 1 then
    ESR[EC] ← 00111
else
    PC ← (rA) + sext(IMM)
    allow following instruction to complete execution
    MSR[IE] ← 1
    MSR[UM] ← MSR[UMS]
    MSR[VM] ← MSR[VMS]

```

Registers Altered

- PC
- MSR[IE], MSR[UM], MSR[VM]
- ESR[EC], in case a privileged instruction exception is generated

Latency

- 2 cycles

Note

Convention is to use general purpose register r14 as rA.

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

rted**Return from Exception**

rted rA, IMM

| | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|----|---|---|----|----|--|--|-----|
| 1 | 0 | 1 | 1 | 0 | 1 | 1 | 0 | 1 | 0 | 0 | | rA | | | IMM |
| 0 | | | | | 6 | | | 11 | | | 16 | | | | 31 |

Description

Return from exception will branch to the location specified by the contents of rA plus the IMM field, sign-extended to 32 bits. The instruction will also enable exceptions after execution.

This instruction always has a delay slot. The instruction following the RTED is always executed before the branch target.

When MicroBlaze is configured to use an MMU ($C_USE_MMU \geq 1$) this instruction is privileged. This means that if the instruction is attempted in User Mode (MSR [UM] = 1) a Privileged Instruction exception occurs.

Pseudocode

```

if MSR[UM] = 1 then
    ESR[EC] ← 00111
else
    PC ← (rA) + sext(IMM)
    allow following instruction to complete execution
    MSR[EE] ← 1
    MSR[EIP] ← 0
    MSR[UM] ← MSR[UMS]
    MSR[VM] ← MSR[VMS]
    ESR ← 0

```

Registers Altered

- PC
- MSR[EE], MSR[EIP], MSR[UM], MSR[VM]
- ESR

Latency

- 2 cycles

Note

Convention is to use general purpose register r17 as rA. This instruction requires that one or more of the MicroBlaze parameters $C_*_EXCEPTION$ are set to 1 or that $C_USE_MMU > 0$.

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

The instruction should normally not be used when MSR[EE] is set, since if the instruction in the delay slot would cause an exception, the exception handler would be entered with exceptions enabled.

Note: Code returning from an exception must first check if MSR[DS] is set, and in that case return to the address in BTR.

rtsd**Return from Subroutine**

rtsd rA, IMM

| 1 | 0 | 1 | 1 | 0 | 1 | 1 | 0 | 0 | 0 | 0 | rA | | IMM |
|---|---|---|---|---|---|---|---|---|---|---|----|--|-----|
| 0 | | | | | | 6 | | | | | 1 | | 1 |

1
6
1
1
6
3
1

Description

Return from subroutine will branch to the location specified by the contents of rA plus the IMM field, sign-extended to 32 bits.

This instruction always has a delay slot. The instruction following the RTSD is always executed before the branch target.

Pseudocode

```
PC ← (rA) + sext(IMM)
allow following instruction to complete execution
```

Registers Altered

- PC

Latency

- 1 cycle (if successful branch prediction occurs)
- 2 cycles (with Branch Target Cache disabled)
- 3 cycles (if branch prediction mispredict occurs)

Note

Convention is to use general purpose register r15 as rA.

A delay slot must not be used by the following: imm, branch, or break instructions. Interrupts and external hardware breaks are deferred until after the delay slot branch has been completed.

sb**Store Byte**

sb rD, rA, rB

| | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|----|----|----|---|----|---|---|---|---|---|---|---|---|----|
| 1 | 1 | 0 | 1 | 0 | 0 | rD | rA | rB | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | | | 6 | | | 11 | | 16 | | 21 | | | | | | | | | 31 |

Description

Stores the contents of the least significant byte of register rD, into the memory location that results from adding the contents of registers rA and rB.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB.

A data storage exception occurs if virtual protected mode is enabled, and access is prevented by no-access-allowed or read-only zone protection. No-access-allowed can only occur in user mode.

Pseudocode

```

Addr ← (rA) + (rB)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 1; ESR[DIZ] ← No-access-allowed
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else
    Mem(Addr) ← (rD) [24:31]

```

Registers Altered

- MSR[UM], MSR[VM], MSR[UMS], MSR[VMS], if an exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

sbi**Store Byte Immediate**

sbi rD, rA, IMM

| | | | | | | | | |
|---|---|---|---|----|---|----|----|-----|
| 1 | 1 | 1 | 1 | 0 | 0 | rD | rA | IMM |
| 0 | | 6 | | 11 | | 16 | | 31 |

Description

Stores the contents of the least significant byte of register rD, into the memory location that results from adding the contents of register rA and the value IMM, sign-extended to 32 bits.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB.

A data storage exception occurs if virtual protected mode is enabled, and access is prevented by no-access-allowed or read-only zone protection. No-access-allowed can only occur in user mode.

Pseudocode

```

Addr ← (rA) + sext(IMM)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 1; ESR[DIZ] ← No-access-allowed
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else
    Mem(Addr) ← (rD)[24:31]

```

Registers Altered

- MSR[UM], MSR[VM], MSR[UMS], MSR[VMS], if an exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

sext16**Sign Extend Halfword****sext16** rD, rA

| | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|----------|----------|----------|----------|----------|----------|-----------|-----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|
| 1 | 0 | 0 | 1 | 0 | 0 | rD | rA | 0 | 1 |
| 0 | | | | 6 | | | 1 | | | | | | 1 | | | | | | | | | | | | | 3 | |

Description

This instruction sign-extends a halfword (16 bits) into a word (32 bits). Bit 16 in rA will be copied into bits 0-15 of rD. Bits 16-31 in rA will be copied into bits 16-31 of rD.

Pseudocode

```
(rD) [0:15] ← (rA) [16]
(rD) [16:31] ← (rA) [16:31]
```

Registers Altered

- rD

Latency

- 1 cycle

sext8

Sign Extend Byte

sext8 rD, rA

| | | | |
|--------------------|-----------|-----------|--|
| 1 0 0 1 0 0 | rD | rA | 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 |
| 0 | 6 | 1 1 | 1 6 1 |

Description

This instruction sign-extends a byte (8 bits) into a word (32 bits). Bit 24 in rA will be copied into bits 0-23 of rD. Bits 24-31 in rA will be copied into bits 24-31 of rD.

Pseudocode

(rD) [0:23] \leftarrow (rA) [24]
(rD) [24:31] \leftarrow (rA) [24:31]

Registers Altered

- rD

Latency

- 1 cycle

sh**Store Halfword**

sh rD, rA, rB

| | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|----|----|----|---|----|---|---|---|---|---|---|---|---|----|
| 1 | 1 | 0 | 1 | 0 | 1 | rD | rA | rB | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | | | 6 | | | 11 | | 16 | | 21 | | | | | | | | | 31 |

Description

Stores the contents of the least significant halfword of register rD, into the halfword aligned memory location that results from adding the contents of registers rA and rB.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB.

A data storage exception occurs if virtual protected mode is enabled, and access is prevented by no-access-allowed or read-only zone protection. No-access-allowed can only occur in user mode.

An unaligned data access exception occurs if the least significant bit in the address is not zero.

Pseudocode

```

Addr ← (rA) + (rB)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 1; ESR[DIZ] ← No-access-allowed
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Addr[31] ≠ 0 then
    ESR[EC] ← 00001; ESR[W] ← 0; ESR[S] ← 1; ESR[Rx] ← rD
else
    Mem(Addr) ← (rD)[16:31]

```

Registers Altered

- MSR[UM], MSR[VM], MSR[UMS], MSR[VMS], if a TLB miss exception or a data storage exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated
- ESR[W], ESR[Rx], if an unaligned data access exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

shi**Store Halfword Immediate**

shi rD, rA, IMM

| 1 | 1 | 1 | 1 | 0 | 1 | rD | rA | IMM |
|---|---|---|---|---|----|----|----|-----|
| 0 | | | 6 | | 11 | | 16 | 31 |

Description

Stores the contents of the least significant halfword of register rD, into the halfword aligned memory location that results from adding the contents of register rA and the value IMM, sign-extended to 32 bits.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB. A data storage exception occurs if virtual protected mode is enabled, and access is prevented by no-access-allowed or read-only zone protection. No-access-allowed can only occur in user mode. An unaligned data access exception occurs if the least significant bit in the address is not zero.

Pseudocode

```

Addr ← (rA) + sext(IMM)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 1; ESR[DIZ] ← No-access-allowed
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Addr[31] ≠ 0 then
    ESR[EC] ← 00001; ESR[W] ← 0; ESR[S] ← 1; ESR[Rx] ← rD
else
    Mem(Addr) ← (rD)[16:31]

```

Registers Altered

- MSR[UM], MSR[VM], MSR[UMS], MSR[VMS], if a TLB miss exception or a data storage exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated
- ESR[W], ESR[Rx], if an unaligned data access exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

sra**Shift Right Arithmetic**

sra rD, rA

| | | | |
|--------------------|-----------|-----------|--|
| 1 0 0 1 0 0 | rD | rA | 0 1 |
| 0 | 6 | 1 1 | 1 6 3 1 |

Description

Shifts arithmetically the contents of register rA, one bit to the right, and places the result in rD. The most significant bit of rA (that is, the sign bit) placed in the most significant bit of rD. The least significant bit coming out of the shift chain is placed in the Carry flag.

Pseudocode

```
(rD)[0] ← (rA)[0]
(rD)[1:31] ← (rA)[0:30]
MSR[C] ← (rA)[31]
```

Registers Altered

- rD
- MSR[C]

Latency

- 1 cycle

SRC**Shift Right with Carry**

src rD, rA

| | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|----|----|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 1 | 0 | 0 | rD | rA | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 1 |
| 0 | | | 6 | | | 1 | | 1 | | | | | | | | | | | | | 3 | | | | | 1 | |

Description

Shifts the contents of register rA, one bit to the right, and places the result in rD. The Carry flag is shifted in the shift chain and placed in the most significant bit of rD. The least significant bit coming out of the shift chain is placed in the Carry flag.

Pseudocode

```
(rD) [0] ← MSR[C]
(rD) [1:31] ← (rA) [0:30]
MSR[C] ← (rA) [31]
```

Registers Altered

- rD
- MSR[C]

Latency

- 1 cycle

srl**Shift Right Logical**

srl rD, rA

| | | | | |
|--------------------|-----------|-----------|--|---|
| 1 0 0 1 0 0 | rD | rA | 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 | |
| 0 | 6 | 1 | 1 | 3 |

1 6 1 1 6 1

Description

Shifts logically the contents of register rA, one bit to the right, and places the result in rD. A zero is shifted in the shift chain and placed in the most significant bit of rD. The least significant bit coming out of the shift chain is placed in the Carry flag.

Pseudocode

```
(rD) [0] ← 0
(rD) [1:31] ← (rA) [0:30]
MSR [C] ← (rA) [31]
```

Registers Altered

- rD
- MSR[C]

Latency

- 1 cycle

SW**Store Word**

sw rD, rA, rB

| | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|----|----|----|---|----|---|---|---|---|---|---|---|---|----|
| 1 | 1 | 0 | 1 | 1 | 0 | rD | rA | rB | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | | | 6 | | | 11 | | 16 | | 21 | | | | | | | | | 31 |

Description

Stores the contents of register rD, into the word aligned memory location that results from adding the contents of registers rA and rB.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB.

A data storage exception occurs if virtual protected mode is enabled, and access is prevented by no-access-allowed or read-only zone protection. No-access-allowed can only occur in user mode.

An unaligned data access exception occurs if the two least significant bits in the address are not zero.

Pseudocode

```

Addr ← (rA) + (rB)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 1; ESR[DIZ] ← No-access-allowed
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Addr[30:31] ≠ 0 then
    ESR[EC] ← 00001; ESR[W] ← 1; ESR[S] ← 1; ESR[Rx] ← rD
else
    Mem(Addr) ← (rD)[0:31]

```

Registers Altered

- MSR[UM], MSR[VM], MSR[UMS], MSR[VMS], if a TLB miss exception or a data storage exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated
- ESR[W], ESR[Rx], if an unaligned data access exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

SWI**Store Word Immediate****swi** rD, rA, IMM

| 1 | 1 | 1 | 1 | 1 | 0 | rD | rA | IMM | |
|---|---|---|---|---|---|----|----|-----|----|
| 0 | | | 6 | | | 11 | | 16 | 31 |

Description

Stores the contents of register rD, into the word aligned memory location that results from adding the contents of registers rA and the value IMM, sign-extended to 32 bits.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB.

A data storage exception occurs if virtual protected mode is enabled, and access is prevented by no-access-allowed or read-only zone protection. No-access-allowed can only occur in user mode.

An unaligned data access exception occurs if the two least significant bits in the address are not zero.

Pseudocode

```

Addr ← (rA) + sext(IMM)
if TLB_Miss(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10010; ESR[S] ← 1
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Access_Protected(Addr) and MSR[VM] = 1 then
    ESR[EC] ← 10000; ESR[S] ← 1; ESR[DIZ] ← No-access-allowed
    MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
else if Addr[30:31] ≠ 0 then
    ESR[EC] ← 00001; ESR[W] ← 1; ESR[S] ← 1; ESR[Rx] ← rD
else
    Mem(Addr) ← (rD)[0:31]

```

Register Altered

- MSR[UM], MSR[VM], MSR[UMS], MSR[VMS], if a TLB miss exception or a data storage exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated
- ESR[W], ESR[Rx], if an unaligned data access exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.

SWX**Store Word Exclusive**

swx rD, rA, rB

| | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|----|----|----|---|----|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|----|
| 1 | 1 | 0 | 1 | 1 | 0 | rD | rA | rB | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | | | 6 | | | 11 | | 16 | | 21 | | | | | | | | | | | | | | | | | 31 |

Description

Conditionally stores the contents of register rD when the reservation bit is set, into the word aligned memory location that results from adding the contents of registers rA and rB. The carry bit (MSR[C]) is set if the store does not occur, otherwise it is cleared. The reservation bit is cleared.

A data TLB miss exception occurs if virtual protected mode is enabled, and a valid translation entry corresponding to the address is not found in the TLB.

A data storage exception occurs if virtual protected mode is enabled, and access is prevented by no-access-allowed or read-only zone protection. No-access-allowed can only occur in user mode.

An unaligned data access exception will *not* occur even if the two least significant bits in the address are not zero.

Pseudocode

```

Addr ← (rA) + (rB)
if Reservation = 0 then
    MSR[C] ← 1
else
    if TLB_Miss(Addr) and MSR[VM] = 1 then
        ESR[EC]← 10010; ESR[S]← 1
        MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
    else if Access_Protected(Addr) and MSR[VM] = 1 then
        ESR[EC] ← 10000; ESR[S]← 1; ESR[DIZ] ← No-access-allowed
        MSR[UMS] ← MSR[UM]; MSR[VMS] ← MSR[VM]; MSR[UM] ← 0; MSR[VM] ← 0
    else
        Mem(Addr) ← (rD)[0:31]; Reservation ← 0; MSR[C] ← 0

```

Registers Altered

- MSR[C], unless an exception is generated
- MSR[UM], MSR[VM], MSR[UMS], MSR[VMS], if a TLB miss exception or a data storage exception is generated
- ESR[EC], ESR[S], if an exception is generated
- ESR[DIZ], if a data storage exception is generated

Latency

- 1 cycle with C_AREA_OPTIMIZED=0
- 2 cycles with C_AREA_OPTIMIZED=1

Note

This instruction is used together with LWX to implement exclusive access, such as semaphores and spinlocks .

wdc**Write to Data Cache**

wdc rA,rB
wdc.flush rA,rB
wdc.clear rA,rB

| | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|----|----|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | rA | rB | 0 | 0 | 0 | 0 | 1 | 1 | F | 0 | 1 | T | 0 |
| 0 | | | 6 | | | 1 | | | 1 | 6 | | | | | 2 | | 3 | | 7 | | 1 |

Description

Write into the data cache tag to invalidate or flush a cache line. The mnemonic wdc.flush is used to set the F bit, and wdc.clear is used to set the T bit.

When C_DCACHE_USE_WRITEBACK is set to 1, the instruction will flush the cache line and invalidate it if the F bit is set, otherwise it will only invalidate the cache line and discard any data that has not been written to memory. If the T bit is set, only a cache line with a matching address is invalidated. Register rA added with rB is the address of the affected cache line.

When C_DCACHE_USE_WRITEBACK is cleared to 0, the instruction will always invalidate the cache line. Register rA contains the address of the affected cache line, and the register rB value is not used.

When MicroBlaze is configured to use an MMU (C_USE_MMU >= 1) the instruction is privileged. This means that if the instruction is attempted in User Mode (MSR [UM] = 1) a Privileged Instruction exception occurs.

Pseudocode

```

if MSR[UM] = 1 then
    ESR[EC] ← 00111
else
    if C_DCACHE_USE_WRITEBACK = 1 then
        address ← (Ra) + (Rb)
    else
        address ← (Ra)
    if C_DCACHE_LINE_LEN = 4 then
        cacheline_mask ← (1 << log2(C_DCACHE_BYTE_SIZE) - 4) - 1
        cacheline ← (DCache Line)[(address >> 4) ^ cacheline_mask]
        cacheline_addr ← address & 0xffffffff
    if C_DCACHE_LINE_LEN = 8 then
        cacheline_mask ← (1 << log2(C_DCACHE_BYTE_SIZE) - 5) - 1
        cacheline ← (DCache Line)[(address >> 5) ^ cacheline_mask]
        cacheline_addr ← address & 0xfffffff0
    if F = 1 and cacheline.Dirty then
        for i = 0 .. C_DCACHE_LINE_LEN - 1 loop
            if cacheline.Valid[i] then
                Mem(cacheline_addr + i * 4) ← cacheline.Data[i]
    if T = 0 then
        cacheline.Tag ← 0
    else if cacheline.Address = cacheline_addr then
        cacheline.Tag ← 0

```

Registers Altered

- ESR[EC], in case a privileged instruction exception is generated

Latency

- 2 cycles for wdc and wdc.clear
- 2 + N cycles for wdc.flush, where N is the number of clock cycles required to flush the cache line to memory when necessary

Note

The wdc, wdc.flush and wdc.clear instructions are independent of data cache enable (MSR[DCE]), and can be used either with the data cache enabled or disabled.

The wdc.clear instruction is intended to invalidate a specific area in memory, for example a buffer to be written by a Direct Memory Access device. Using this instruction ensures that other cache lines are not inadvertently invalidated, erroneously discarding data that has not yet been written to memory.

The address of the affected cache line is always the physical address, independent of the parameter C_USE_MMU and whether the MMU is in virtual mode or real mode.

When using wdc.flush in a loop to flush the entire cache, the loop can be optimized by using Ra as the cache base address and Rb as the loop counter:

```
addik      r5,r0,C_DCACHE_BASEADDR
addik      r6,r0,C_DCACHE_BYTE_SIZE-C_DCACHE_LINE_LEN*4
loop:   wdc.flush r5,r6
        bgtid    r6,loop
        addik    r6,r6,-C_DCACHE_LINE_LEN*4
```

When using wdc.clear in a loop to invalidate a memory area in the cache, the loop can be optimized by using Ra as the memory area base address and Rb as the loop counter:

```
addik      r5,r0,memory_area_base_address
addik      r6,r0,memory_area_byte_size-C_DCACHE_LINE_LEN*4
loop:   wdc.clear r5,r6
        bgtid    r6,loop
        addik    r6,r6,-C_DCACHE_LINE_LEN*4
```

wic**Write to Instruction Cache**

wic rA,rB

| | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|----|----|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | rA | rB | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 1 | 0 | 0 | 0 |
| 0 | | | 6 | | | 1 | | | 1 | 6 | | | | | 3 | | | 1 | | | |

Description

Write into the instruction cache tag to invalidate a cache line. The register rB value is not used. Register rA contains the address of the affected cache line.

When MicroBlaze is configured to use an MMU (`C_USE_MMU >= 1`) this instruction is privileged. This means that if the instruction is attempted in User Mode (`MSR [UM] = 1`) a Privileged Instruction exception occurs.

Pseudocode

```

if MSR [UM] = 1 then
    ESR[EC] ← 00111
else
    if C_ICACHE_LINE_LEN = 4 then
        cacheline_mask ← (1 << log2(C_CACHE_BYTE_SIZE) - 4) - 1
        (ICache Line)[((Ra) >> 4) ^ cacheline_mask].Tag ← 0
    if C_ICACHE_LINE_LEN = 8 then
        cacheline_mask ← (1 << log2(C_CACHE_BYTE_SIZE) - 5) - 1
        (ICache Line)[((Ra) >> 5) ^ cacheline_mask].Tag ← 0

```

Registers Altered

- ESR[EC], in case a privileged instruction exception is generated

Latency

- 2 cycles

Note

The WIC instruction is independent of instruction cache enable (`MSR[ICE]`), and can be used either with the instruction cache enabled or disabled.

The address of the affected cache line is always the physical address, independent of the parameter `C_USE_MMU` and whether the MMU is in virtual mode or real mode.

XOR**Logical Exclusive OR****xor** rD, rA, rB

| 1 0 0 0 1 0 | rD | rA | rB | 0 0 0 0 0 0 0 0 0 0 0 0 | | | | | | | |
|-------------|----|----|----|-------------------------|---|---|--|--|--|--|--|
| 0 | 6 | 1 | 1 | 2 | 3 | 1 | | | | | |

Description

The contents of register rA are XORed with the contents of register rB; the result is placed into register rD.

Pseudocode
$$(rD) \leftarrow (rA) \oplus (rB)$$
Registers Altered

- rD

Latency

- 1 cycle

xori**Logical Exclusive OR with Immediate**

xori rD, rA, IMM

| 1 | 0 | 1 | 0 | 1 | 0 | rD | rA | IMM |
|---|---|---|---|---|---|----|----|-----|
| 0 | | 6 | | 1 | | 1 | 1 | 6 |

Description

The IMM field is extended to 32 bits by concatenating 16 0-bits on the left. The contents of register rA are XOR'ed with the extended IMM field; the result is placed into register rD.

Pseudocode

$(rD) \leftarrow (rA) \oplus \text{sext(IMM)}$

Registers Altered

- rD

Latency

- 1 cycle

Note

By default, Type B Instructions will take the 16-bit IMM field value and sign extend it to 32 bits to use as the immediate operand. This behavior can be overridden by preceding the Type B instruction with an imm instruction. See the instruction “imm,” page 161 for details on using 32-bit immediate values.